


# DEEP LEARNING CRASH COURSE

- Single Layer Perceptron
- Multiple Layer Perceptron
- Convolutional Neural Net



M.A. Nielsen. Neural Networks and Deep Learning, 2015  
<http://neuralnetworksanddeeplearning.com/>

# ARTIFICIAL INTELLIGENCE




1997: Deep Blue beats chess World Champion




2016: AlphaGo beats go world champion

# LINEAR CLASSIFICATION




$$f(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + \mathbf{b} \geq 0, \\ 0 & \text{otherwise.} \end{cases}$$

# SINGLE LAYER PERCEPTRON




$$f(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x} + \mathbf{b})$$

Step function




$\sigma$ :

Sigmoid function



# MULTILAYER PERCEPTRON



$$\mathbf{f}(\mathbf{x}) = \sigma(\mathbf{W} \cdot \mathbf{x} + \mathbf{B})$$

For each node  $j$  in layer  $l$ ,

$$a_j^l = \sigma \left( b_j^l + \sum_k w_{j,k}^l a_k^{l-1} \right),$$

where  $a$  is the activation of the node.

- The network can be trained to produce a desired output given a specific input.
- In practice, this means learning the  $b$  and  $w$  parameters by minimizing a loss function on a training set.
- Often done on GPUs, which are much faster.

# BINARY LOSS FUNCTION



In the binary case,

$$L(\mathbf{w}, \mathbf{b}) = -\frac{1}{N} \sum_1^N [y_n \log(\hat{y}_n) + (1 - y_n) \log(1 - \hat{y}_n)]$$

where  $\hat{y}_n = f_{\mathbf{w}, \mathbf{b}}(x_n)$ .

# MULTICLASS LOSS FUNCTION



In the multiclass case, the probability that input vector  $\mathbf{x}$  belongs to class  $i$  can be written as

$$P(Y = i|\mathbf{x}, \mathbf{w}, \mathbf{b}) = \frac{f_i(\mathbf{x})}{\sum_j f_j(\mathbf{x})}$$

The class assigned to vector  $\mathbf{x}$  is taken to be


$$\hat{y} = \arg \max_i P(Y = i|\mathbf{x}, \mathbf{w}, \mathbf{b})$$

Given a set of  $N$  training samples  $(\mathbf{x}_n, y_n)_{1 \leq n \leq N}$ , the loss function can be written as

$$L(\mathbf{w}, \mathbf{b}) = \sum_n \log(P(Y = y_n|\mathbf{x}_n, \mathbf{w}, \mathbf{b}))$$

→  $L$  is a differentiable function of  $\mathbf{w}$  and  $\mathbf{b}$  and can be optimised using *back propagation*, that is, gradient descent.


# MNIST



- The network takes as input 28x28 images represented as 784D vectors.
- The output is a 10D vector giving the probability of the image representing any of the 10 digits.
- There are 50'000 training pairs of images and the corresponding label, 10'000 validation pairs, and 5'000 testing pairs.




# DEEP LEARNING




- The descriptive power of the net increases with the number of layers.
- Since every neuron is connected to every other in adjacent layers, the number of parameters to be learned increases quickly.
- Does not take into account the specific topology of images.

# CONVOLUTIONAL LAYER




$$\sigma \left( b + \sum_{x=0}^{n_x} \sum_{y=0}^{n_y} w_{i,j} a_{i+x,j+y} \right)$$

# FEATURE MAPS




# POOLING LAYER

hidden neurons (output from feature map)




max-pooling units




- Reduce the number of inputs by replacing all activations in a neighbourhood by a single one.
- Can be thought as asking if a particular feature is present in that neighbourhood while ignoring the exact location.

# ADDING THE POOLING LAYERS




The output size is reduced by the pooling layers.

# ADDING A FULLY CONNECTED LAYER




- Each neuron in the final fully connected layer is connected to all neurons in the preceding one.
- Deep architecture with many parameters to learn but still far fewer than an equivalent multilayer perceptron.

# HAND POSE ESTIMATION




Input: Depth image.

Output: 3D pose vector.



Oberweger et al. , ICCV'15

# OPTIMIZATION



Network parameters are found by minimizing and objective function of the form


$$\min_{\mathbf{W}_l, \mathbf{B}_l} \sum_i ||\mathbf{F}(\mathbf{x}_i, \mathbf{W}_1, \dots, \mathbf{W}_L, \mathbf{b}_1, \dots, \mathbf{b}_L) - \mathbf{y}_i||^2$$

using


- stochastic gradient descent on mini-batches,
- dropout,
- hard example mining,
- .....



# FEATURE MAPS LEARNED FOR IMAGE CLASSIFICATION




First convolutional layer




Second convolutional layer

- Some of the convolutional masks seem very similar to oriented Gaussian or Gabor filters!
- Much ongoing work to better understand this.

# ROC HUNTING



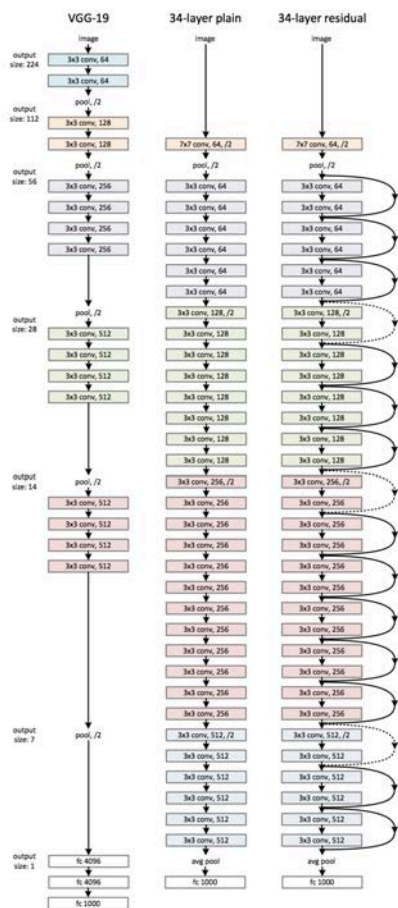


DeepFace  
Taigman et al. 2014





Deep Edge Detection  
Shen et al. 2015

# DEEPER AND DEEPER



He et al. , CVPR'16

# VISUAL CORTEX




# AlphaGo



- Uses Deep Nets to find the most promising locations to focus on.
- Performs Tree based search when possible.
- Relies on reinforcement learning and other ML techniques to train.



# ADVERSARIAL IMAGES



Szegedy et al. 2013

# IN SHORT



- Deep Belief Networks in general and Convolutional Neural Nets in particular outperform conventional Computer Vision algorithms on many benchmarks.
  - It is not fully understood why and unexpected failure cases have been demonstrated.
  - They require a lot of manual tuning to perform well and performance is hard to predict.
- > Many questions are still open and there is much work left to do.