

SHAPE FROM X

One image:

- **Shading**
- Texture

Two images or more:

- Stereo
- Contours
- Motion



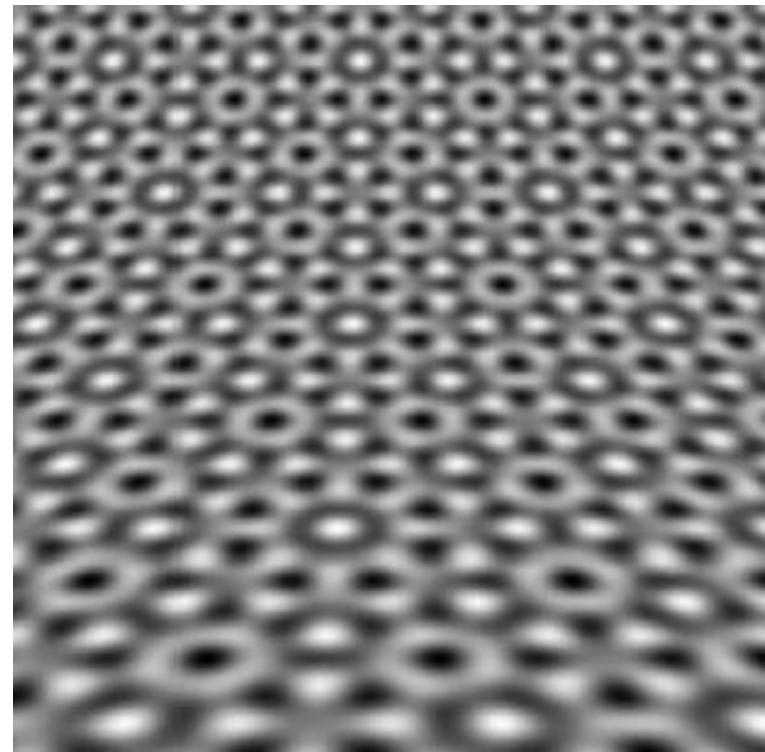
SHAPE FROM X

One image:

- Shading
- **Texture**

Two images or more:

- Stereo
- Contours
- Motion



SHAPE FROM X

One image:

- Shading
- Texture

Two images or more:

- **Stereo**
- Contours
- Motion



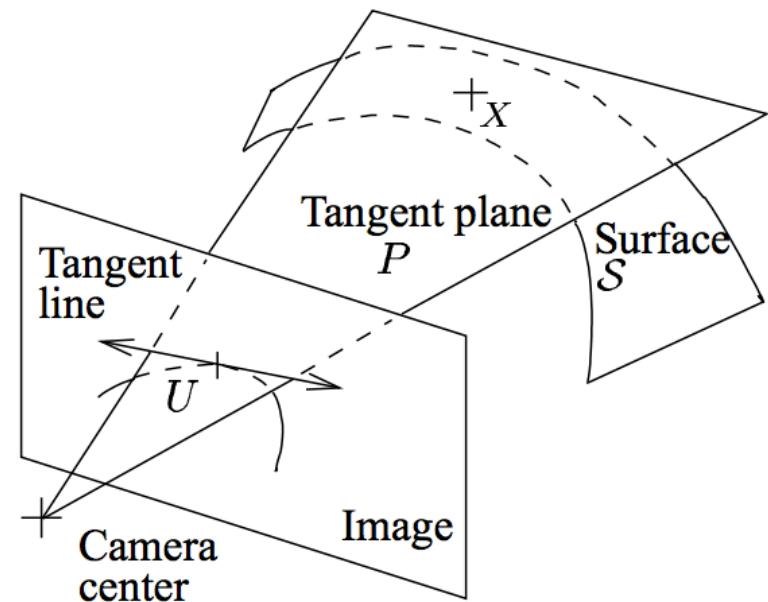
SHAPE FROM X

One image:

- Shading
- Texture

Two images or more:

- Stereo
- **Contours**
- Motion



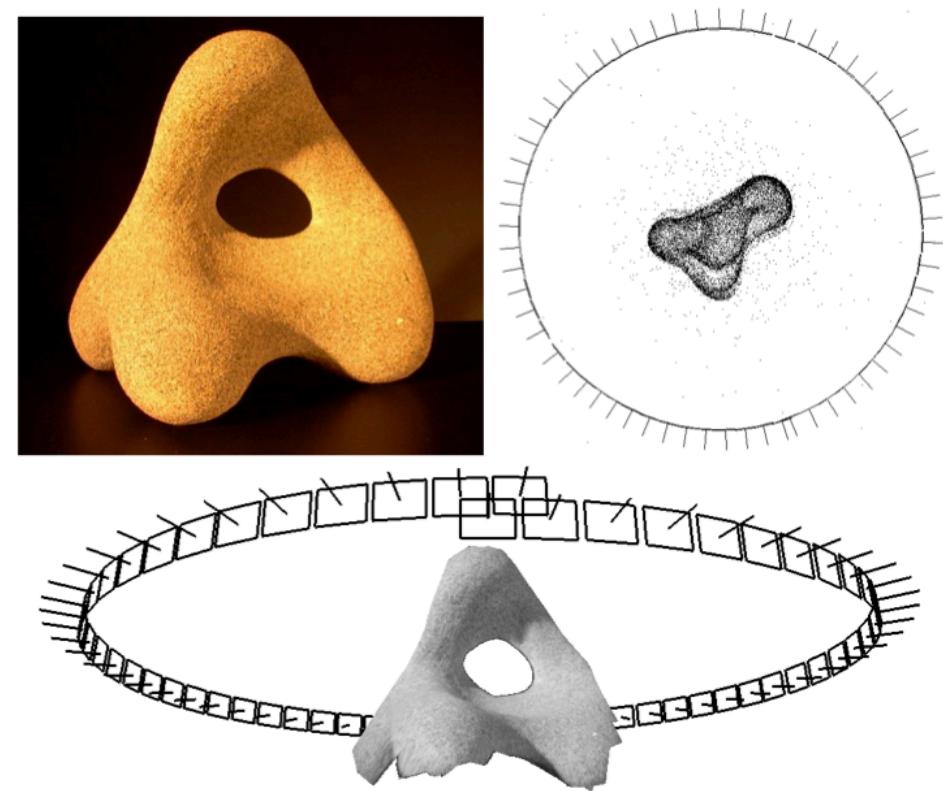
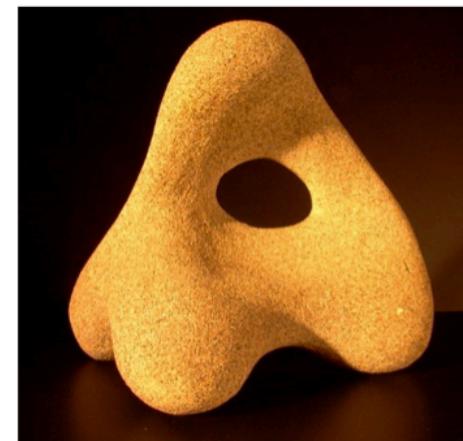
SHAPE FROM X

One image:

- Shading
- Texture

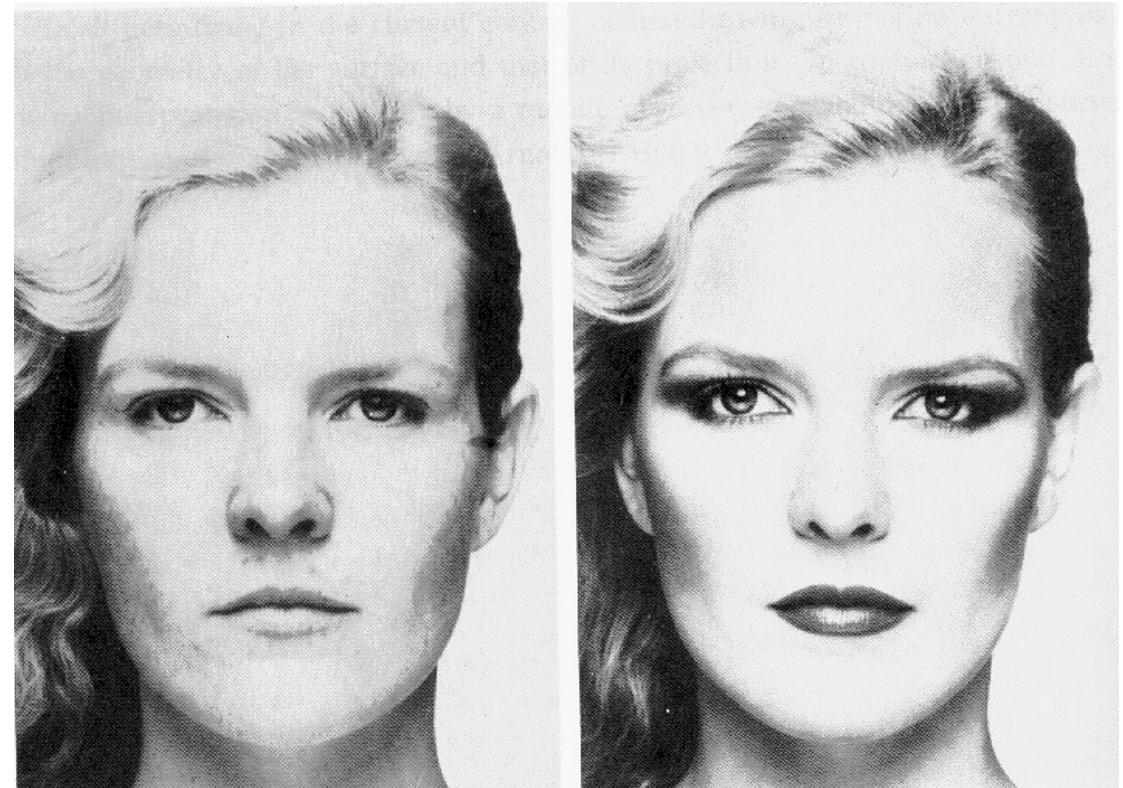
Two images or more:

- Stereo
- Contours
- Motion

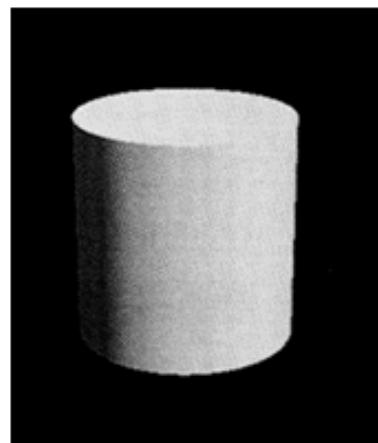
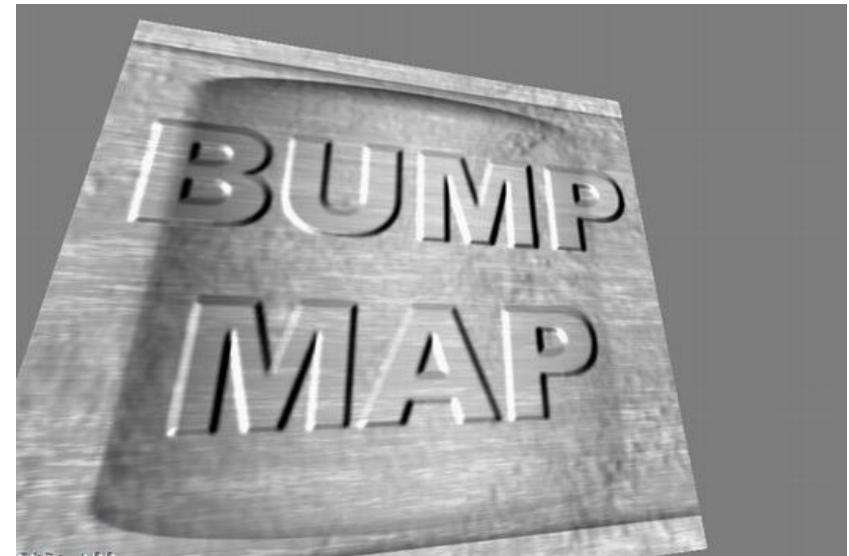
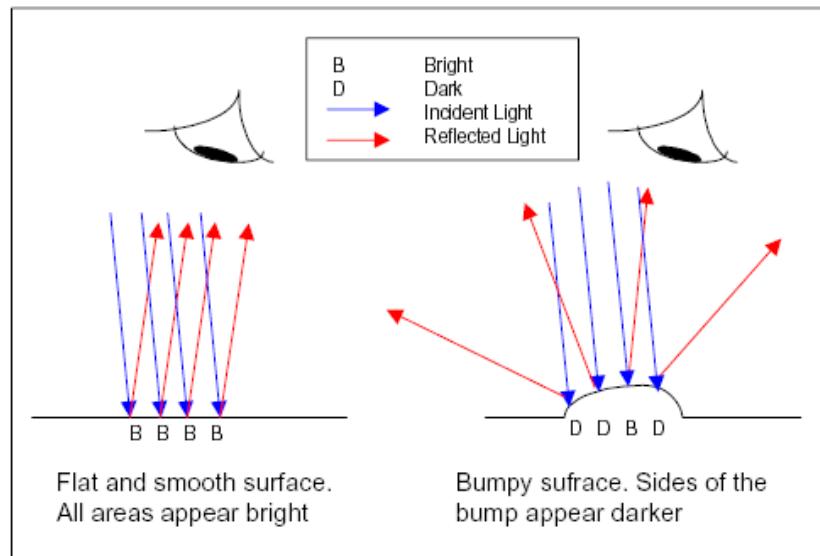


SHADING

- Shading models
- Shape from shading
 - Variational Methods
 - Photometric Stereo

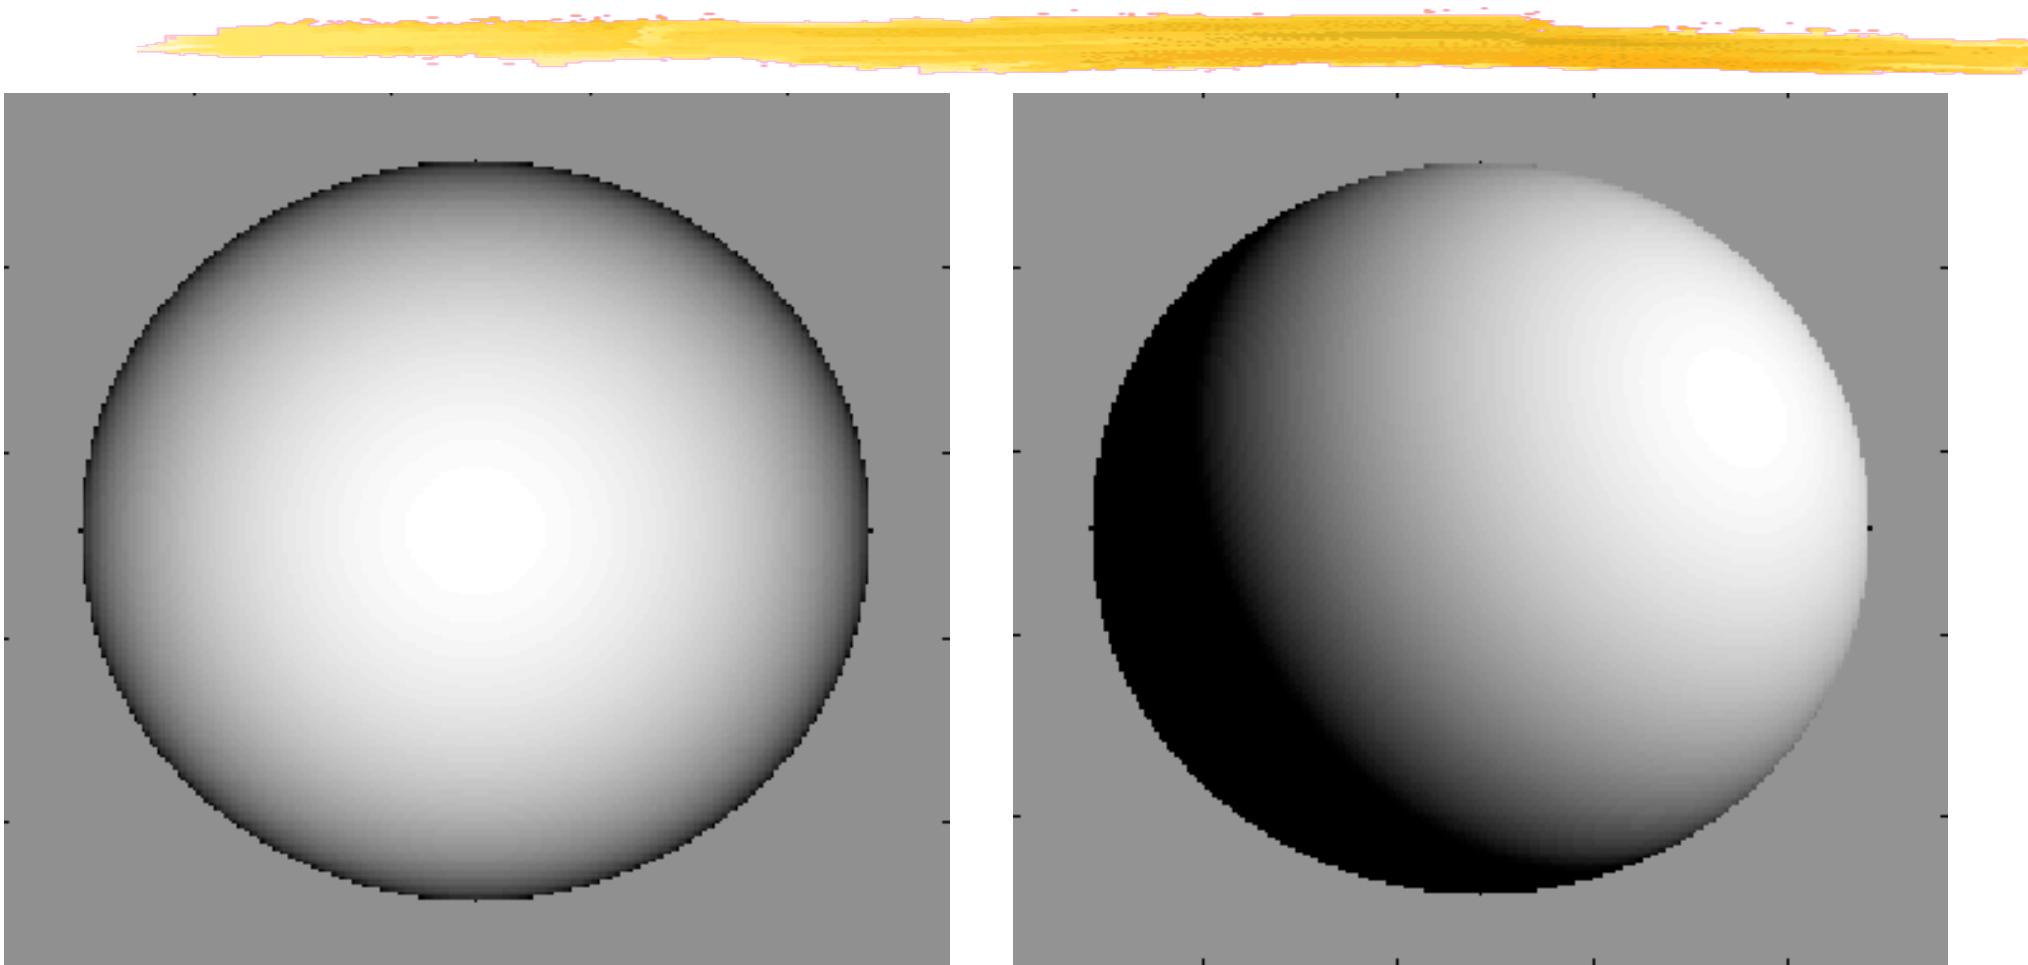


BUMP MAPPING



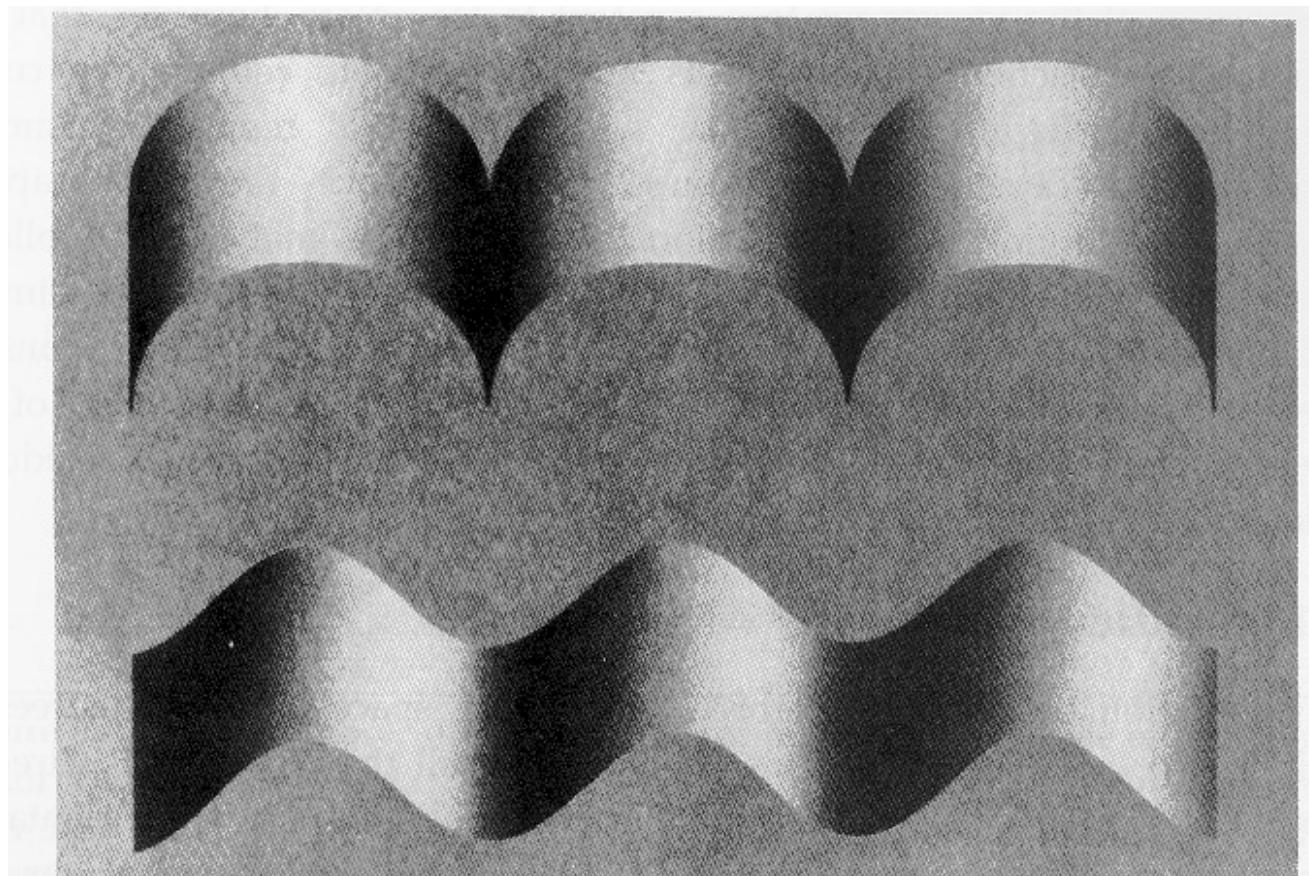
Simple mesh + 2D bump map = Complex looking object

LAMBERTIAN HALF-SPHERE



Gray level changes are interpreted as changes in the direction of the surface normal

BOUNDARY CONDITIONS



todo: what does the
boundary condition
here mean

Shading gives information about surface normals
→ Boundary conditions required for a complete interpretation.

BEETHOVEN

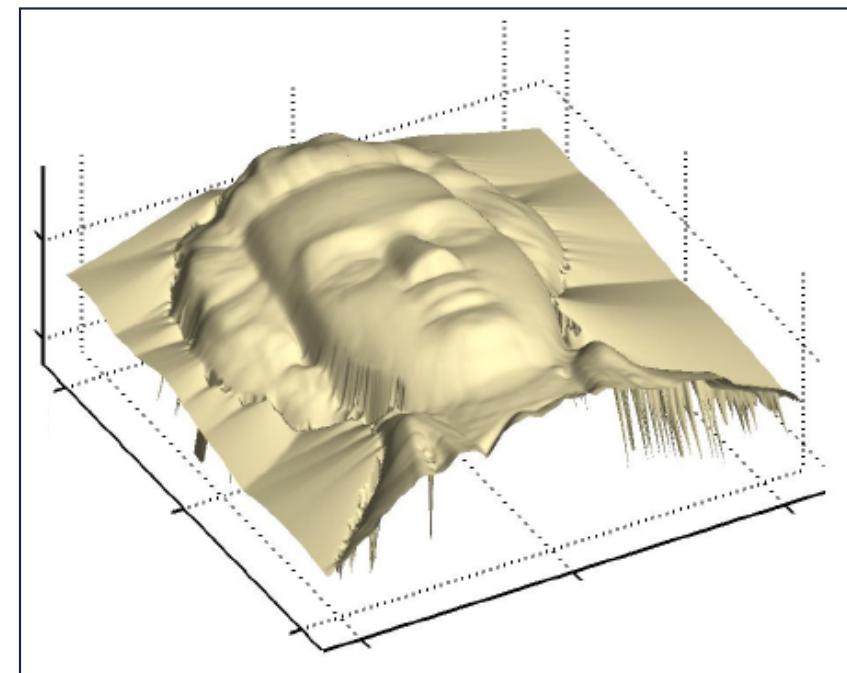
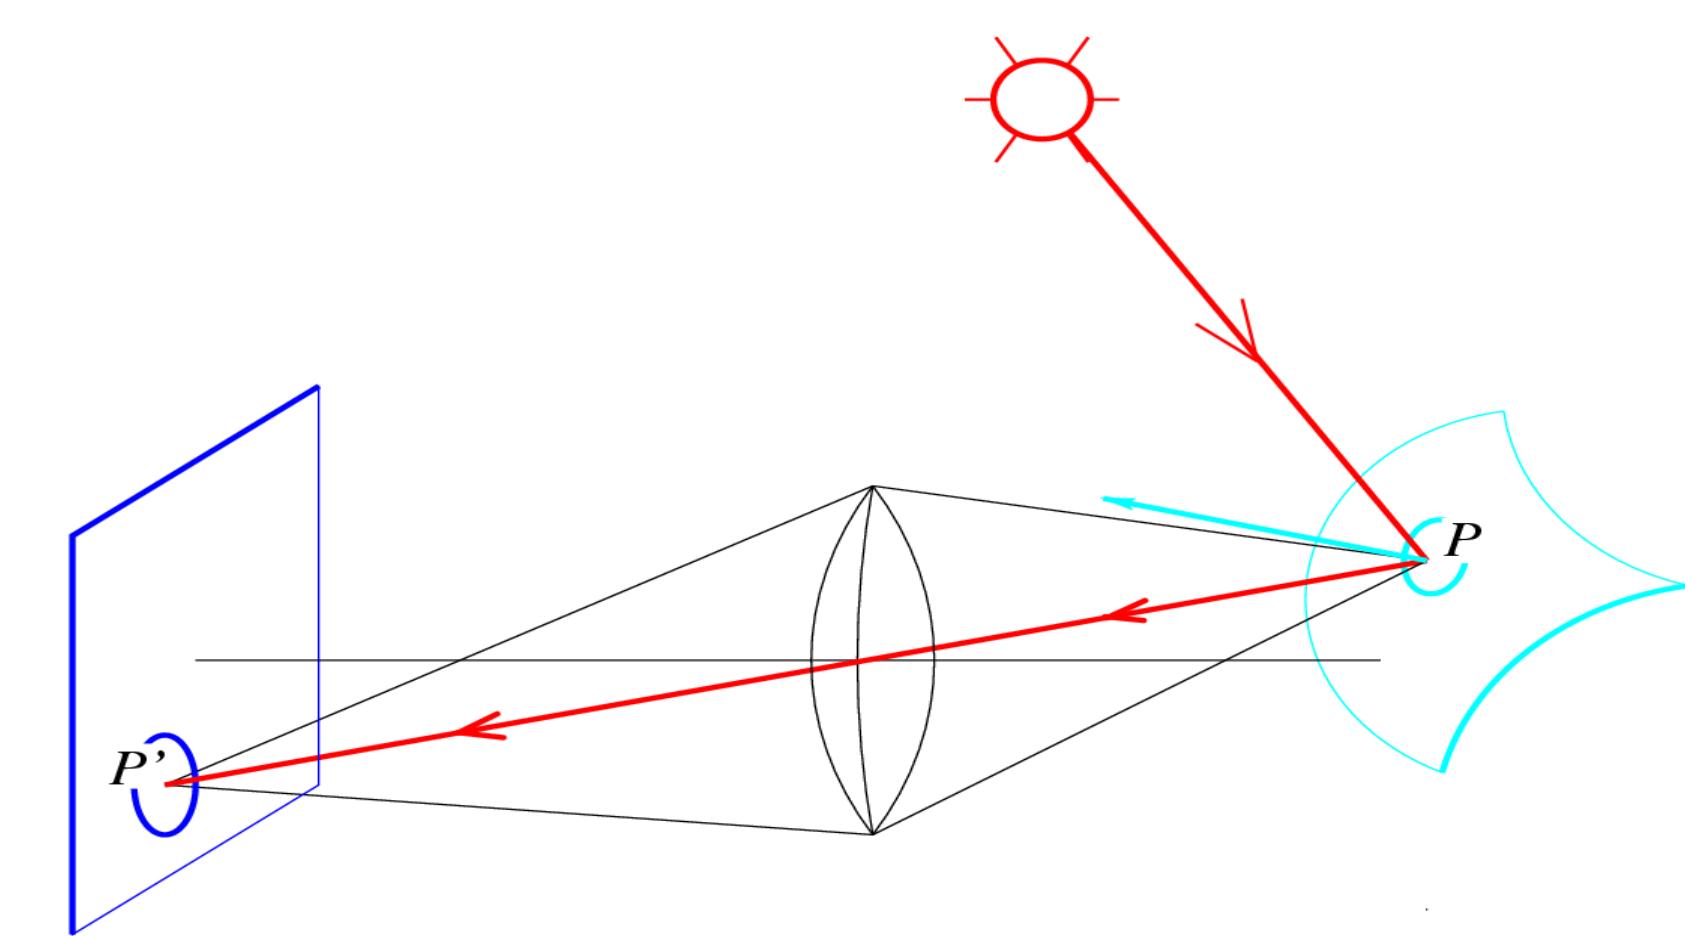
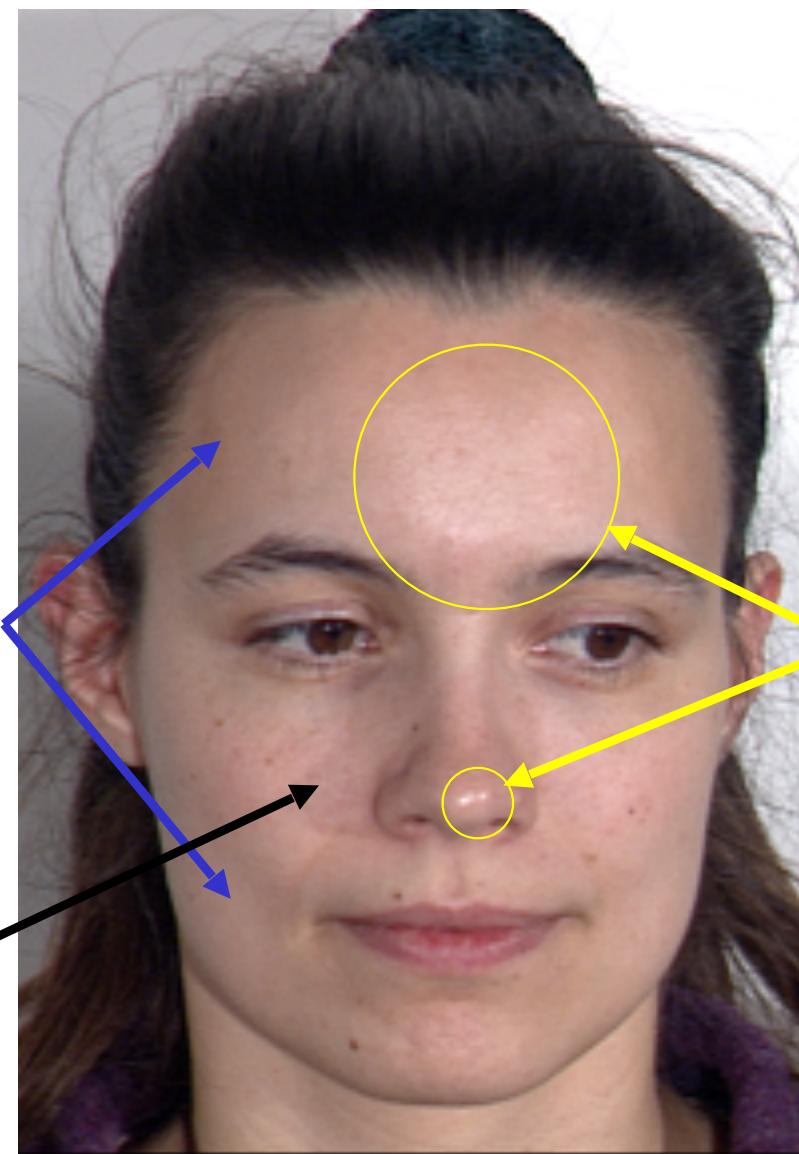


IMAGE FORMATION



DIRECT LIGHTING



Diffuse reflection

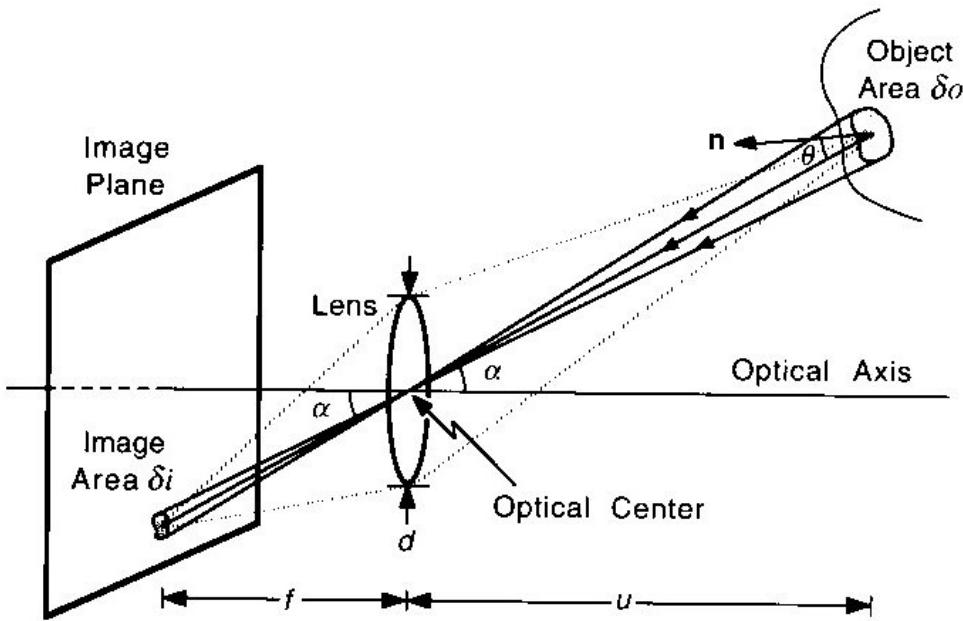
Shadow

Specularities

INDIRECT LIGHTING



FUNDAMENTAL RADIOMETRIC EQUATION



Radiance: amount of light emitted in a given direction
Irradiance: amount of light received on a surface

$$Irr = \frac{p}{4} \left(\frac{d}{f} \right)^2 \cos^4(a) Rad$$

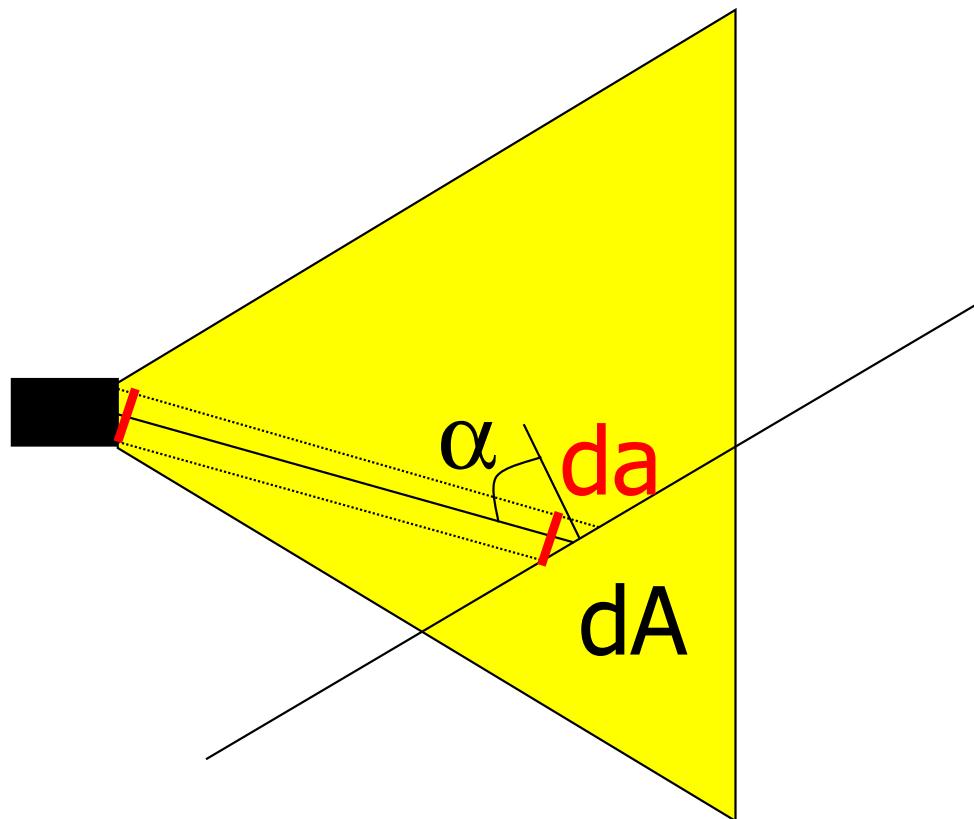
$$\Rightarrow I \propto Irr \propto Rad$$

Image intensity

Image Intensity is simply proportional to radiance and irradiance

when the camera is photometrically calibrated.

LOCAL SHADING MODEL

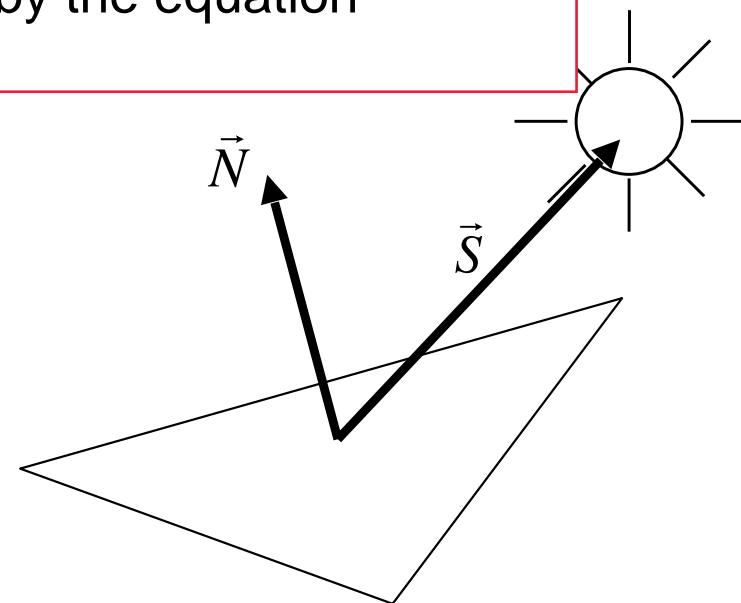


Foreshortening:
 $da = \cos(\alpha)dA$

The effect of a distant source on a surface patch depends on its **apparent surface**.

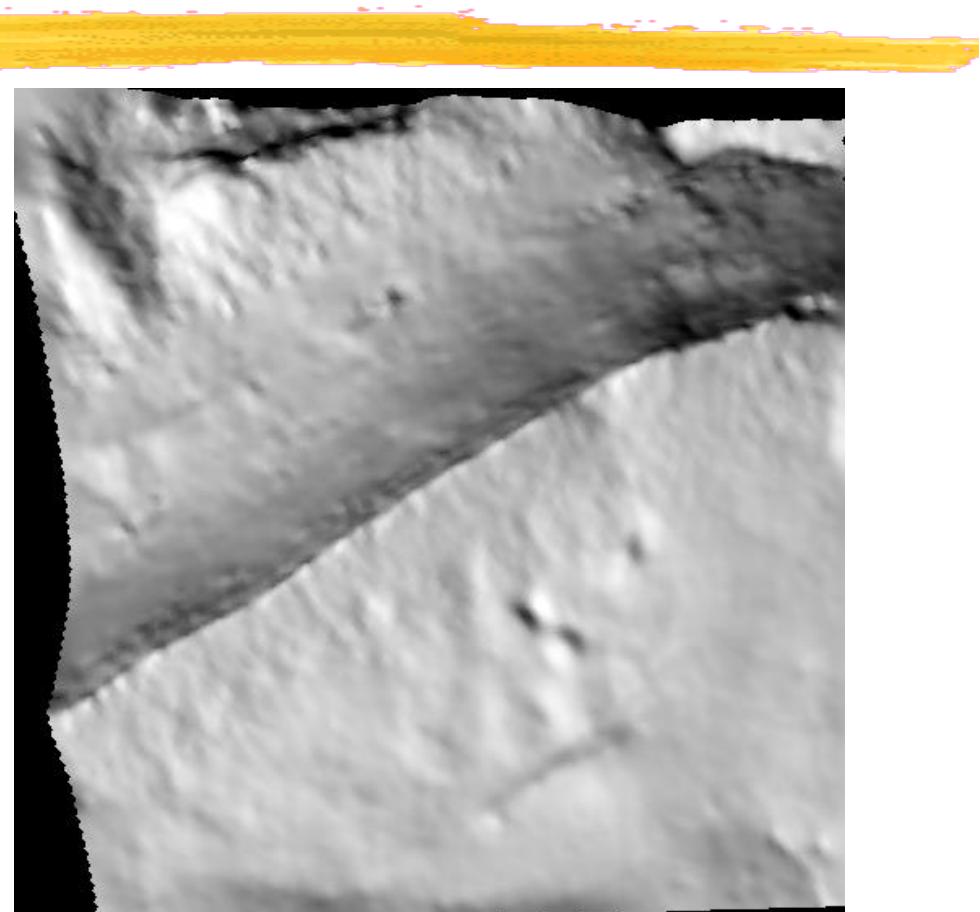
LAMBERTIAN SURFACE

Lambertian object means no specularity i.e. it is completely defined by the equation



$$I = \max(Albedo \cdot (\mathbf{N} \cdot \mathbf{S}), 0)$$

albedo = portion of light being reflected. 0 = black



Perfectly matte surface: The radiance depends only on angle of incidence and not on viewing direction.

ESTIMATED ALBEDO

PER PIXEL PARAMETERS

p - Albedo (1 unknown)

N - Surface Normal (2 unknowns since it has 2 DoF)

S_i - Directional Source (known)

$$I_i = p (N \cdot S_i)$$

$$I_i = S_i \cdot (pN)$$

$$I_i = S_i \cdot B$$

$$I_1(u,v) = | S_1 S_1 S_1 | | p(u,v) N_x(u,v) |$$

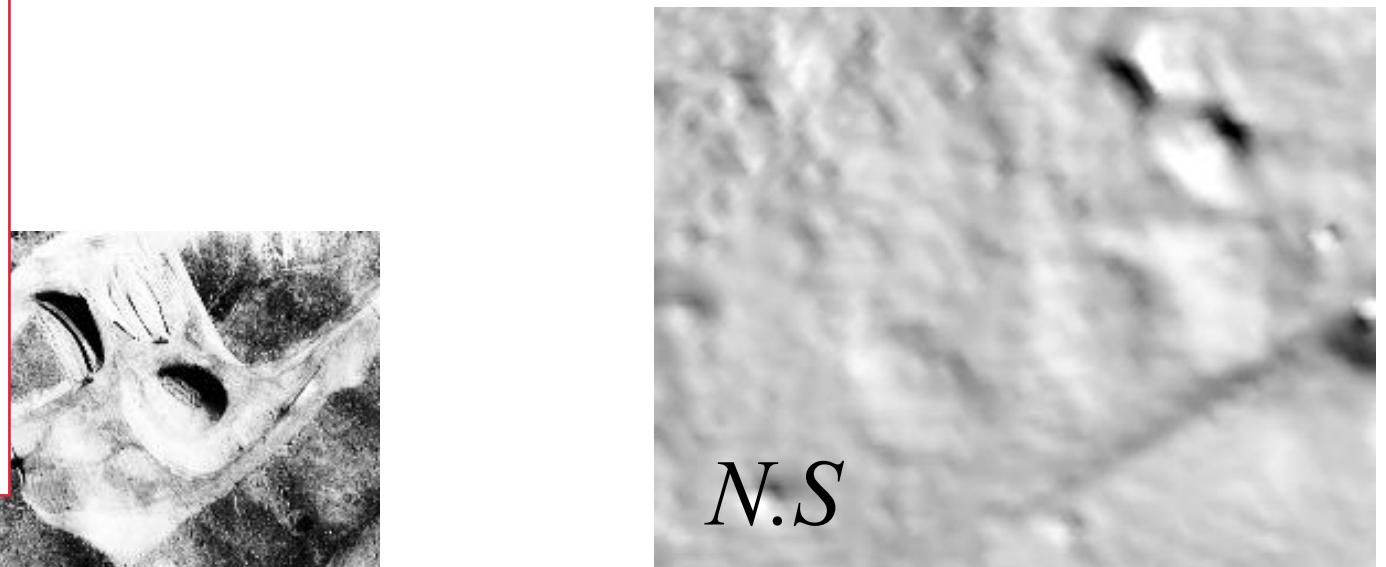
$$I_2(u,v) = | S_2 S_2 S_2 | | p(u,v) N_y(u,v) |$$

$$I_3(u,v) = | S_3 S_3 S_3 | | p(u,v) N_z(u,v) |$$

Solve via least squares with at least 3 equations. We can then extract the albedo and the surface normals.

Let $B(u,v) = p(u,v) N(u,v)$. Once it is solved, $N(u,v)$ is $B(u,v) / \| B(u,v) \|$ since $N(u,v)$ must be unit normal.

Then $p(u,v) = B(u,v) / N(u,v)$



$N.S$

=

*



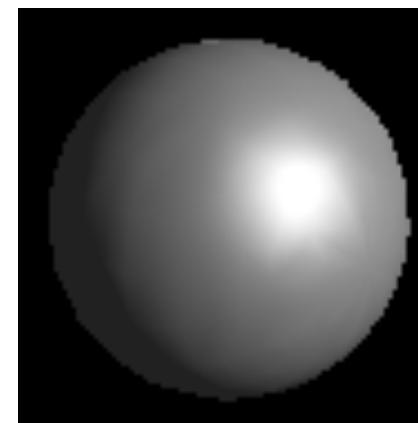
Albedo

SECONDARY ILLUMINATION

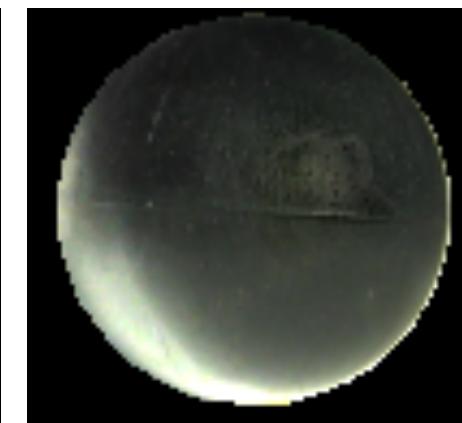
Reflections produce indirect lighting.



Unique light source assumption does not allow correct albedo recovery.

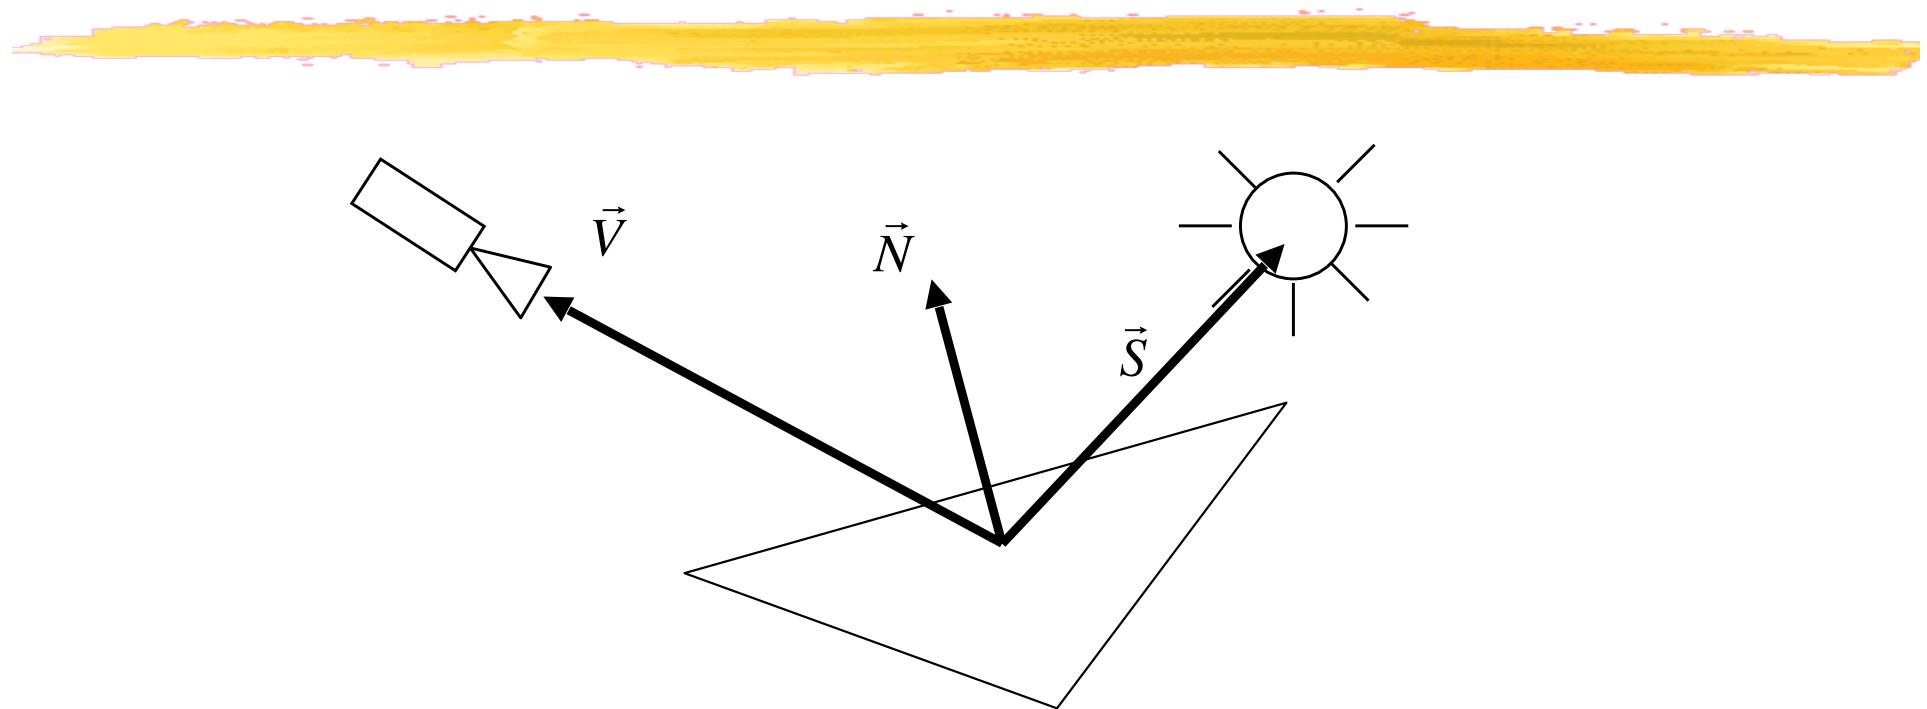


Recovered shading



Recovered albedoes

SIMPLIFYING ASSUMPTIONS



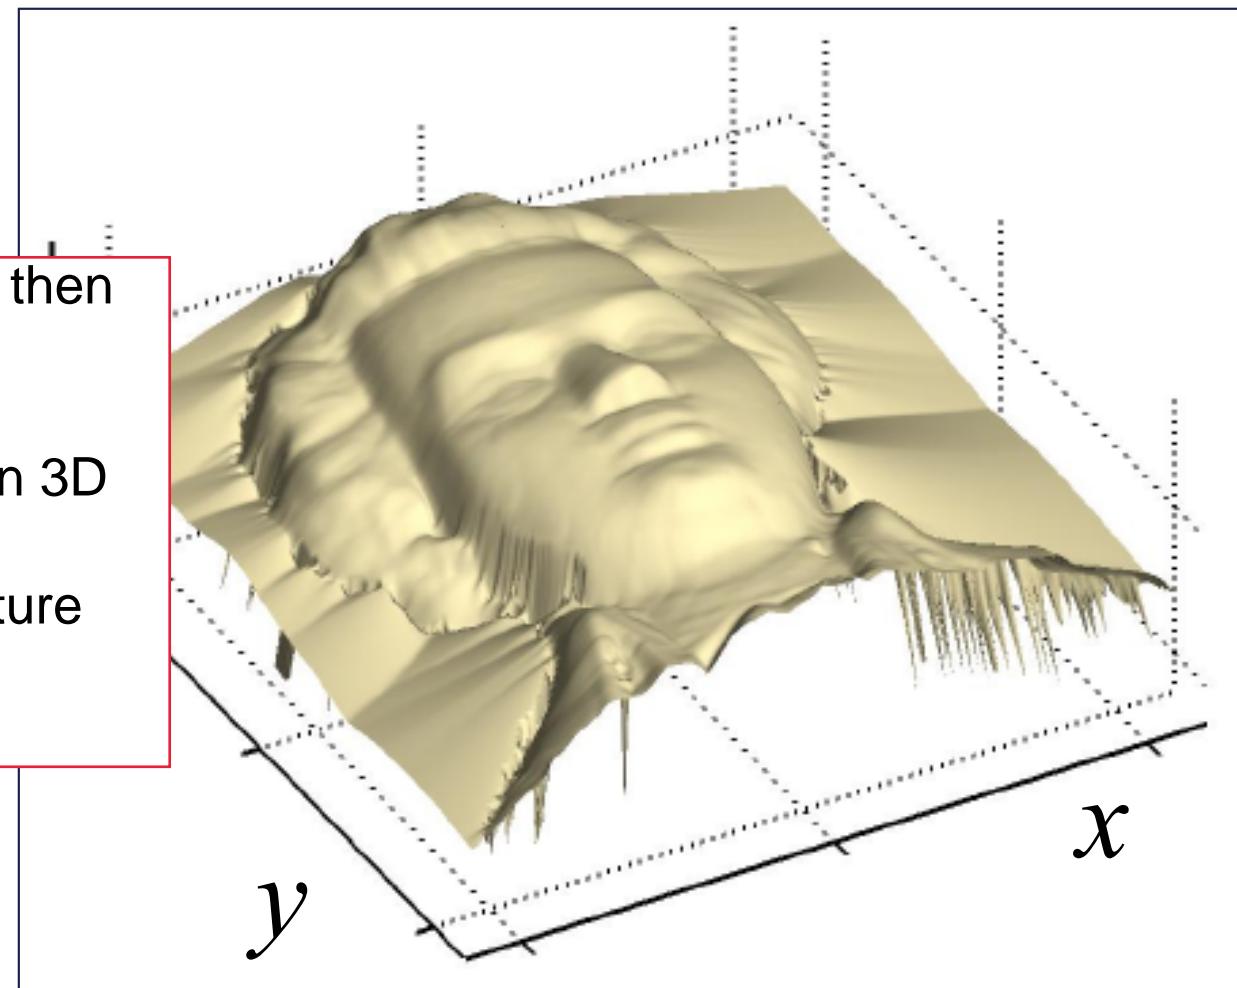
- The illumination sources are distant from the imaged surfaces
- Secondary illumination is not significant
- There are no cast shadows

MONGE SURFACE

$$z = f(x, y)$$

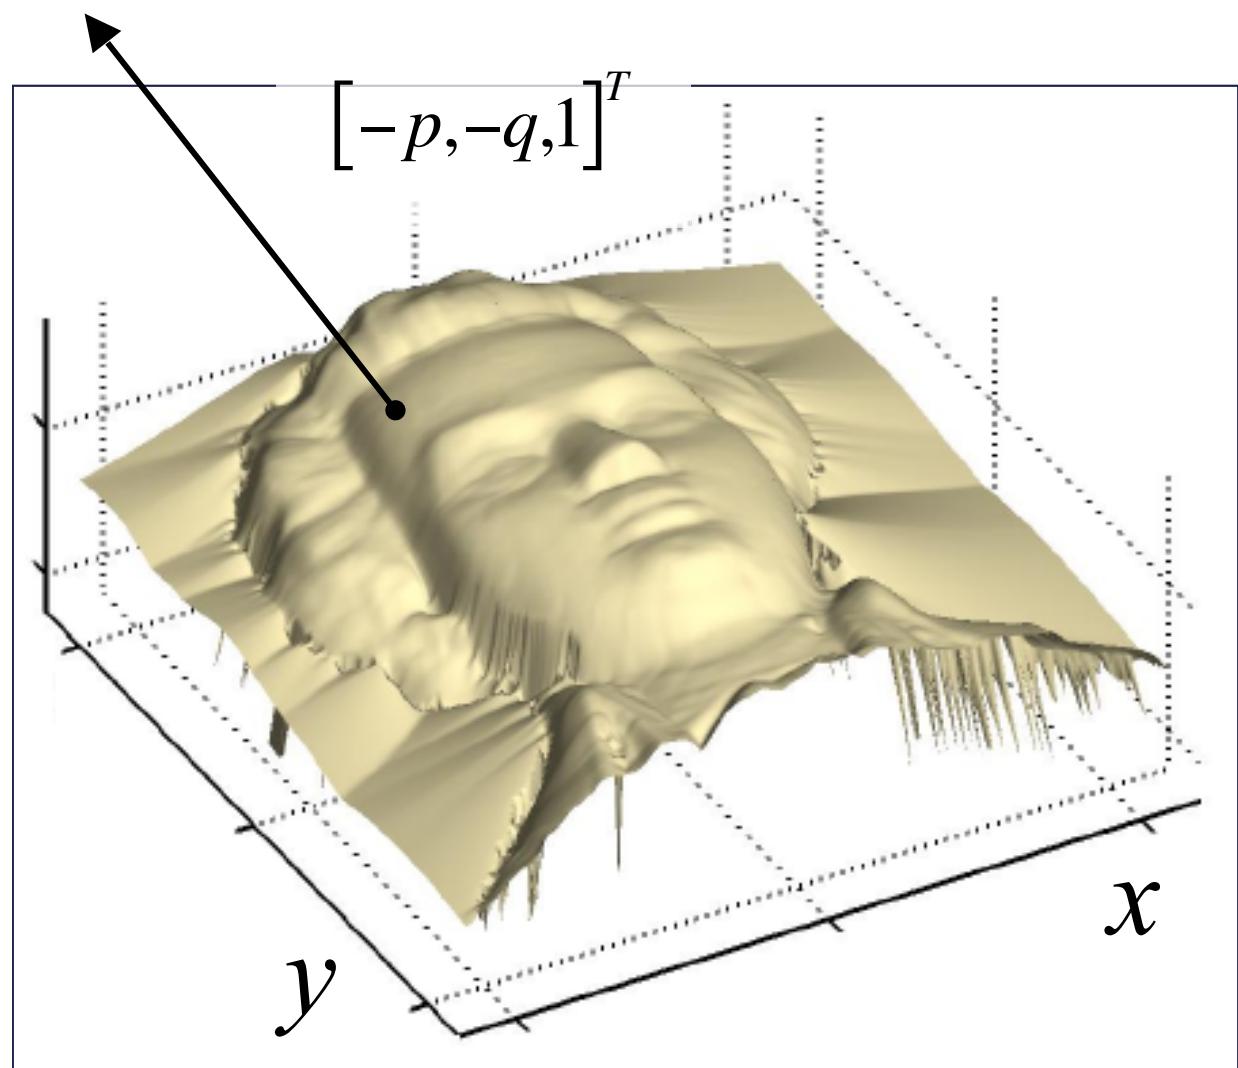
if we pretend z is the intensity, then
 dz/dx = edge in x
 dz/dy = edge in y
which can give us the normal in 3D

thus intensity = normal / curvature



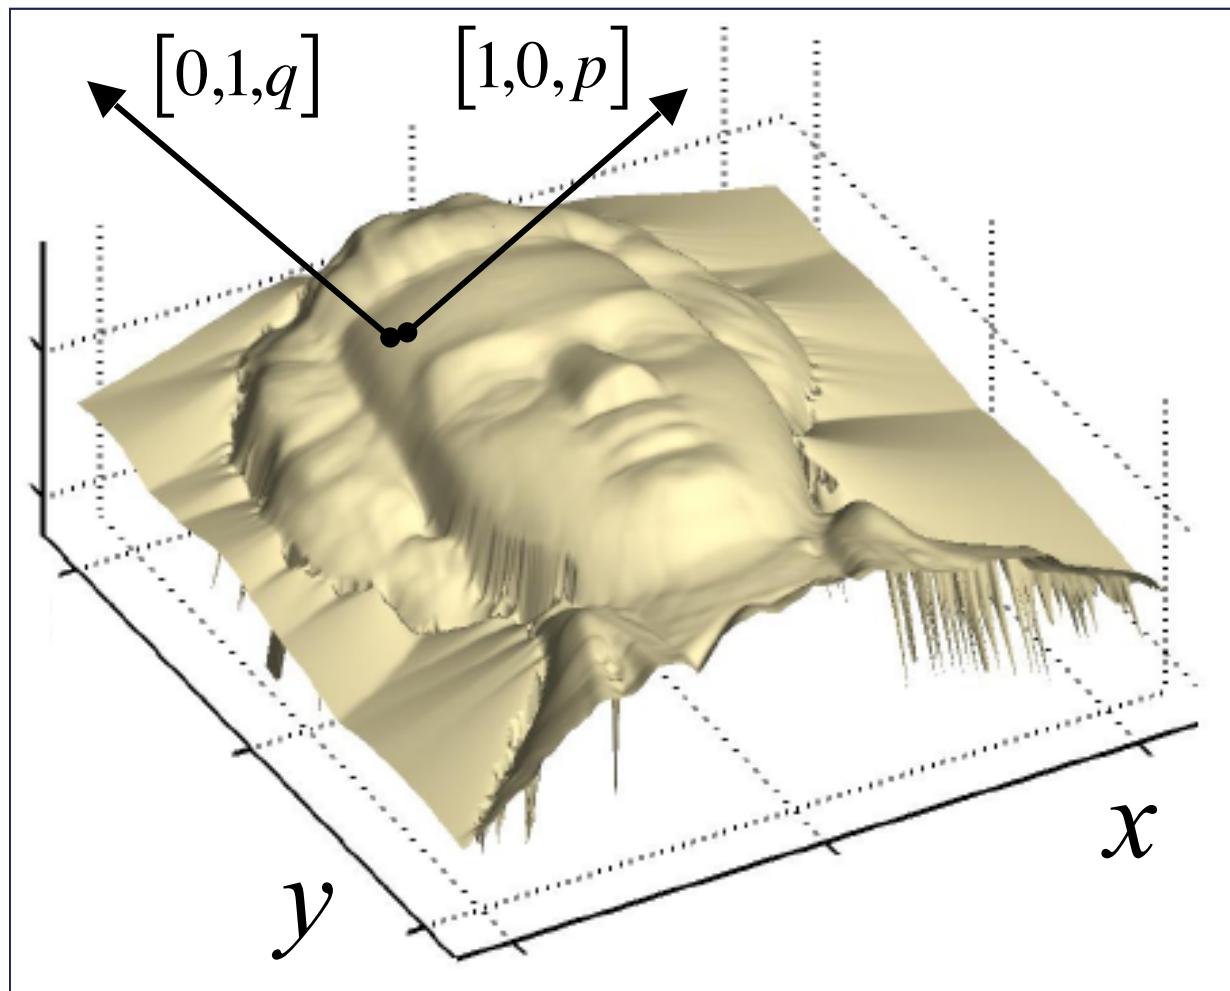
SURFACE NORMALS

$$z = f(x, y)$$
$$p = \frac{\delta z}{\delta x}$$
$$q = \frac{\delta z}{\delta y}$$



TANGENT VECTORS

$$\begin{aligned}z &= f(x, y) \\p &= \frac{\delta z}{\delta x} \\q &= \frac{\delta z}{\delta y}\end{aligned}$$



PROJECTION

Elevation and Normal :

$$z = f(x, y)$$

Surface
normal

$$\mathbf{N}(x, y) = \frac{1}{\sqrt{1 + f_x^2 + f_y^2}} \begin{bmatrix} -f_x \\ -f_y \\ 1 \end{bmatrix}$$

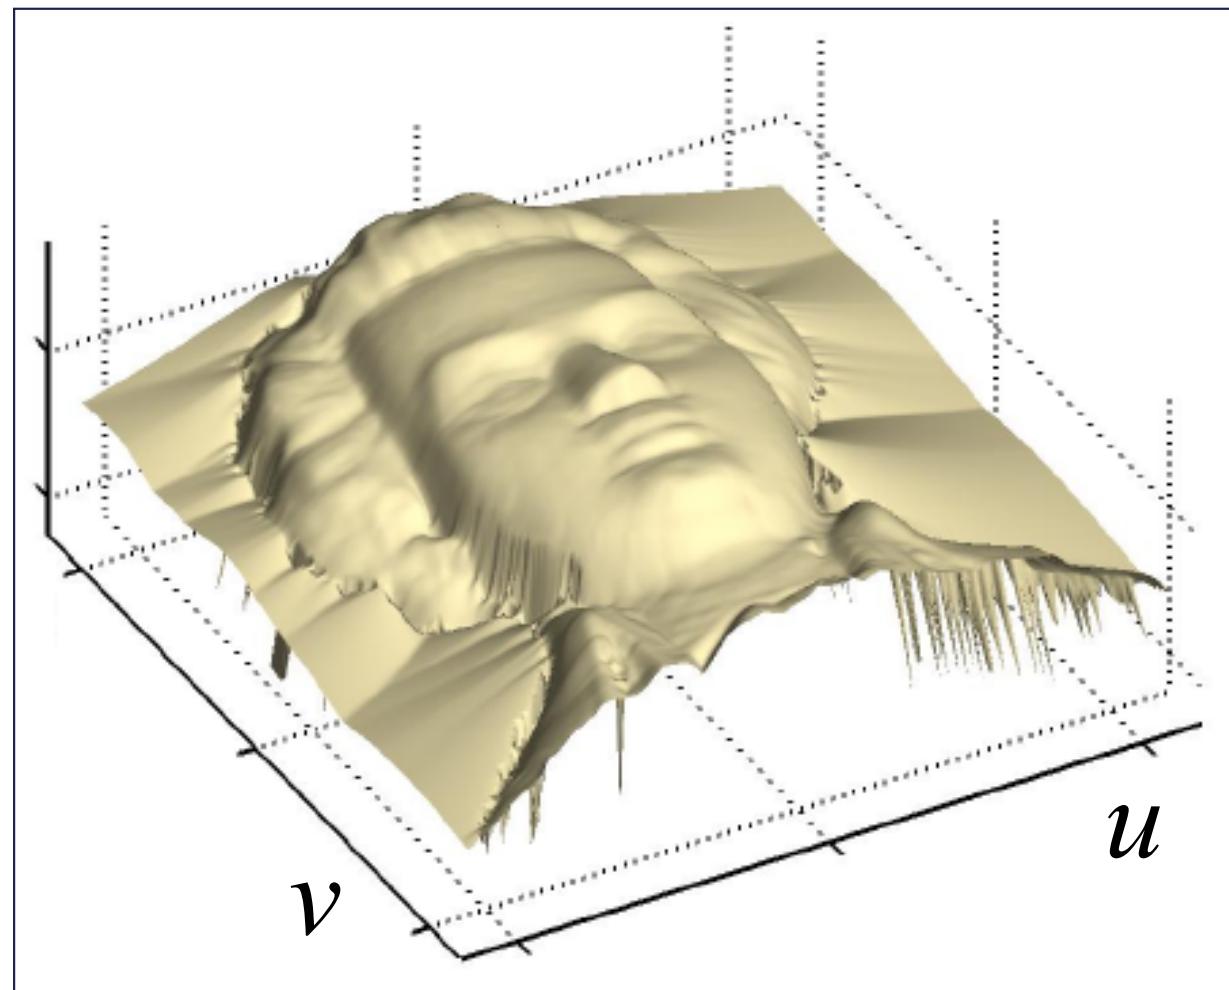
Orthographic Projection :

$$u = sx$$

$$v = sy$$

RE-PARAMETRIZATION

$$z = f(u, v)$$



SHAPE FROM NORMALS

gradient = normal divide by normal

$$\begin{aligned} N &= \frac{1}{\sqrt{1 + \frac{\delta z}{\delta x}^2 + \frac{\delta z}{\delta y}^2}} \begin{bmatrix} -\frac{\delta z}{\delta x} \\ -\frac{\delta z}{\delta y} \\ 1 \end{bmatrix} \propto \begin{bmatrix} -\frac{1}{s} \frac{\delta z}{\delta u} \\ -\frac{1}{s} \frac{\delta z}{\delta v} \\ 1 \end{bmatrix} \\ \Rightarrow \frac{\delta z}{\delta u} &= -s \frac{n_x}{n_z} \text{ and } \frac{\delta z}{\delta v} = -s \frac{n_y}{n_z} \\ \Rightarrow \frac{\delta \bar{z}}{\delta u} &= -\frac{n_x}{n_z} = n_1 \text{ and } \frac{\delta \bar{z}}{\delta v} = -\frac{n_y}{n_z} = n_2, \text{ with } \bar{z} = \frac{z}{s} \end{aligned}$$

FINITE DIFFERENCES

we already have
normals.

goal is to
estimate z ?
depth ?

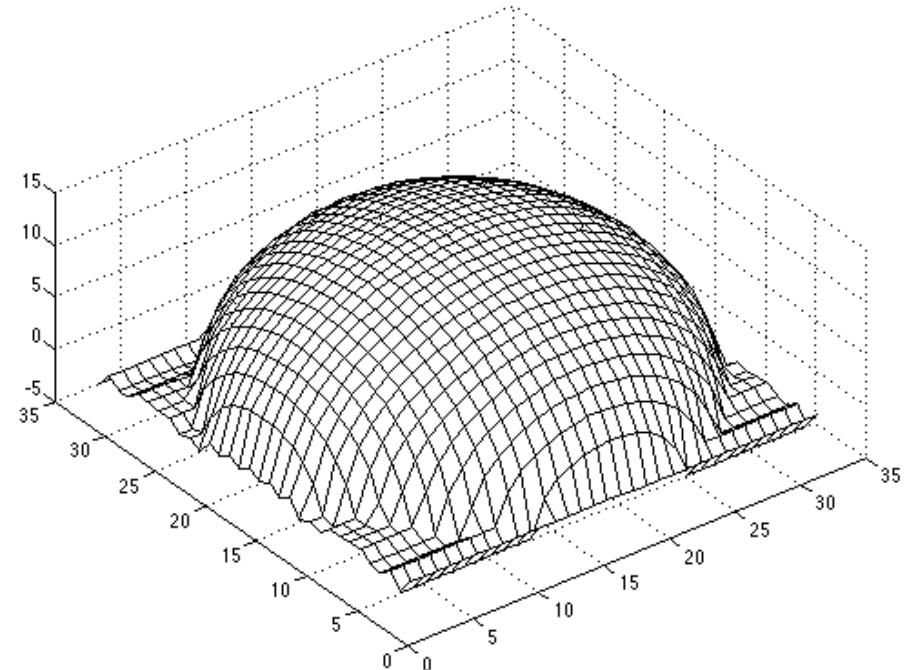
$\forall u, v$

$$\bar{z}(u+1, v) - \bar{z}(u, v) = n_1(u, v)$$

$$\bar{z}(u, v+1) - \bar{z}(u, v) = n_2(u, v)$$

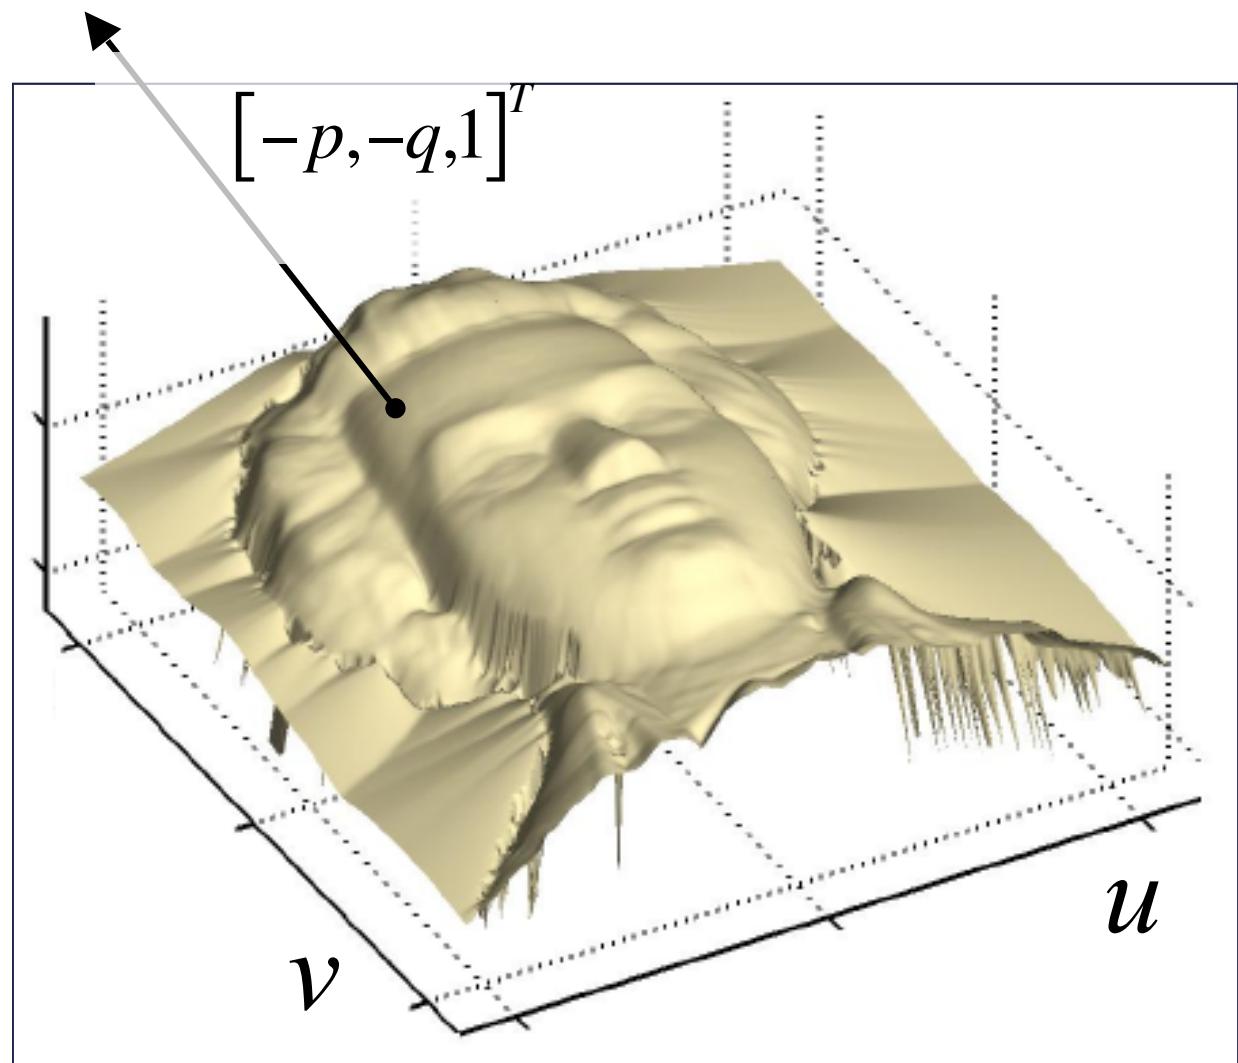
\Rightarrow Twice as many equations as there are unknown.

\Rightarrow Least square solution up to a scale factor.



GRADIENT SPACE

$$\begin{aligned}z &= f(u, v) \\p &= \frac{\delta z}{\delta u} \\q &= \frac{\delta z}{\delta v}\end{aligned}$$



REFLECTANCE MAP



Reflectance:

Amount of light reflected towards the camera.

Reflectance i.e. estimated intensity

Albedo:

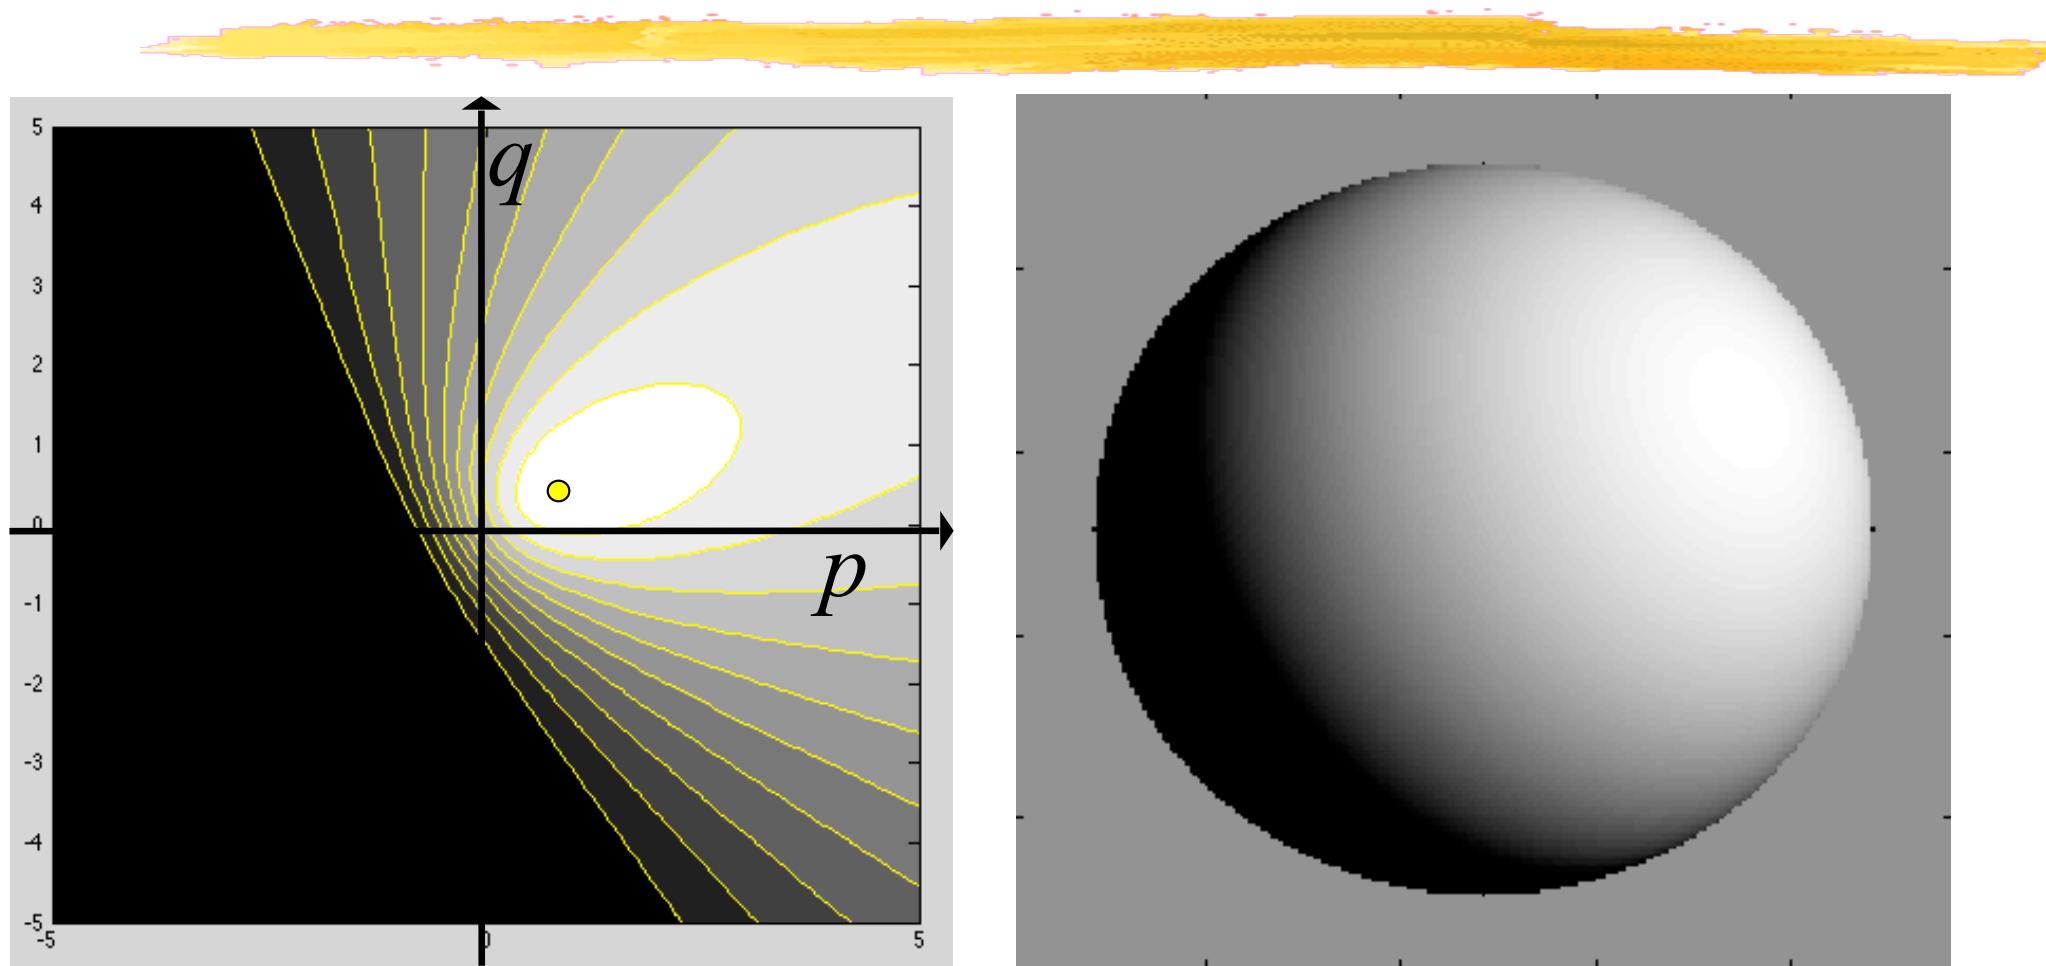
Fraction of the light incident on the surface that is reflected over all directions.

→ In the Lambertian case and for a constant albedo

$$I(u,v) \propto \text{Ref}(p(u,v),q(u,v))$$

$$\propto [p(u,v),q(u,v),-1] \cdot S$$

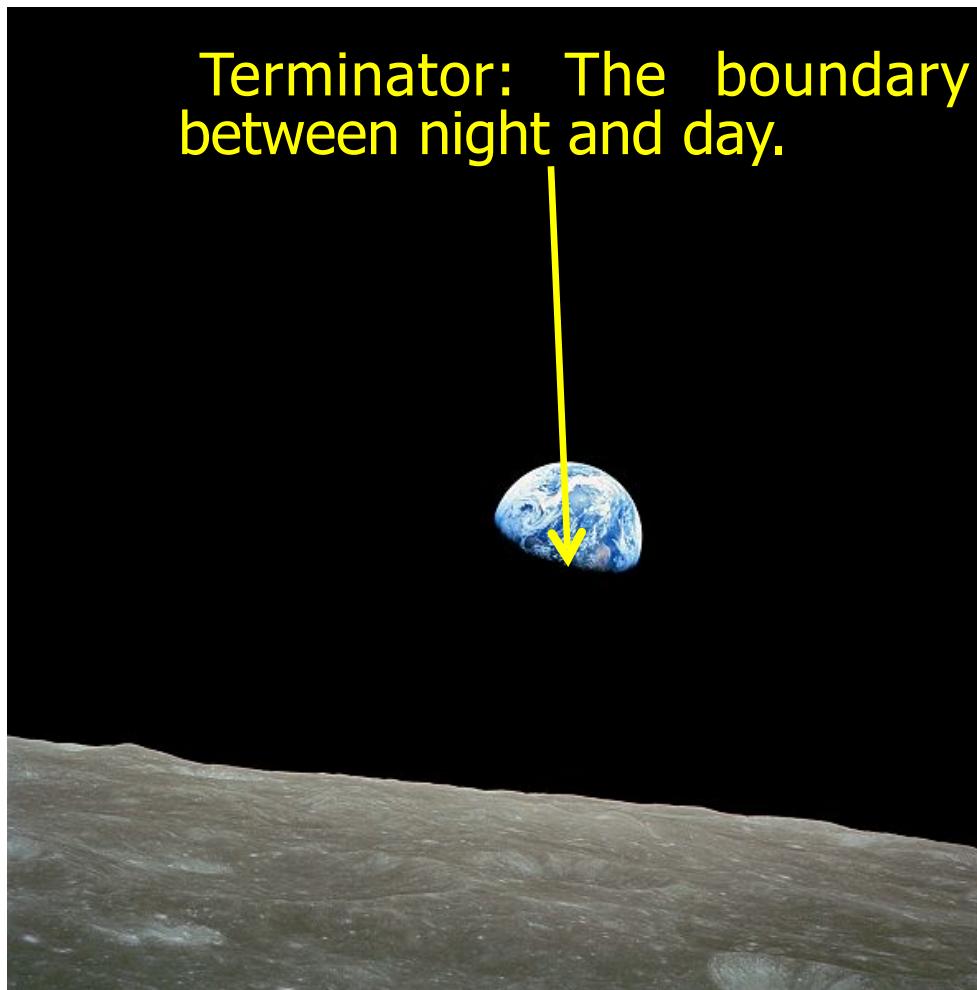
LAMBERTIAN REFLECTANCE MAP



Reflectance map and shaded surface for Lambertian surface illuminated in the direction $[-1 \ -0.5 \ -1]$.

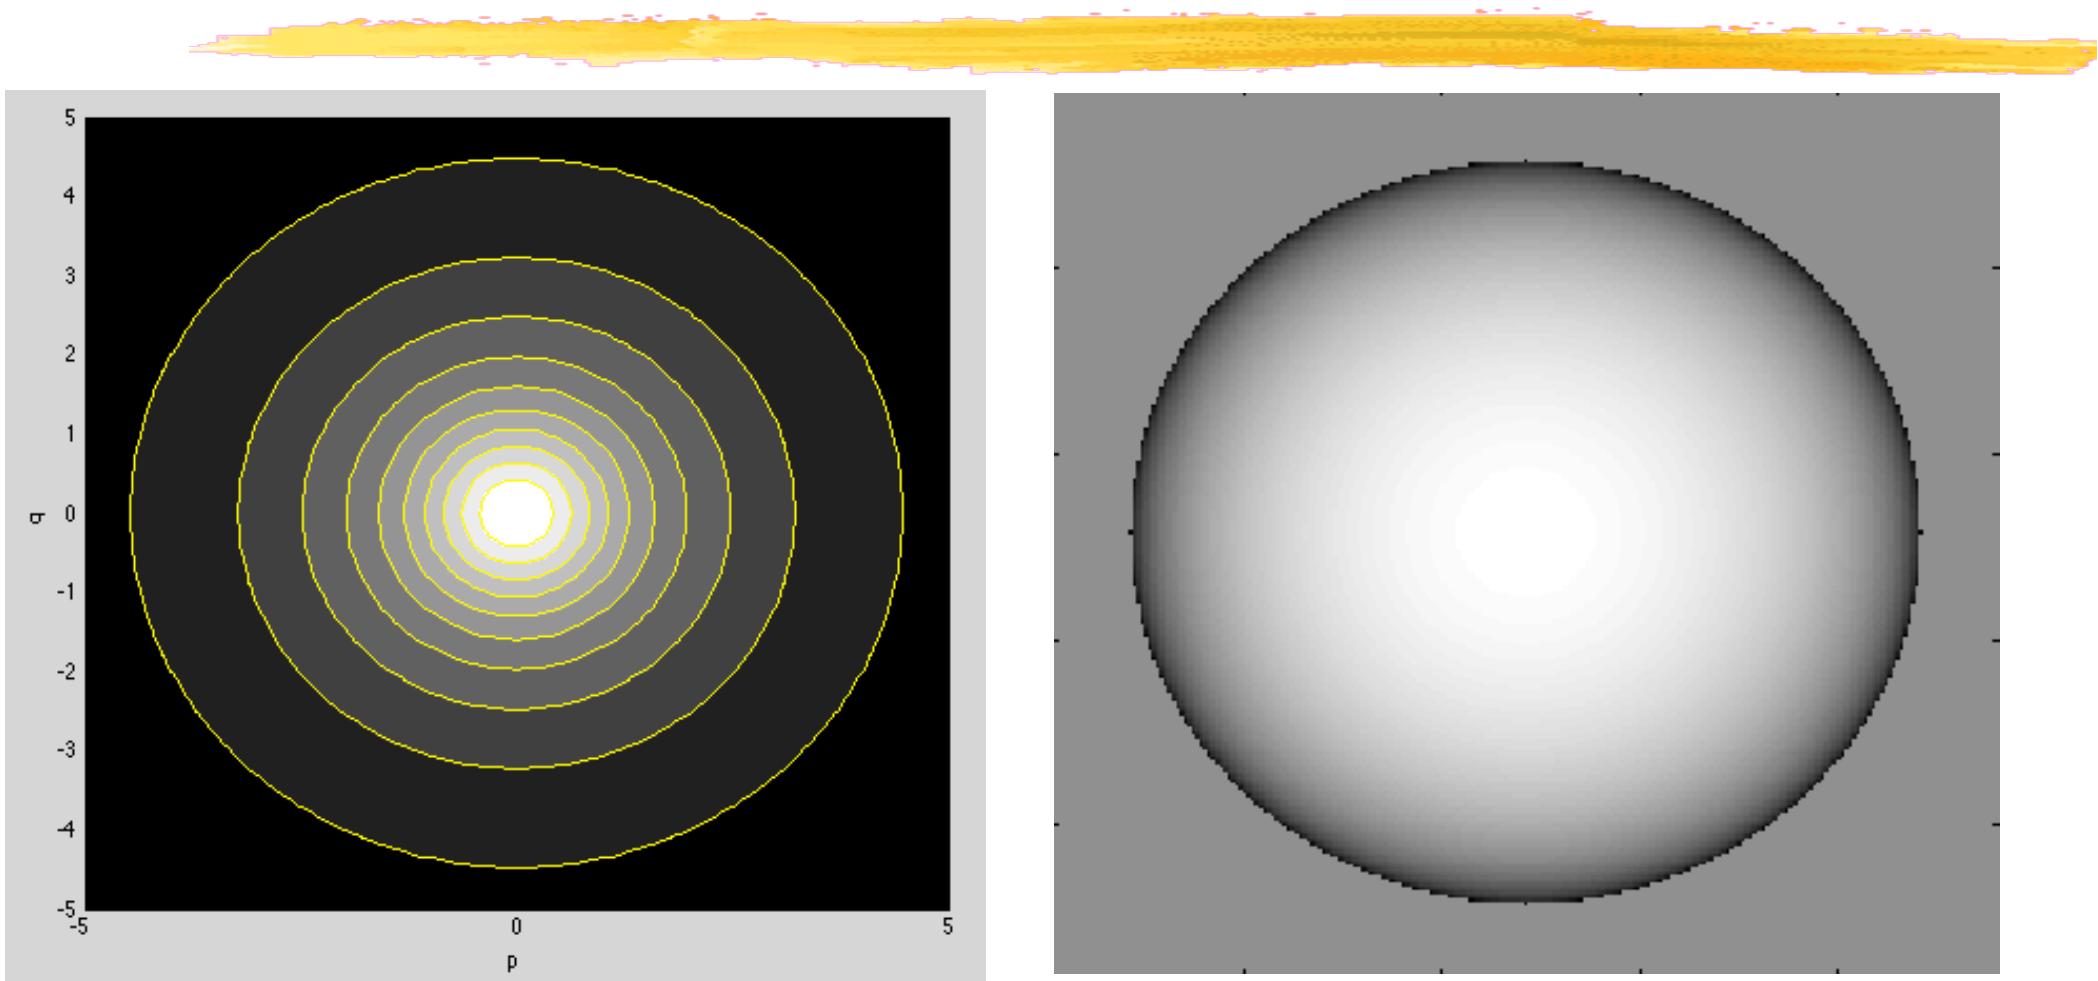
EARTH RISE

Terminator: The boundary between night and day.



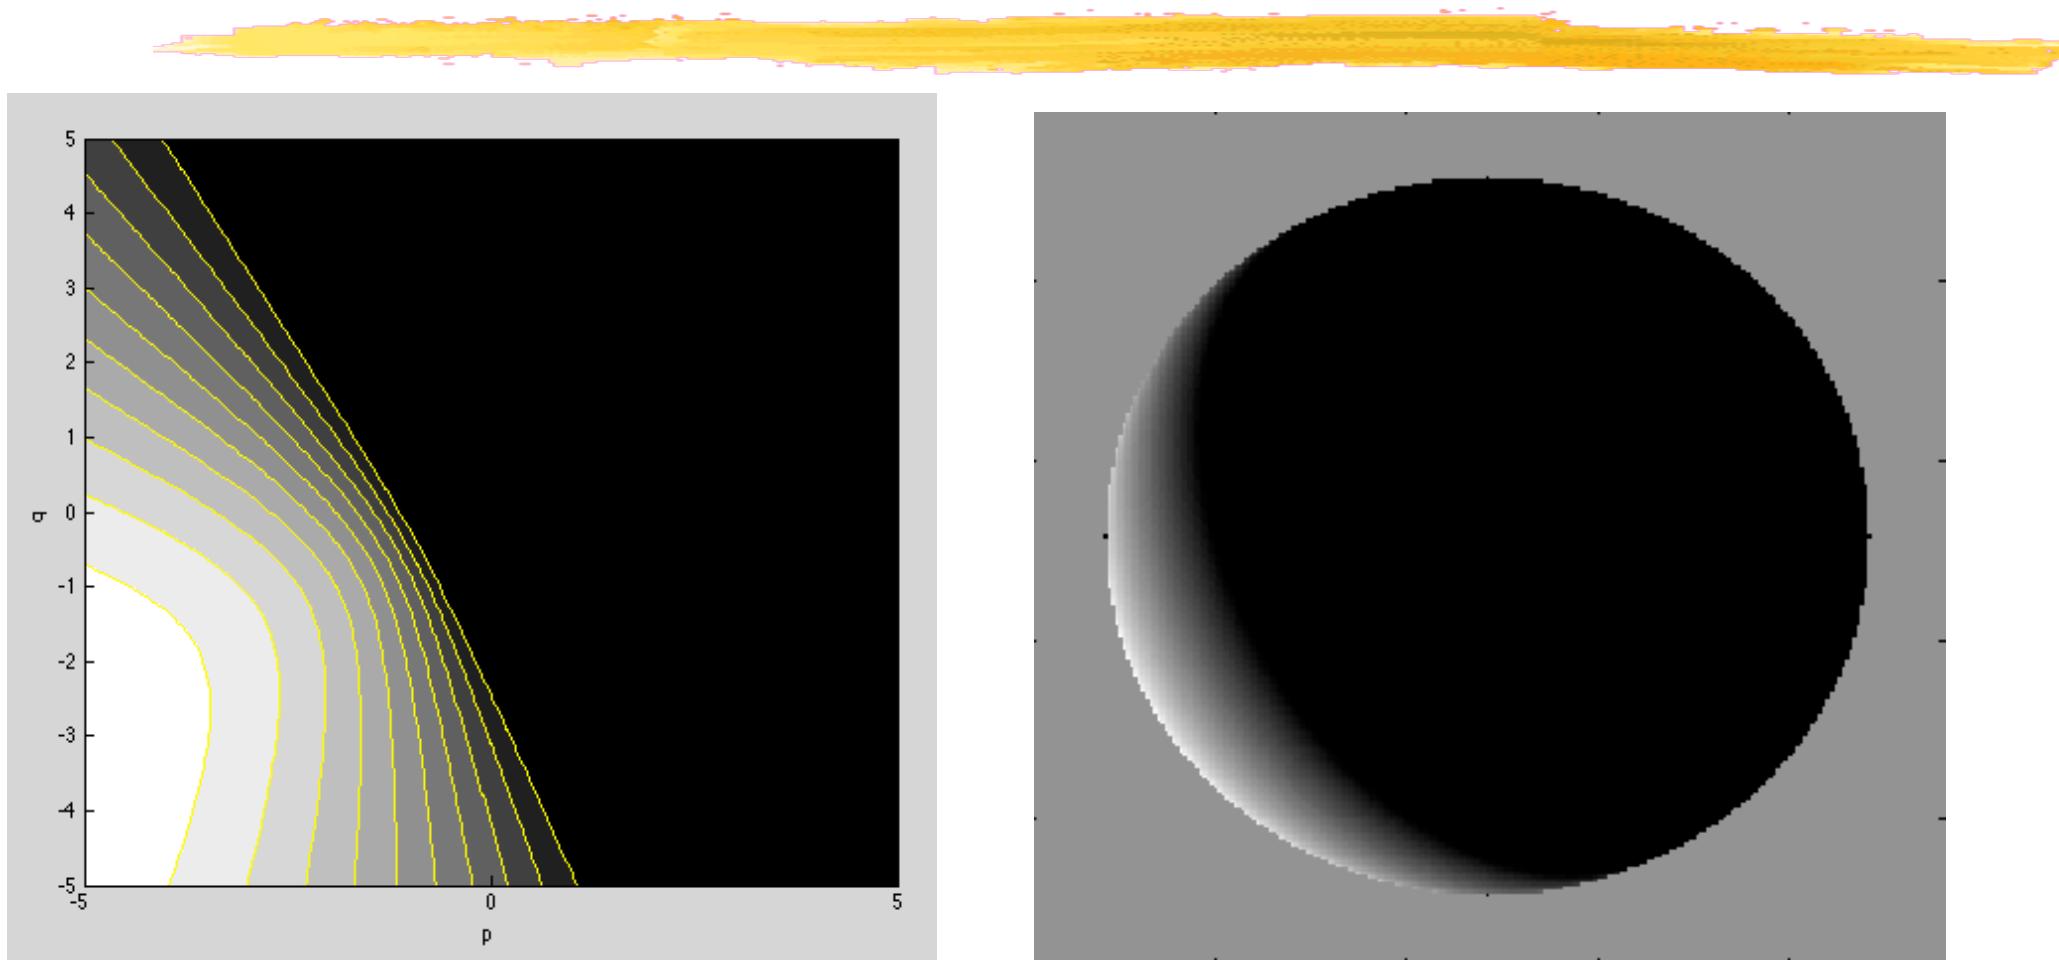
The earth seen from the moon. Apollo 8, 1968.

LAMBERTIAN REFLECTANCE MAP



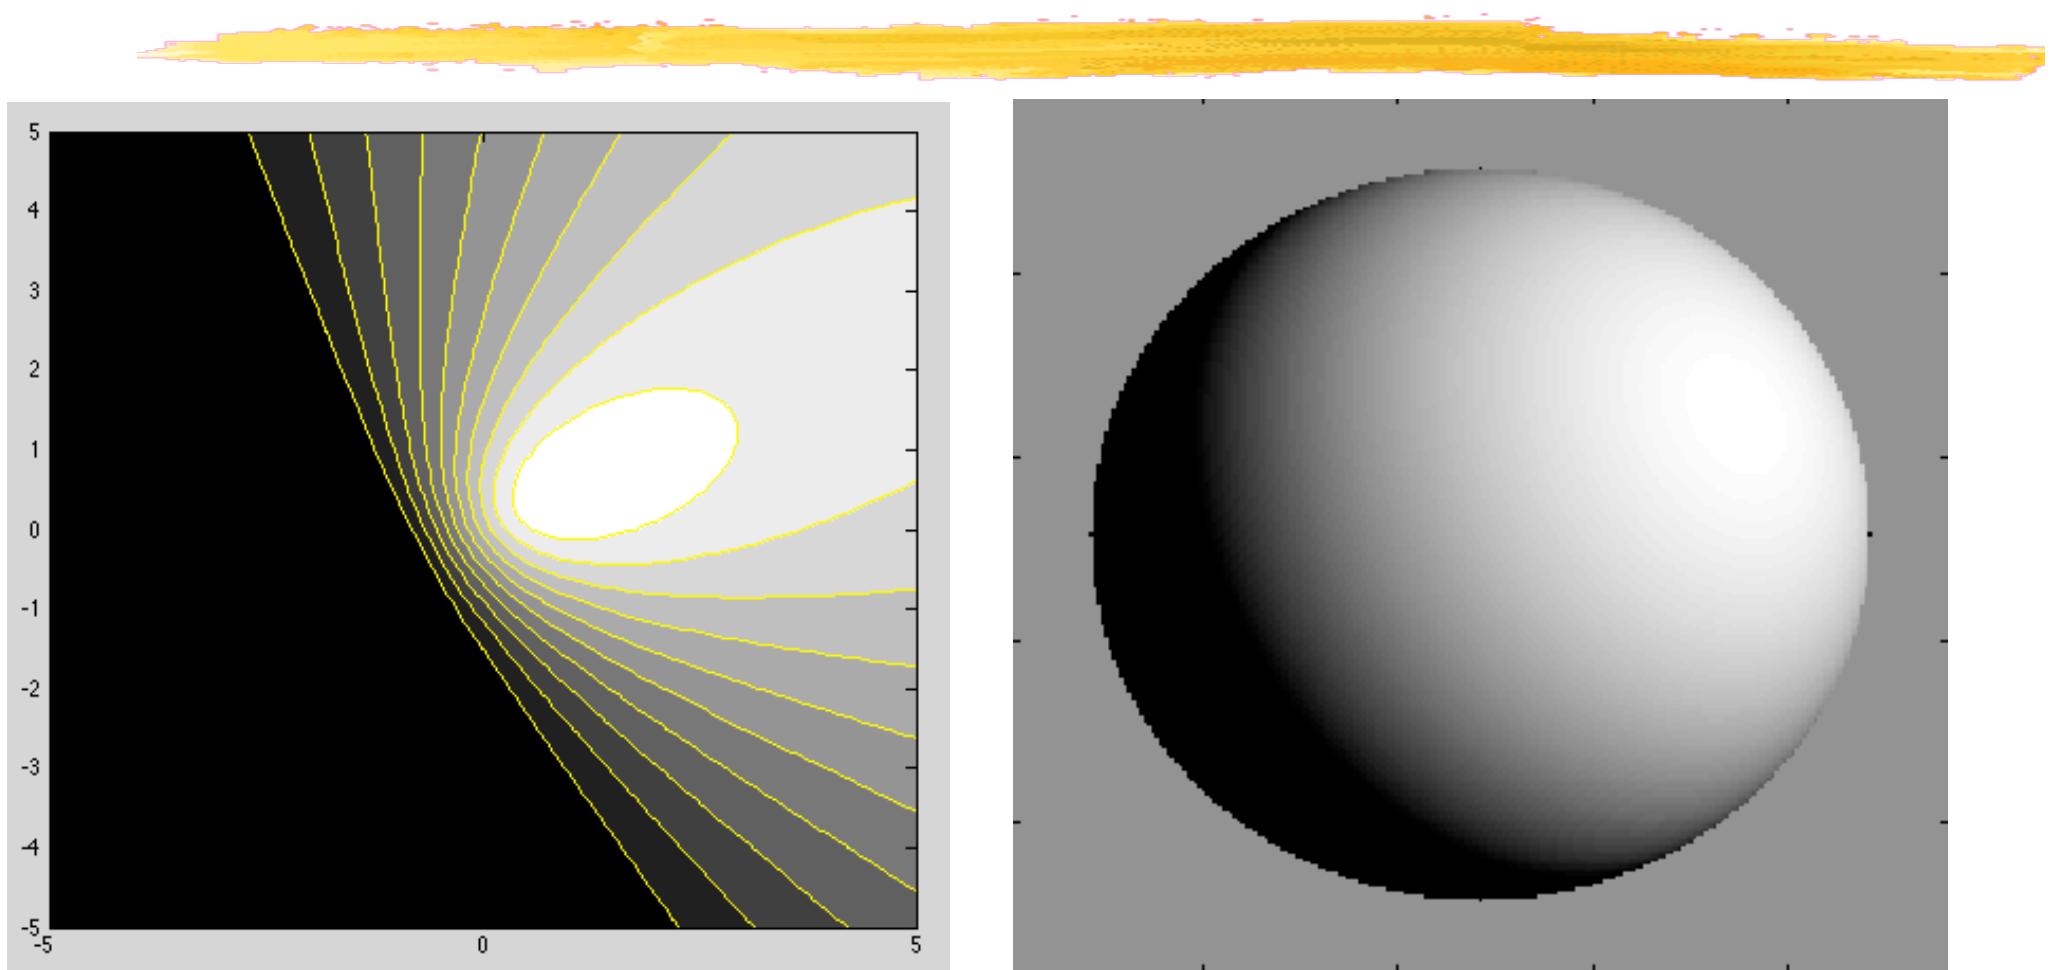
Reflectance map and shaded surface for Lambertian surface illuminated in the direction [0 0 -1].

LAMBERTIAN REFLECTANCE MAP



Reflectance map and shaded surface for Lambertian surface illuminated in the direction [1 0.5 -1].

CAN WE DETERMINE (P, Q) UNIQUELY FOR EACH IMAGE POINT INDEPENDENTLY?



NO -> Global optimization required.

VARIATIONAL METHODS

Minimize:

$$\iint \left(\left[I(u, v) - Ref\left(\frac{\delta z}{\delta u}, \frac{\delta z}{\delta v}\right) \right]^2 + \lambda \left[\left(\frac{\delta^2 z}{\delta u^2} \right)^2 + \left(\frac{\delta^2 z}{\delta u \delta v} \right)^2 + \left(\frac{\delta^2 z}{\delta v^2} \right)^2 \right] \right) dudv$$

observed intensity

intensity given current estimate of the shape of the surface

or:
Brightness constraint

Smoothness term

$$\iint \left([I(u, v) - Ref(p, q)]^2 + \lambda \left[\left(\frac{\delta p}{\delta u} \right)^2 + \left(\frac{\delta p}{\delta v} \right)^2 + \left(\frac{\delta q}{\delta u} \right)^2 + \left(\frac{\delta q}{\delta v} \right)^2 \right] + \mu \left[\frac{\delta p}{\delta v} - \frac{\delta q}{\delta u} \right]^2 \right) dudv$$

u_1
$v_1 z_1 z_2 $

$ z_3 z_4 $

the double integral is actually a summation for the rows and columns. at every u and v we have an unknown z

so the surface is smooth. second order penalizes surfaces that are not smooth.
minimize curvature

Integrability constraint

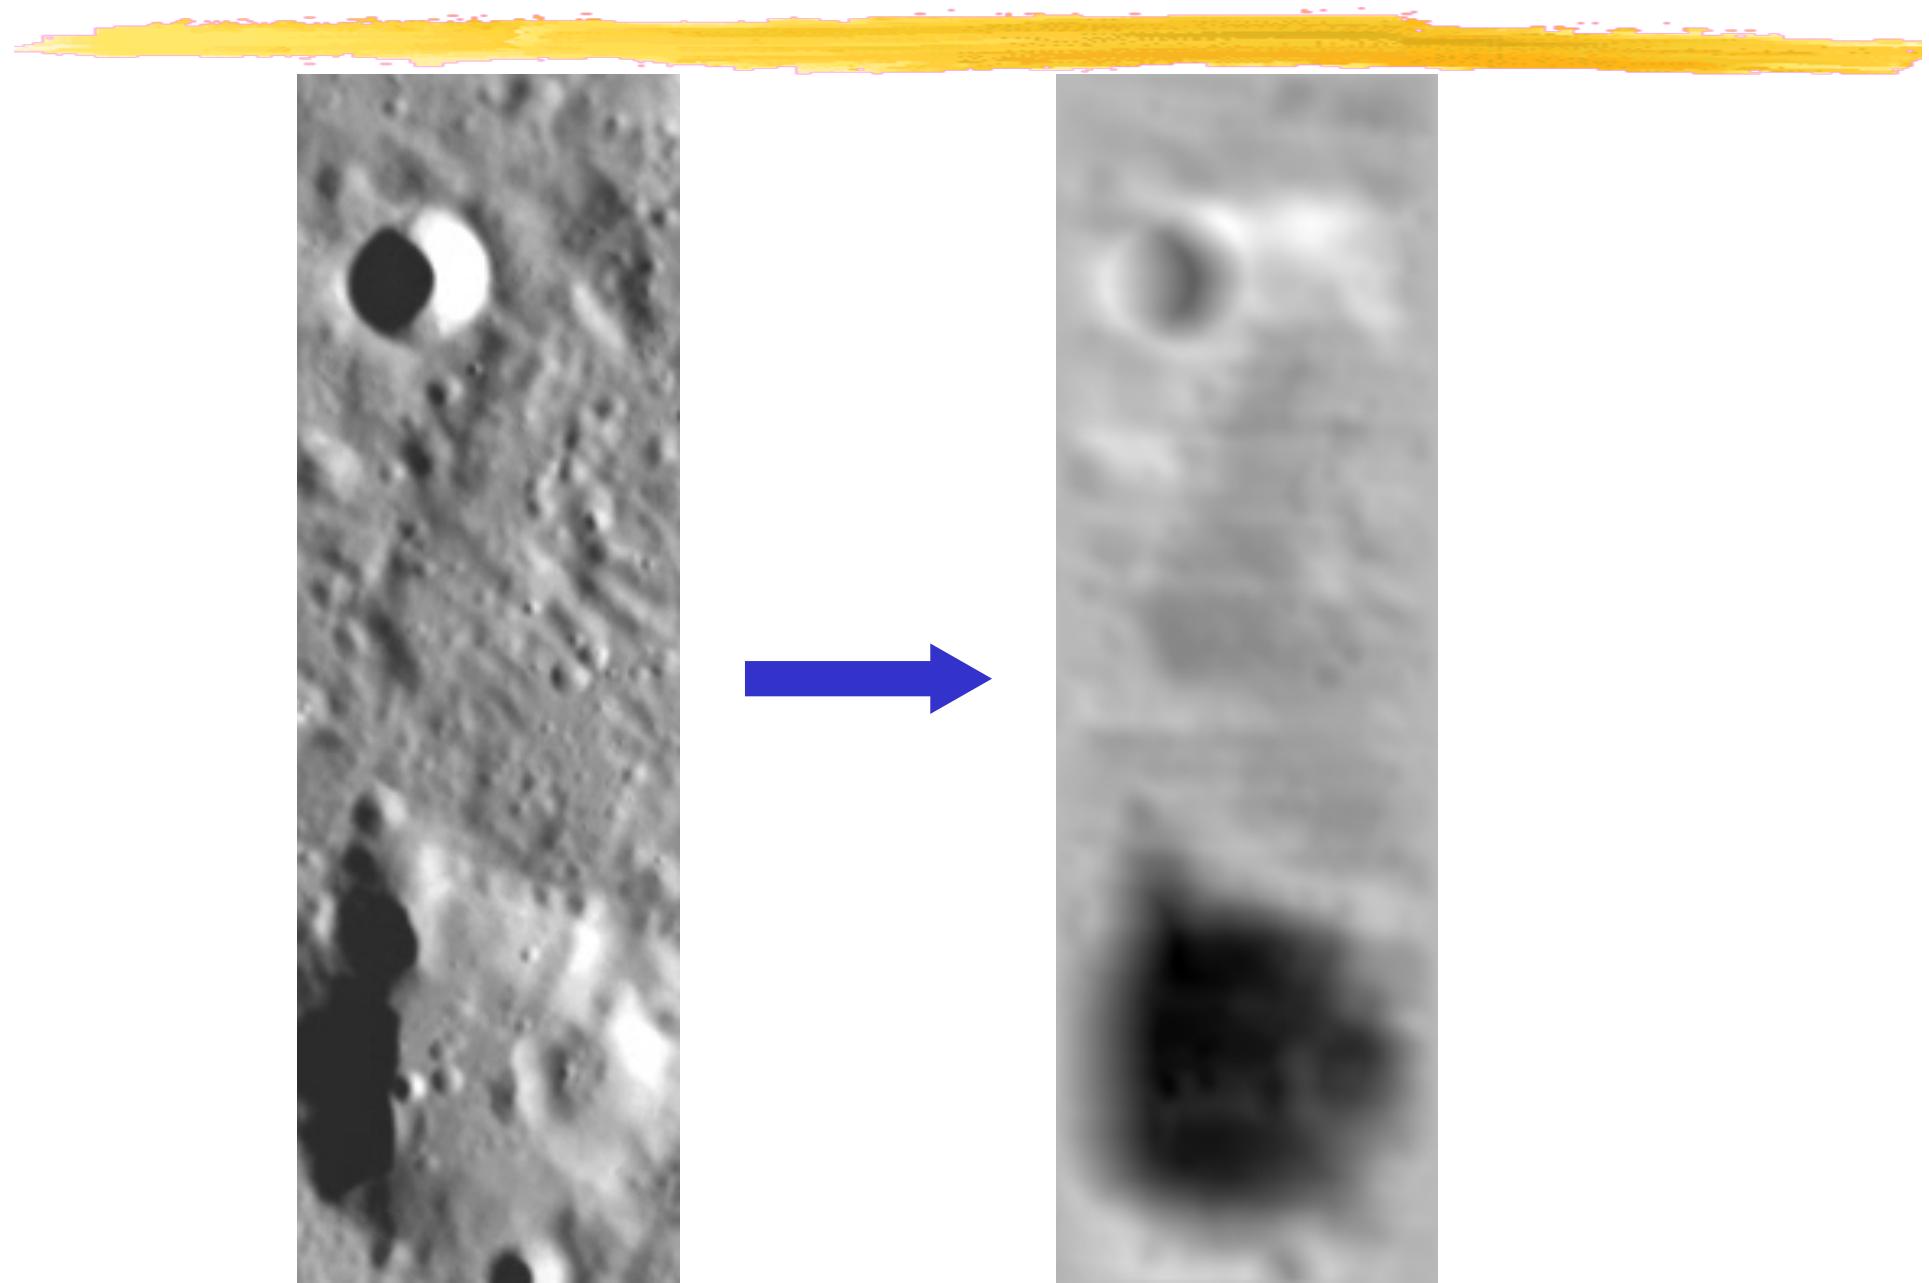
VARIATIONAL METHOD

$$\int \int \left([I(u, v) - R e f(p, q)]^2 + \lambda \left[\left(\frac{\delta p}{\delta u} \right)^2 + \left(\frac{\delta p}{\delta v} \right)^2 + \left(\frac{\delta q}{\delta u} \right)^2 + \left(\frac{\delta q}{\delta v} \right)^2 \right] + \mu \left[\frac{\delta p}{\delta v} - \frac{\delta q}{\delta u} \right]^2 \right) dudv$$

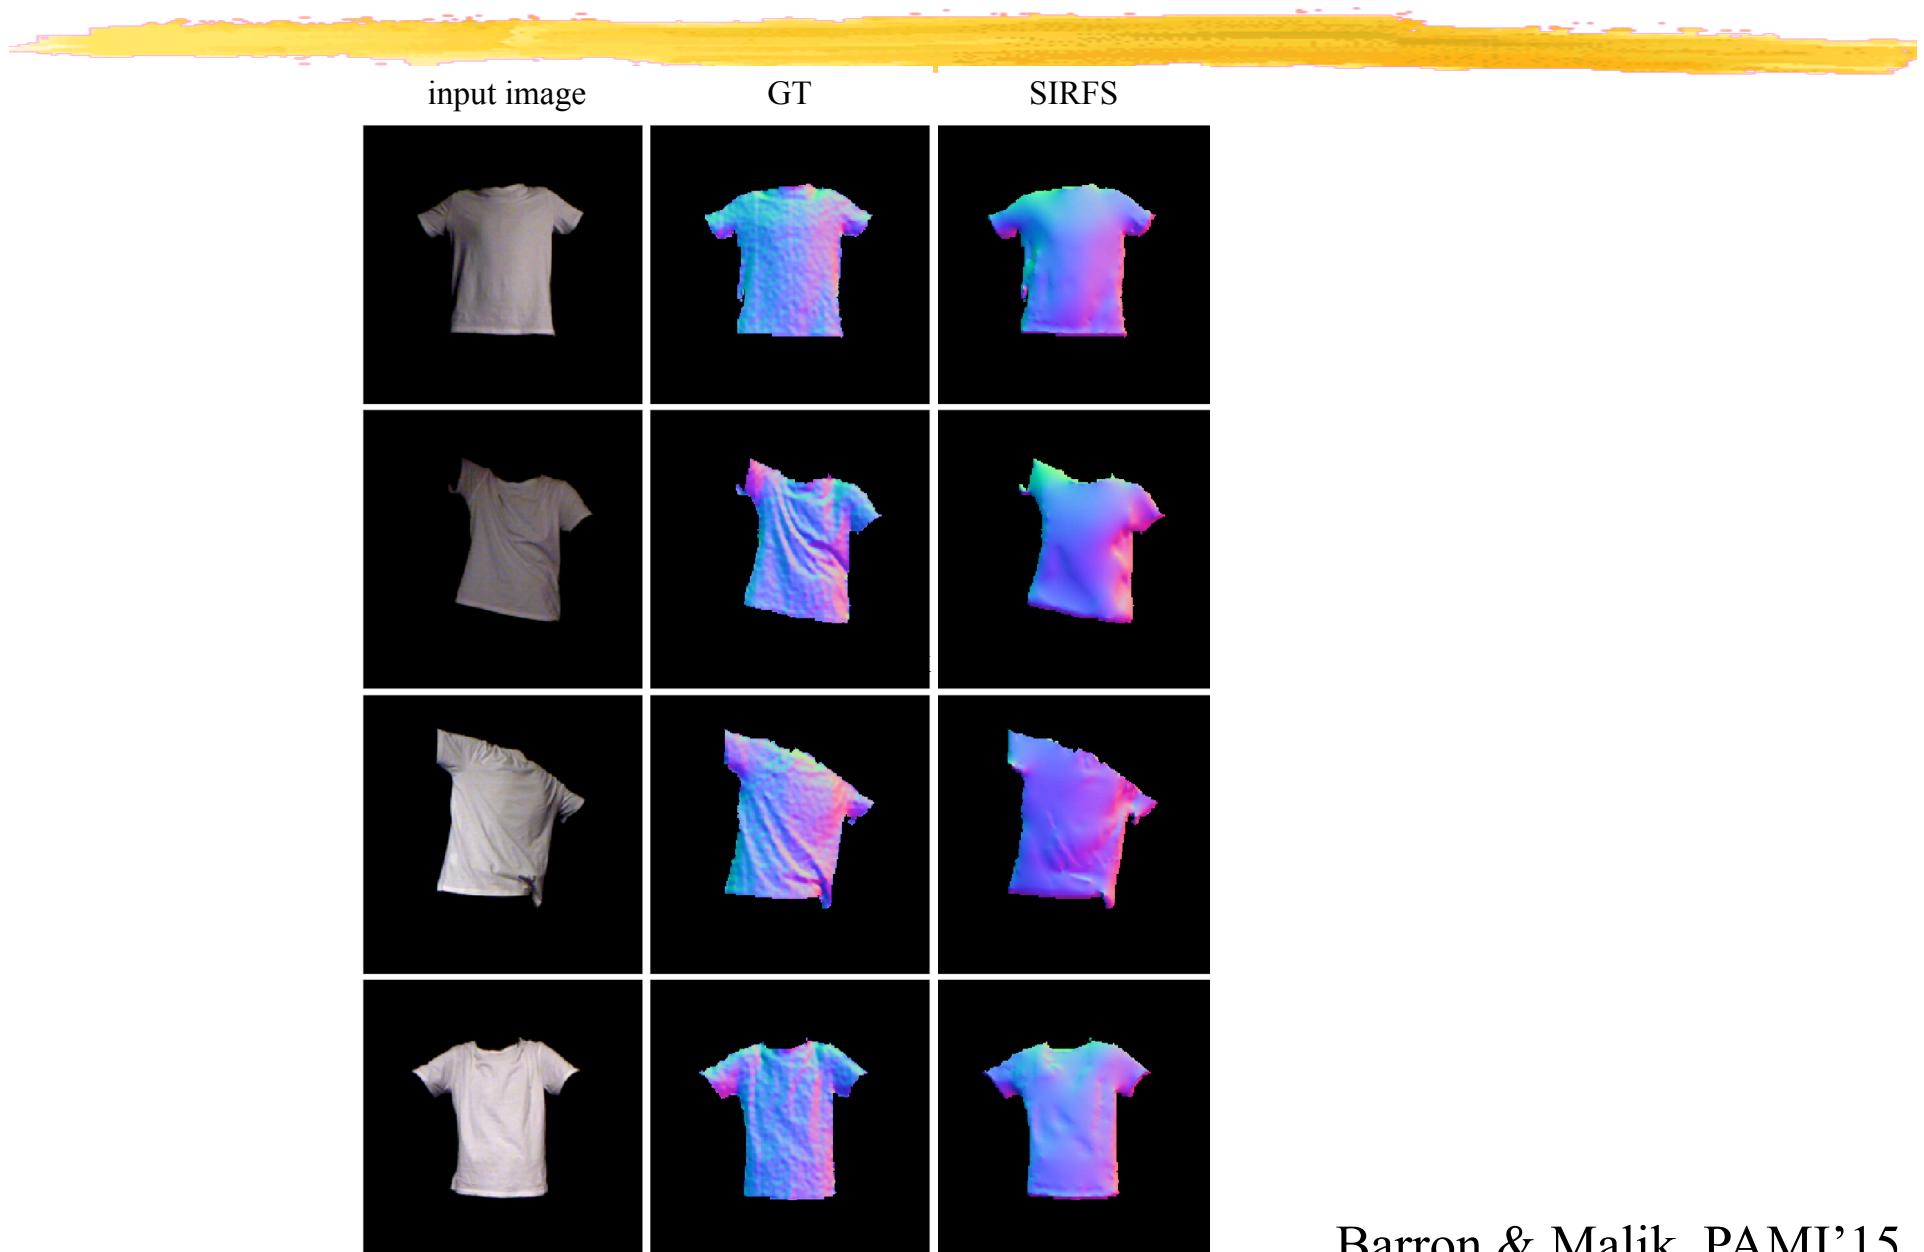
Once p and q have been estimated, integrate to recover f .

→ Need to know the boundary conditions.

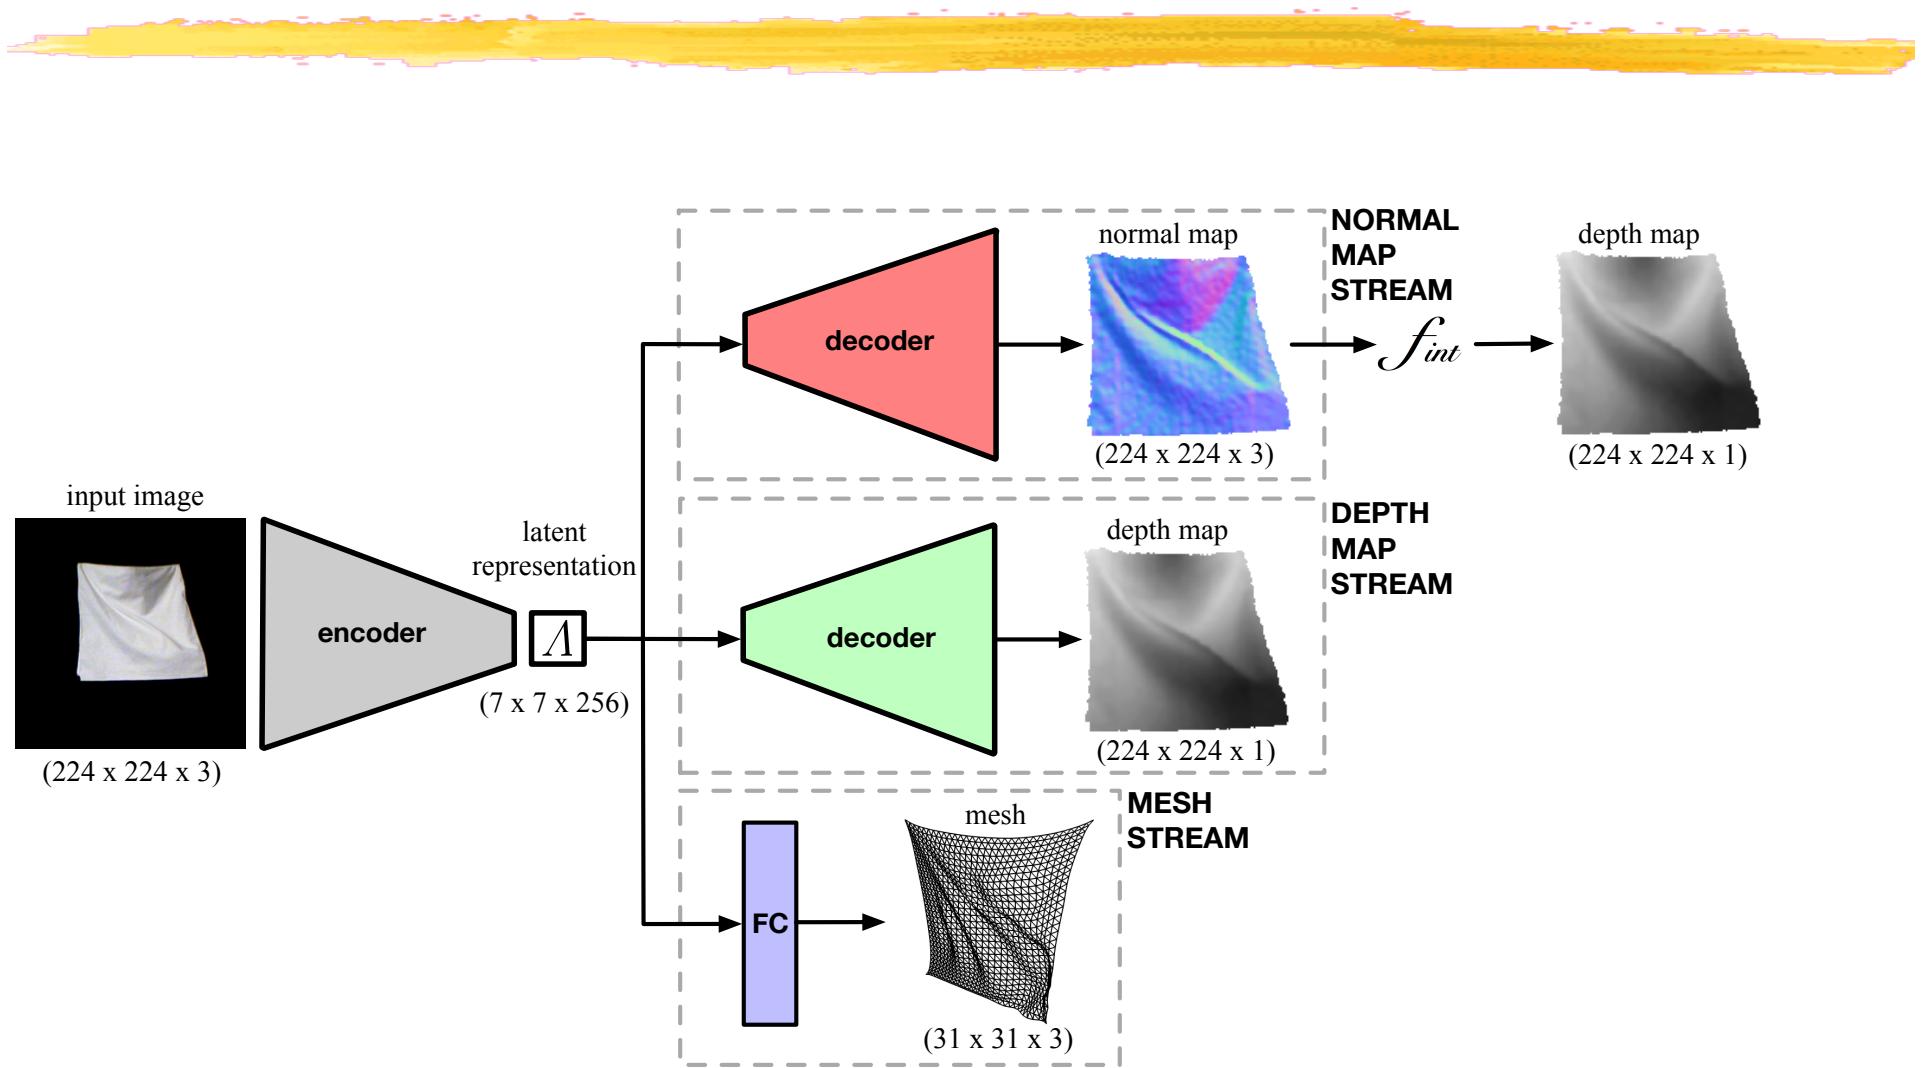
MOONSCAPE



CLOTH



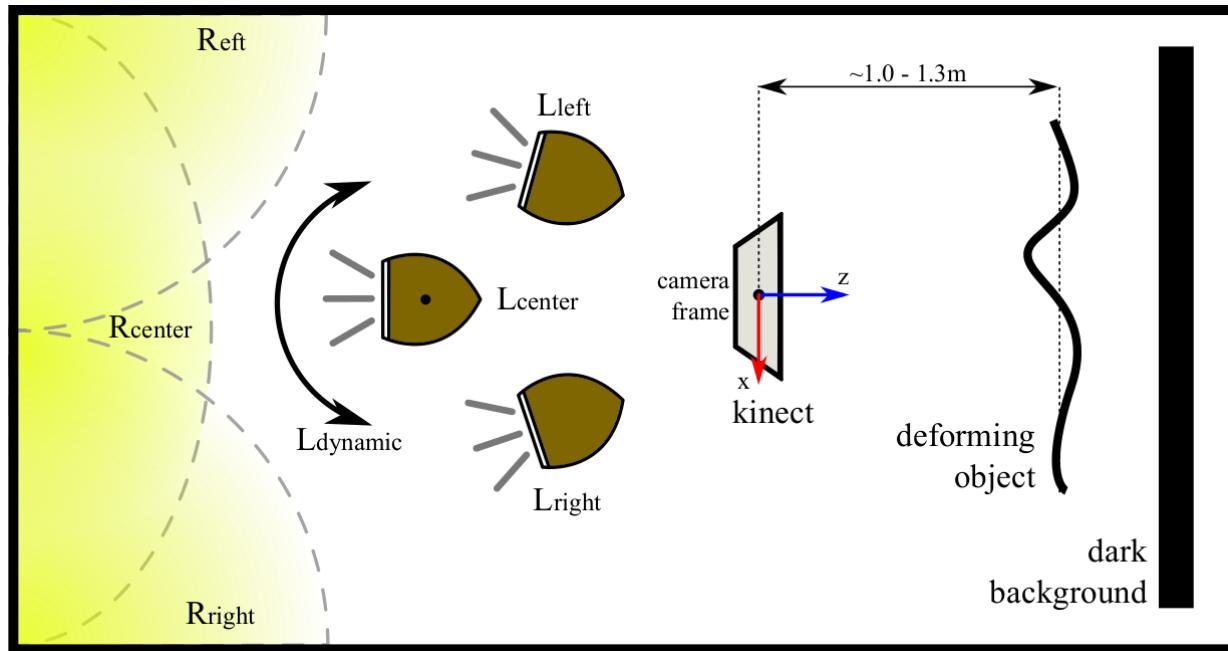
DEEP NETS



DEEP NETS

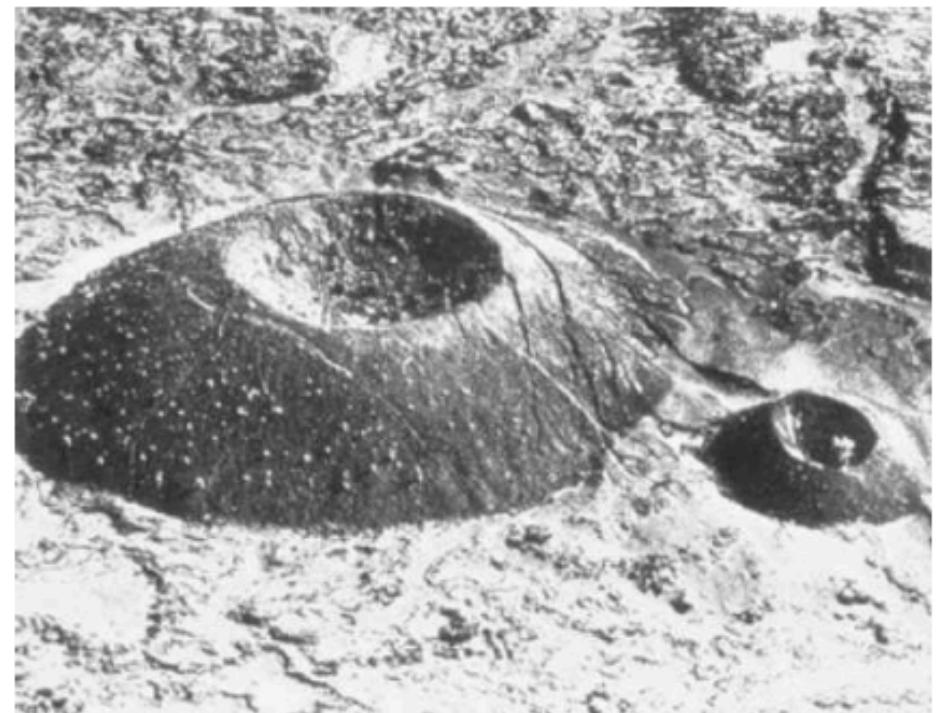


DATA ACQUISITION SETUP



- 3 fixed light sources.
- 1 mobile one.
→ Still a constrained environment.

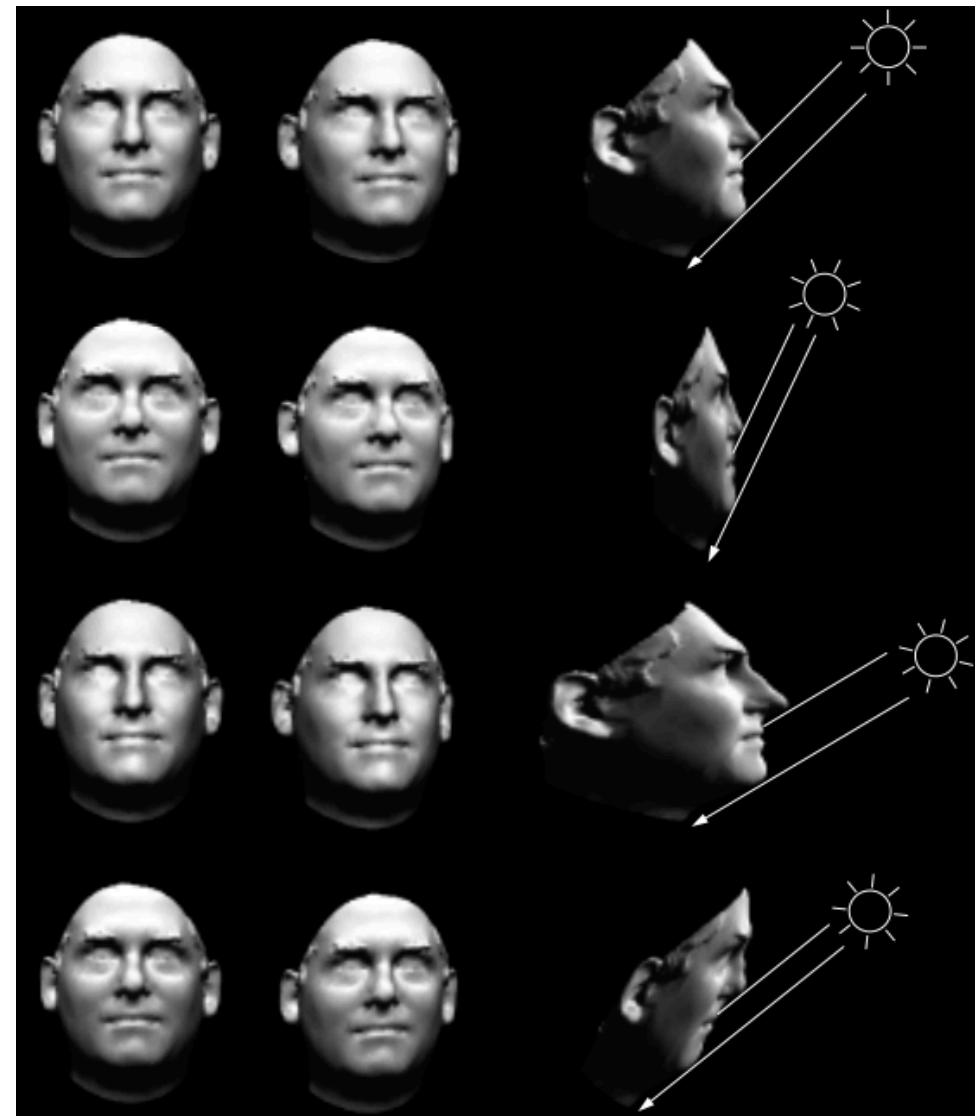
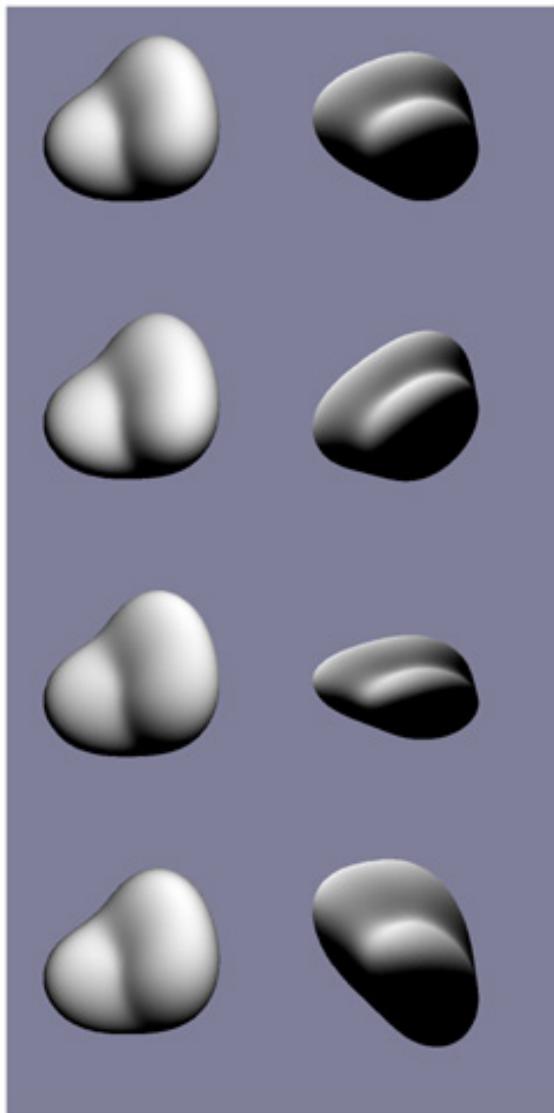
AMBIGUITIES



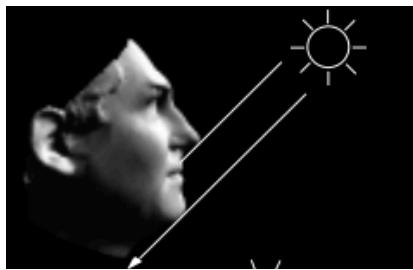
THE BAS-RELIEF AMBIGUITY



MORE GENERALLY

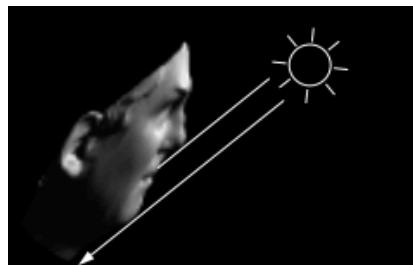


BAS-RELIEF AMBIGUITY



$$Ref = \mathbf{N} \cdot \mathbf{S}$$

For any invertible 3×3 linear transformation A :



$$\mathbf{N} \cdot \mathbf{S} = N^T S = (A\mathbf{N})^T A^{-T} \mathbf{S}$$

But for a valid surface $z=f(u,v)$, we should have:

$$\left. \begin{array}{lcl} \frac{\delta z}{\delta u} & = & -\frac{n_x^*}{n_z^*} \\ \frac{\delta z}{\delta v} & = & -\frac{n_y^*}{n_z^*} \end{array} \right\} \Rightarrow \frac{\delta \frac{n_x^*}{n_z^*}}{\delta v} = \frac{\delta \frac{n_y^*}{n_z^*}}{\delta u} \text{ with } \begin{bmatrix} n_x^* \\ n_y^* \\ n_z^* \end{bmatrix} = \mathbf{A} \begin{bmatrix} n_x \\ n_y \\ n_z \end{bmatrix}$$

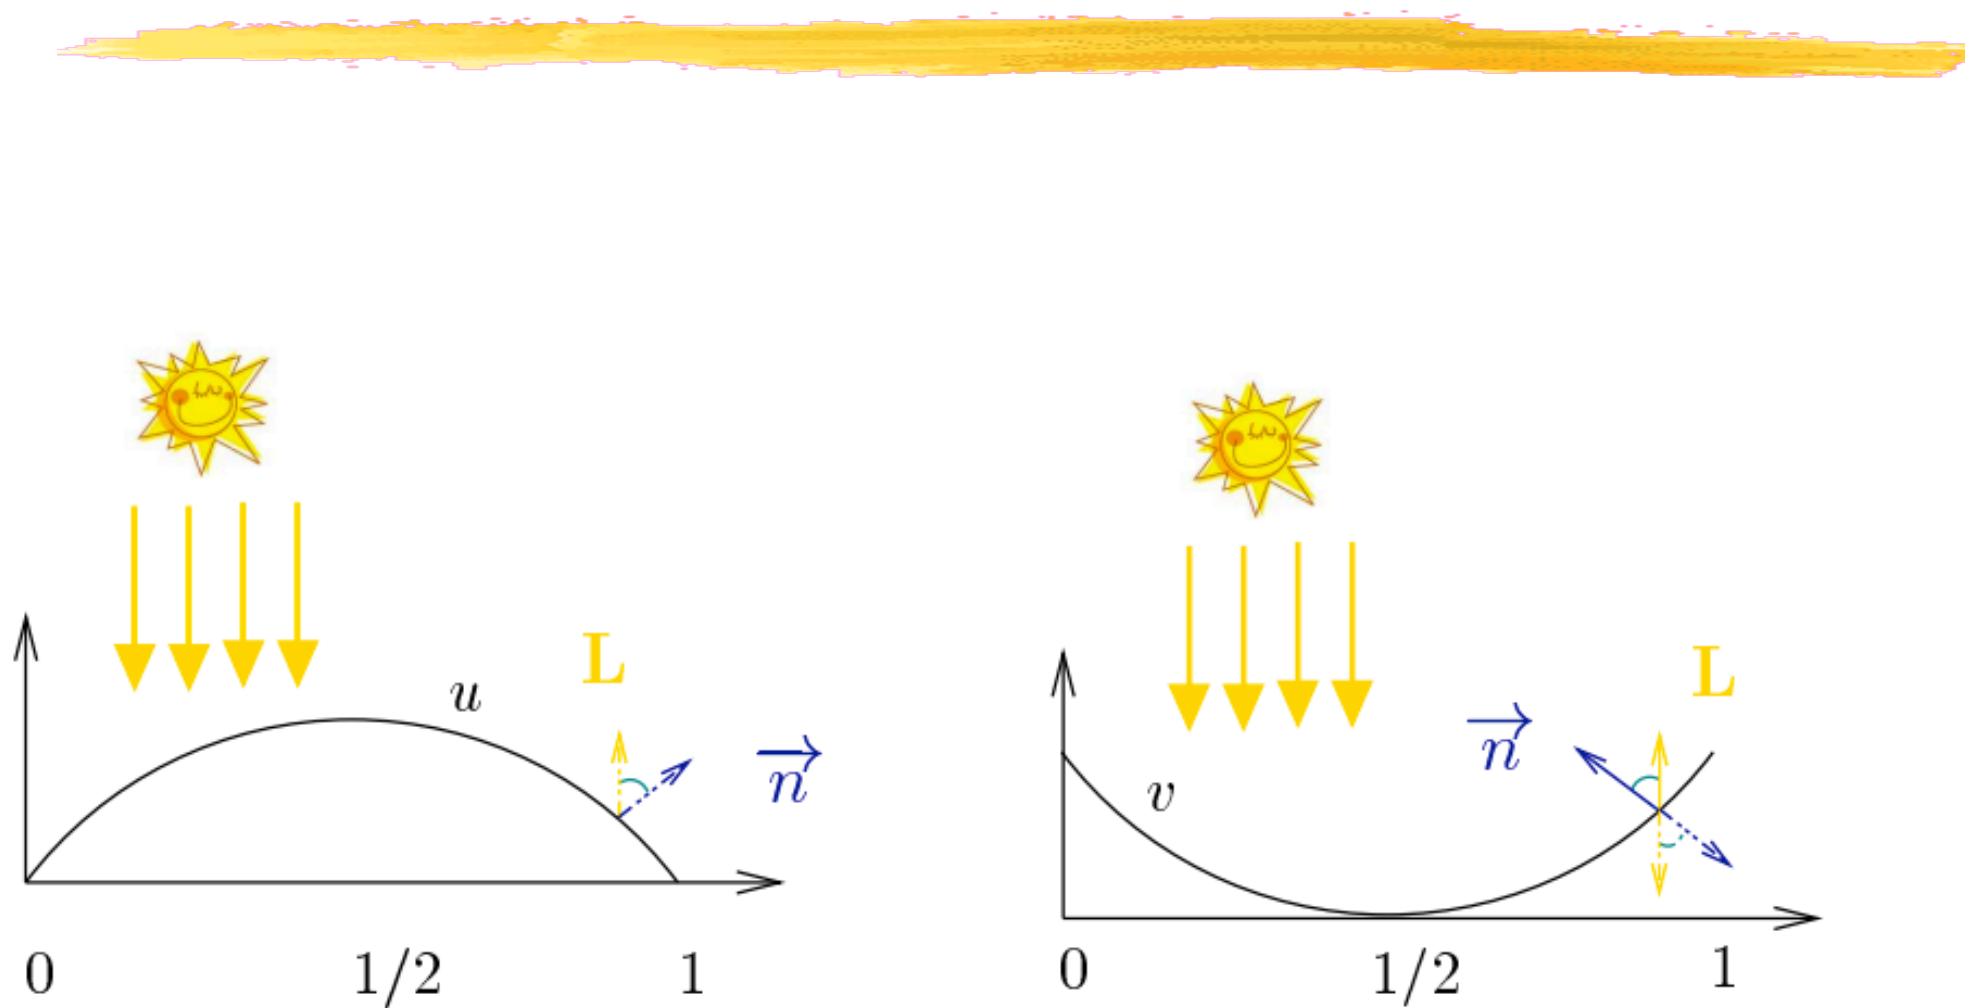
BAS-RELIEF AMBIGUITY

$$\left. \begin{array}{l} \frac{\delta z}{\delta u} = -\frac{n_x^*}{n_z^*} \\ \frac{\delta z}{\delta v} = -\frac{n_y^*}{n_z^*} \end{array} \right\} \Rightarrow \frac{\delta \frac{n_x^*}{n_z^*}}{\delta v} = \frac{\delta \frac{n_y^*}{n_z^*}}{\delta u} \text{ with } \begin{bmatrix} n_x^* \\ n_y^* \\ n_z^* \end{bmatrix} = \mathbf{A} \begin{bmatrix} n_x \\ n_y \\ n_z \end{bmatrix}$$

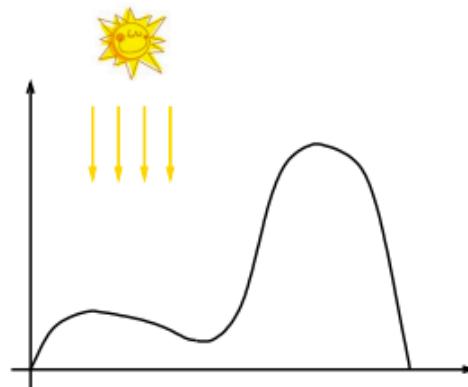
$$\Rightarrow \mathbf{A} \text{ restricted to } \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ \mu & \nu & 1 \end{bmatrix}$$

➤ The surface $f(u, v)$ can be changed into $\lambda f(u, v) + \mu u + \nu v$ and still produce the same image.

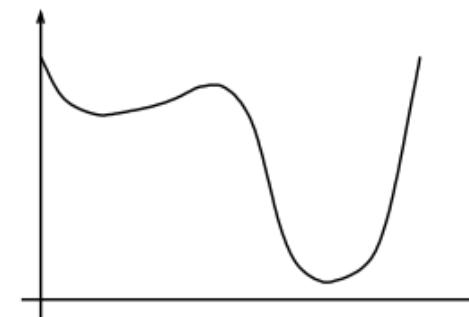
MORE AMBIGUITIES EVEN WHEN THE LIGHT SOURCE IS KNOWN...



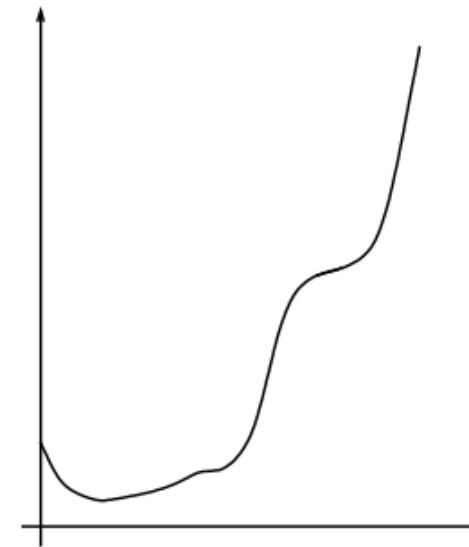
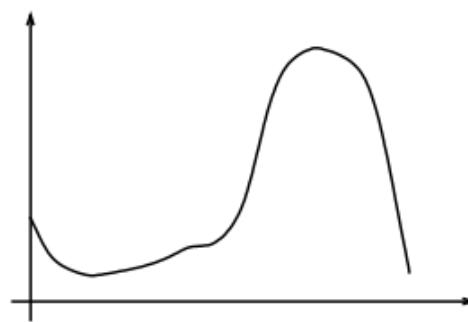
MORE AMBIGUITIES EVEN WHEN THE LIGHT SOURCE IS KNOWN...



a)



b)



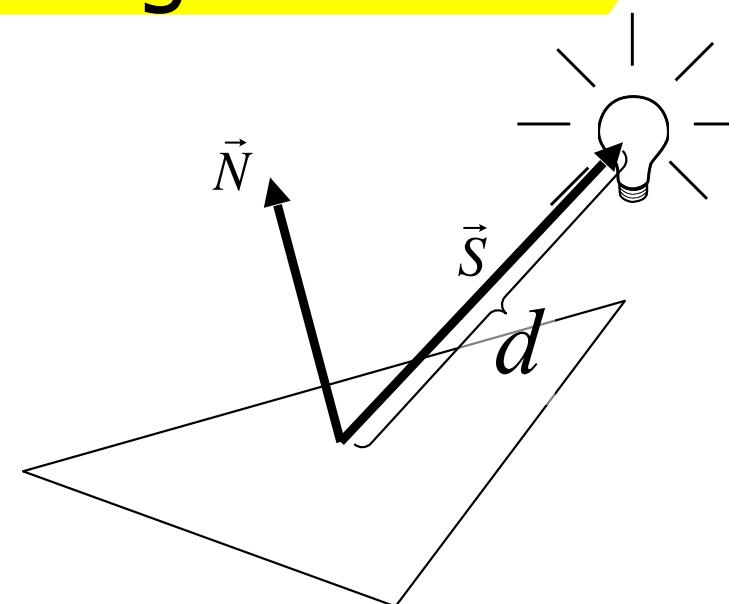
MAKING THE PROBLEM WELL-POSED

- Use perspective projection model;
- Radiance depends on distance to light source:

$$I = \frac{\text{Albedo} \cdot (\mathbf{N} \cdot \mathbf{S})}{d^2}$$

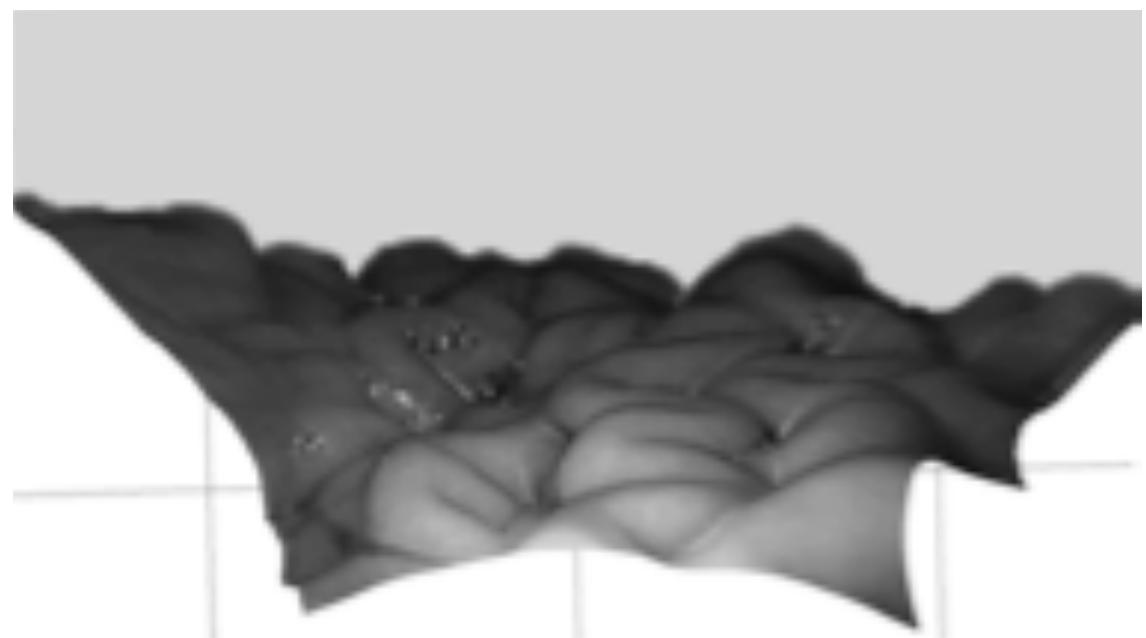
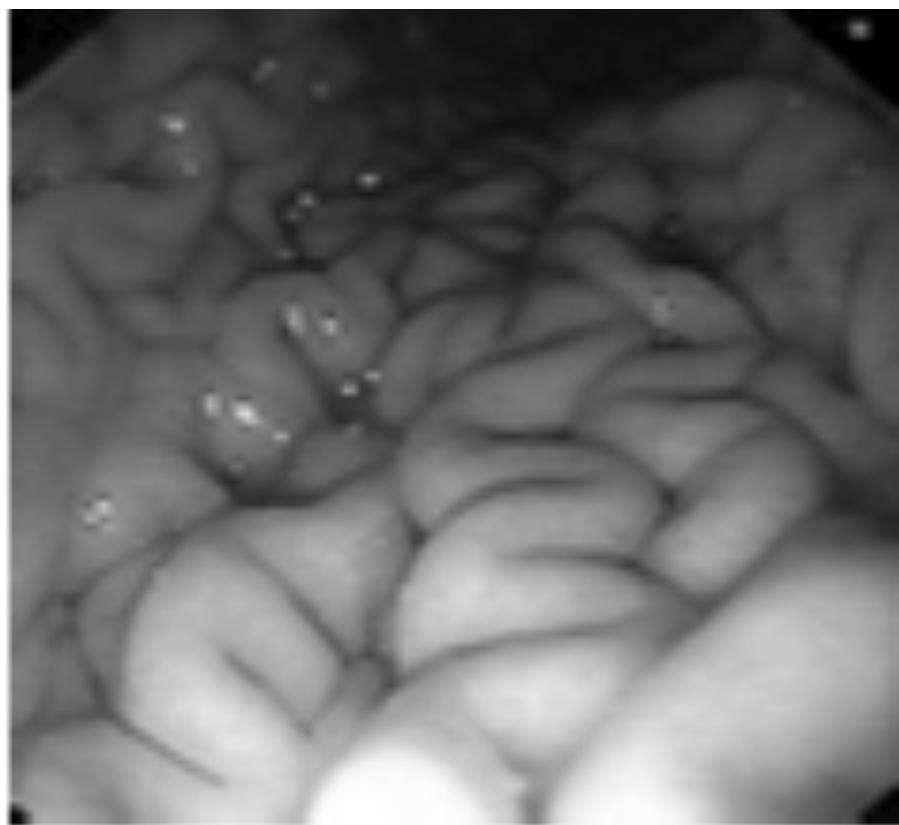
instead of

$$I = \text{Albedo} \cdot (\mathbf{N} \cdot \mathbf{S})$$

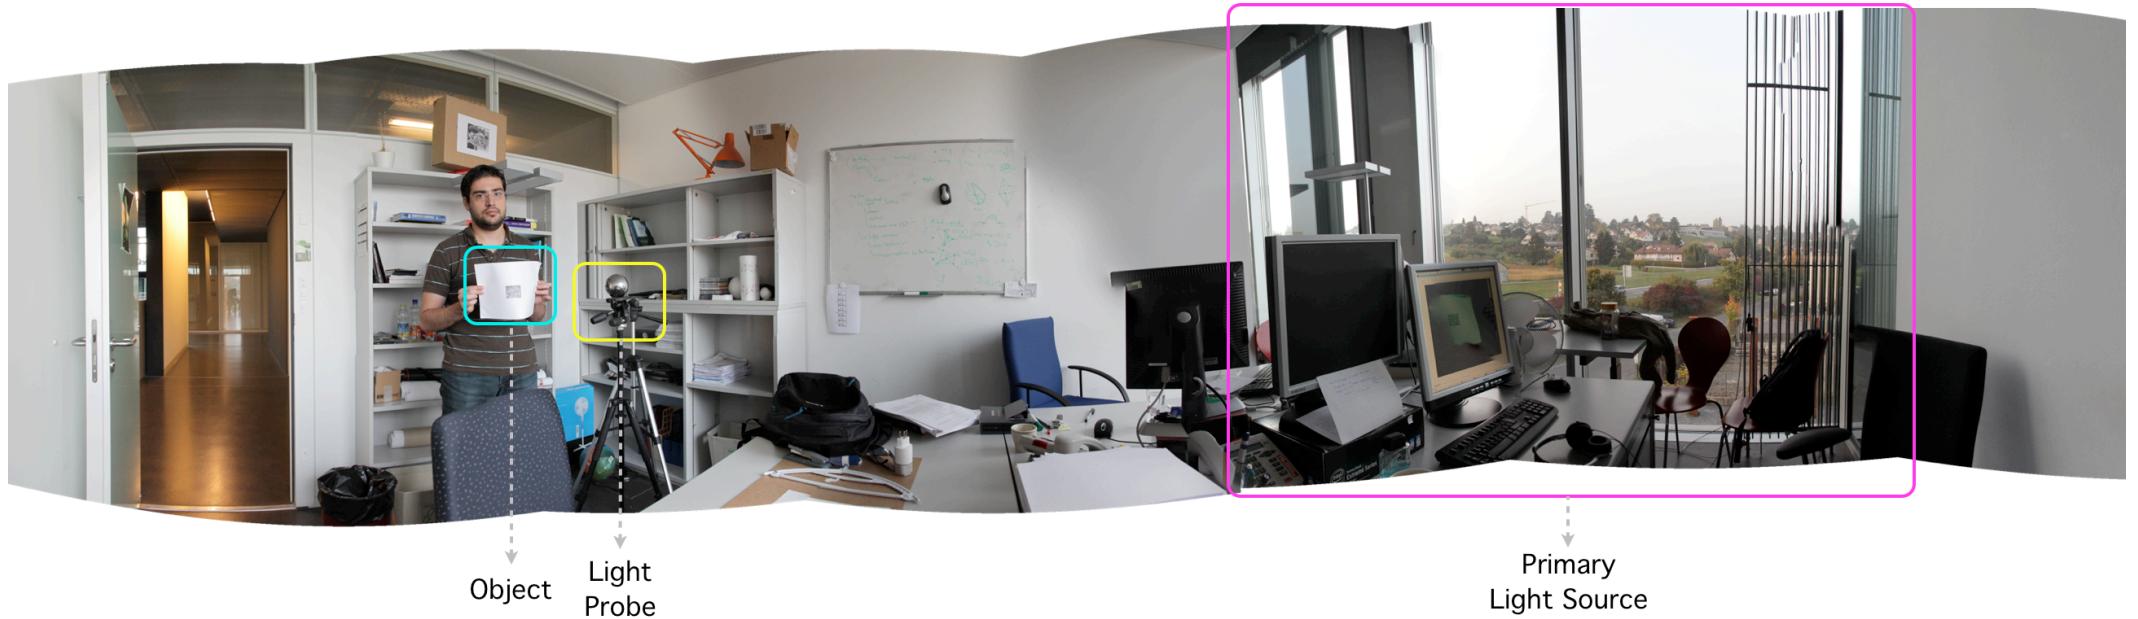


- Light source located at the optical center.
-> Unique solution.

ENDOSCOPY



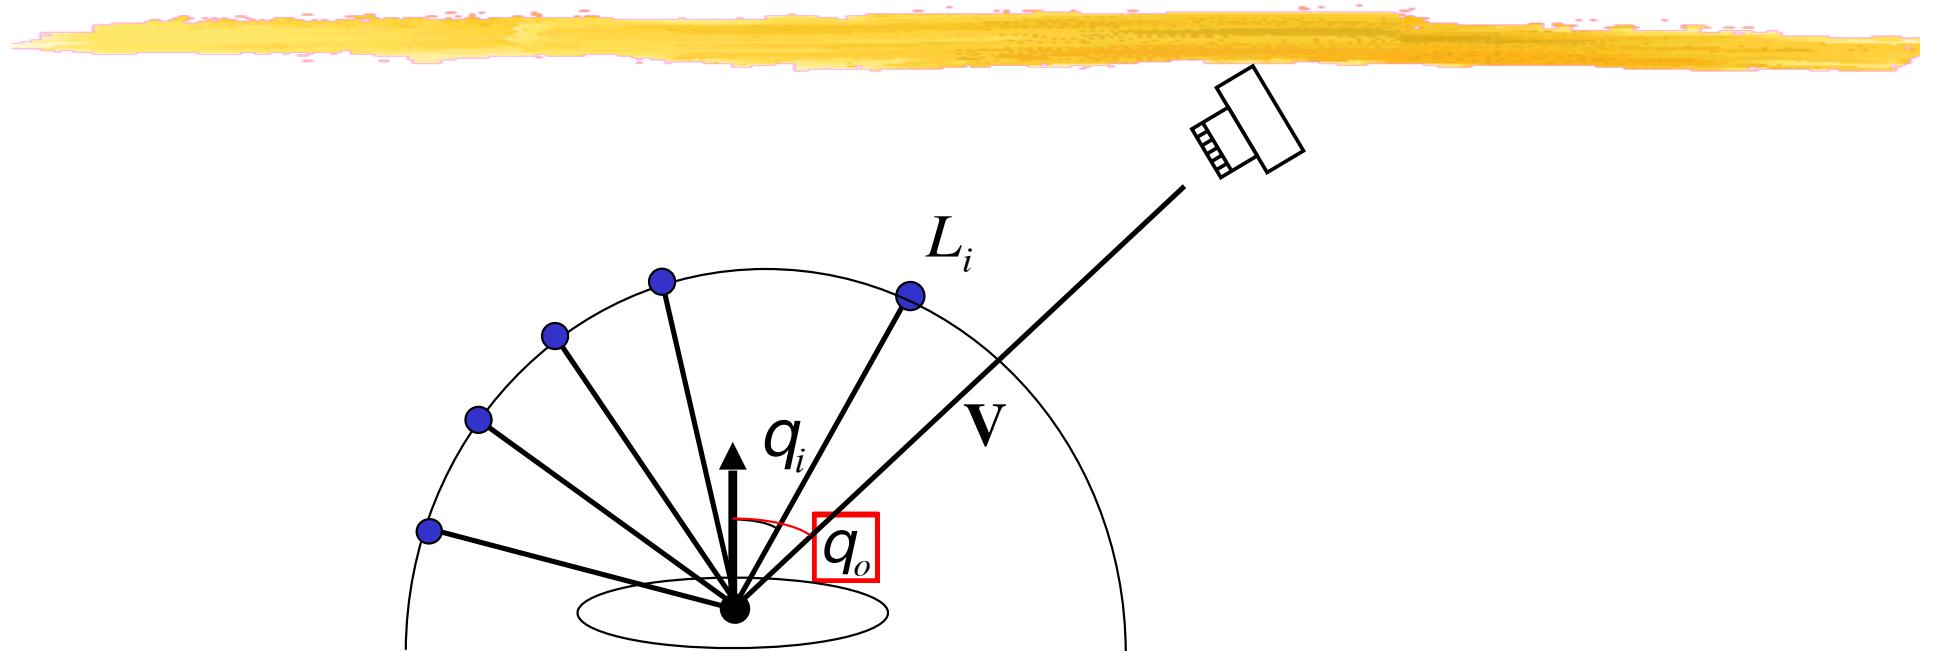
EVERYDAY SETTING



Multiple extended light sources:

- Illumination modeled as a weighted sum of spherical harmonics.
- Illumination parameters estimated using the light probe.

ILLUMINATION HEMISPHERE



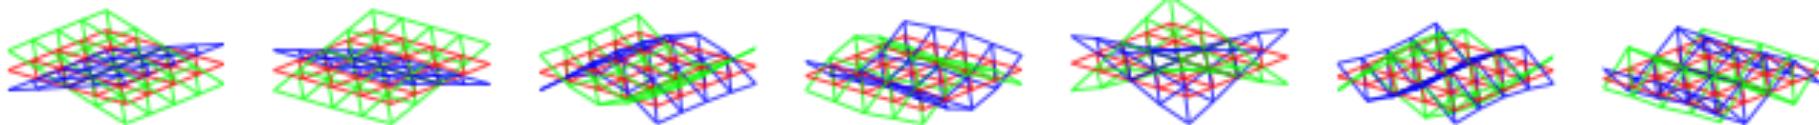
$$\begin{aligned}
 L_o(\mathbf{P}, q_o, f_o) &= \int_{\Omega_i} r_{bd}(q_o, f_o, q_i, f_i) L_i(\mathbf{P}, q_i, f_i) \cos(q_i) d\omega_i \\
 &= \int_{\Omega_i} L_i(q_i, f_i) r_{bd}(q_o, f_o, q_i, f_i) \max(0, \cos(q_i)) d\omega_i \\
 &= \int_{\Omega_i} L_i(q_i, f_i) r^*(q_i, f_i) d\omega_i
 \end{aligned}$$

with $r^*(q_i, f_i) = r_{bd}(q_o, f_o, q_i, f_i) \max(0, \cos(q_i))$ is the BRDF product function.

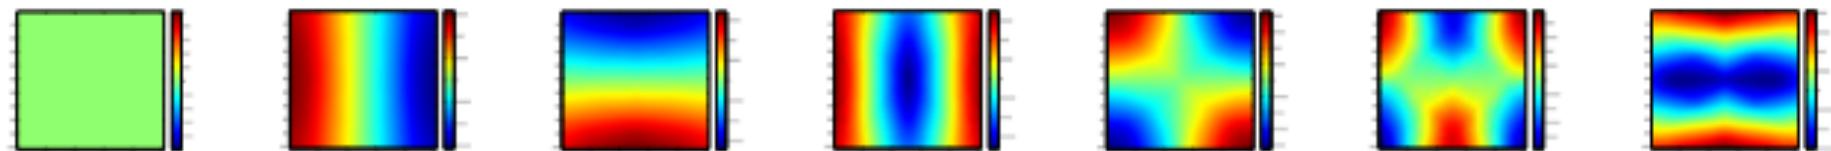
PATCH INTENSITY AND DEFORMATION MODES



Synthesize a training database containing deformed patches and the corresponding intensity patterns. Compute:

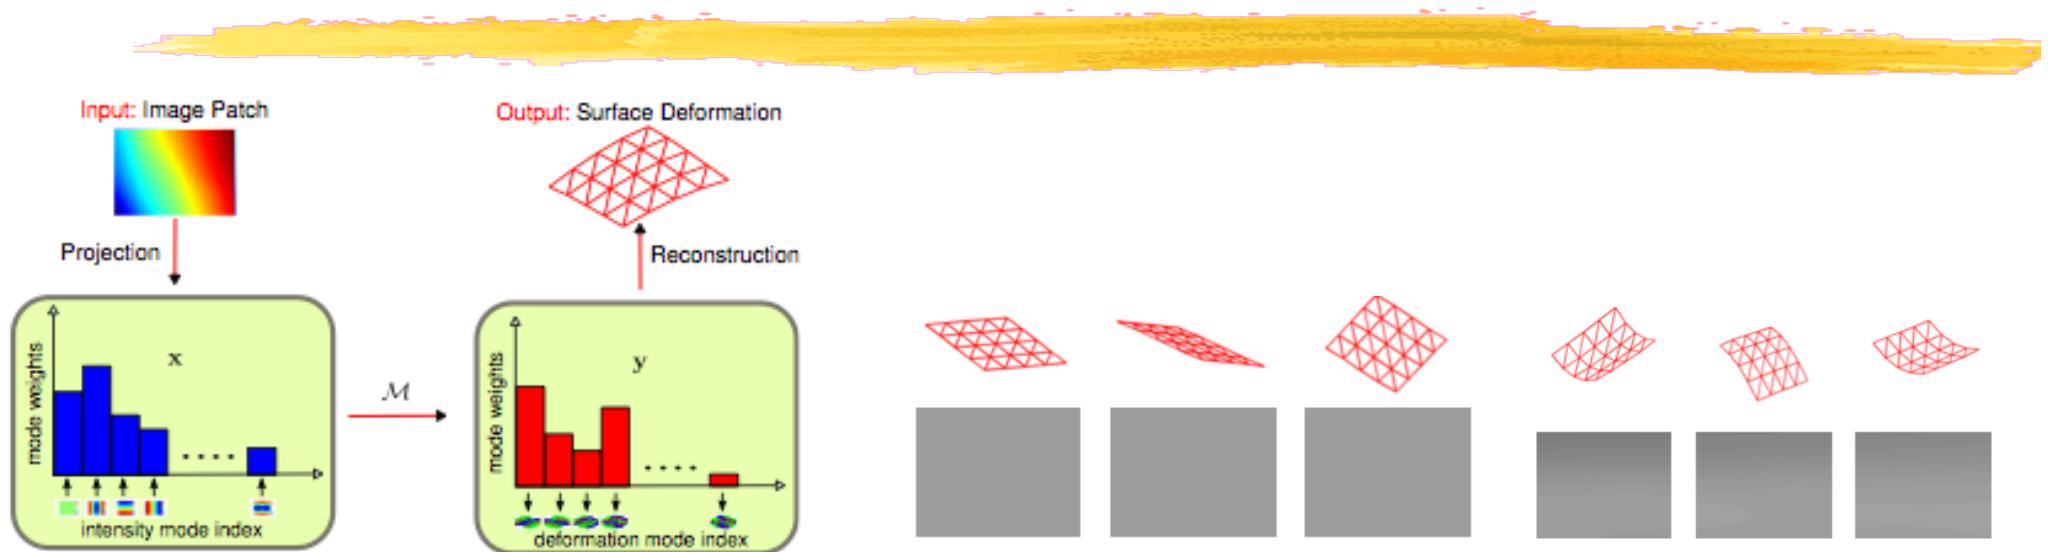


$$\text{Deformation modes: } D = D_0 + \sum_i y_i D_i$$



$$\text{Intensity modes: } I = I_0 + \sum_i x_i I_i$$

AMBIGUOUS MAPPING



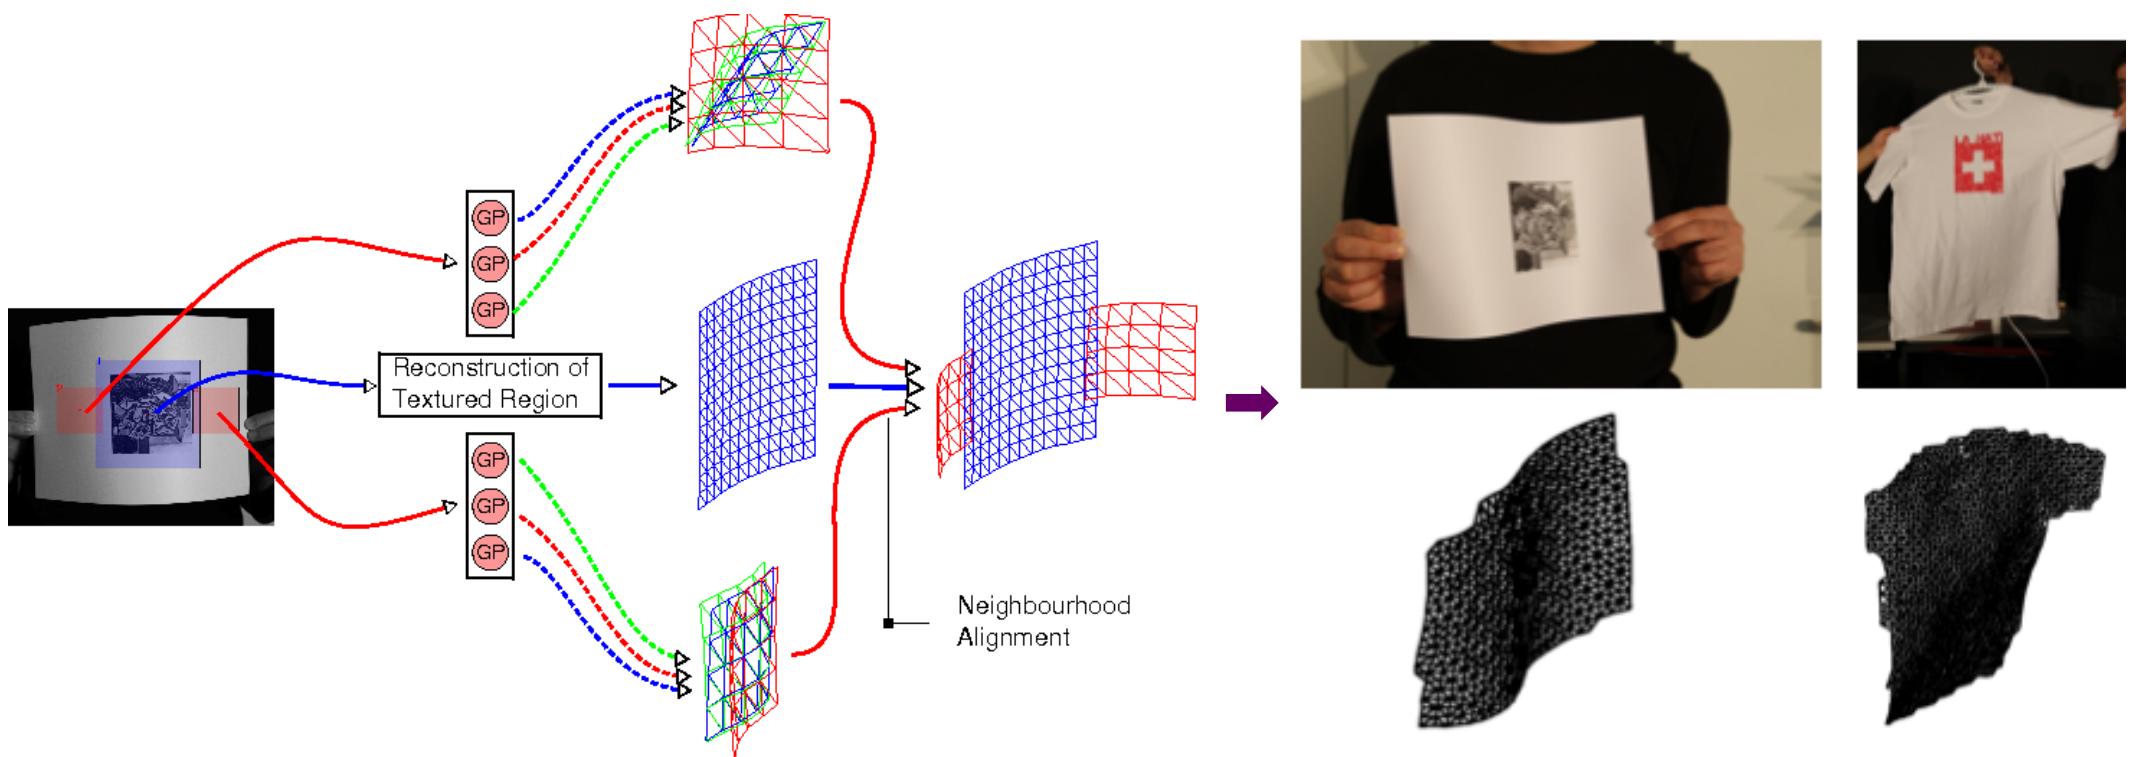
Mapping from intensity to surface deformation.

Unfortunately, the mapping is not one to one.

Algorithm:

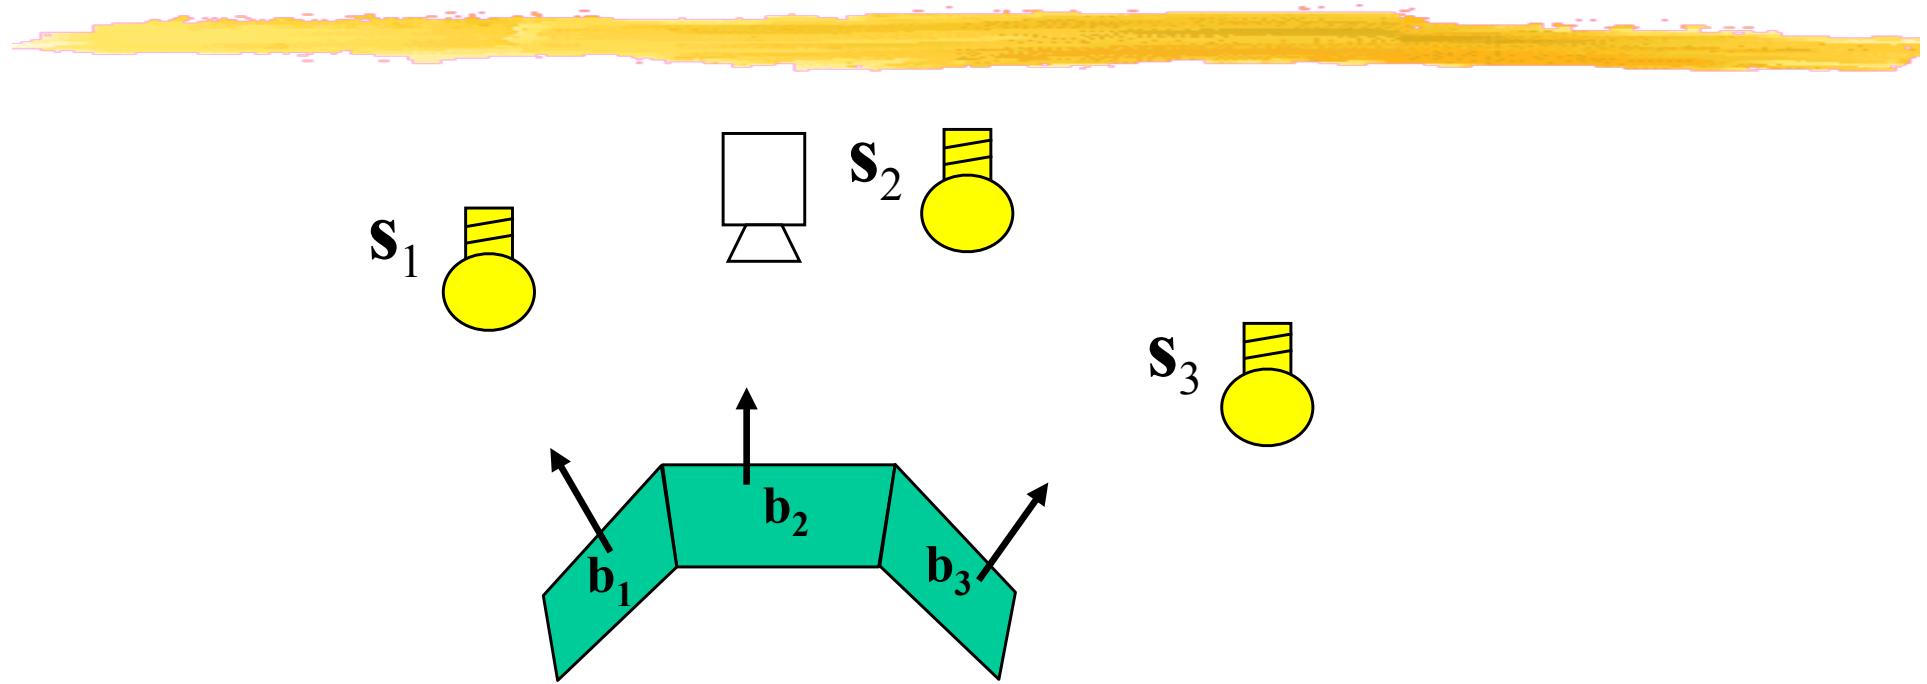
- Partition the training database according to average normal to learn an unambiguous mapping.
- At run-time, predict several potential shapes for each image-patch and use a Markov-Random field to pick a set of consistent interpretations.

ALGORITHMIC FLOW



→ The light environment and the camera need to be very carefully calibrated but there is hope ...

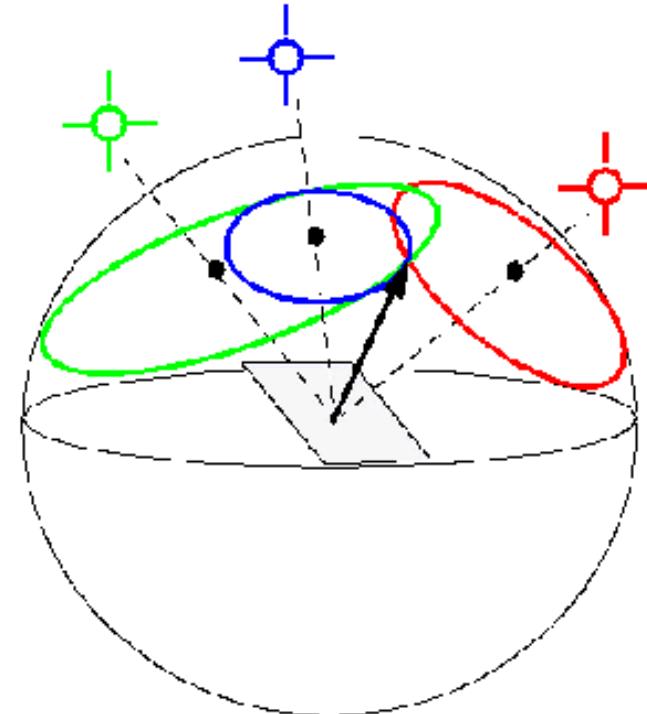
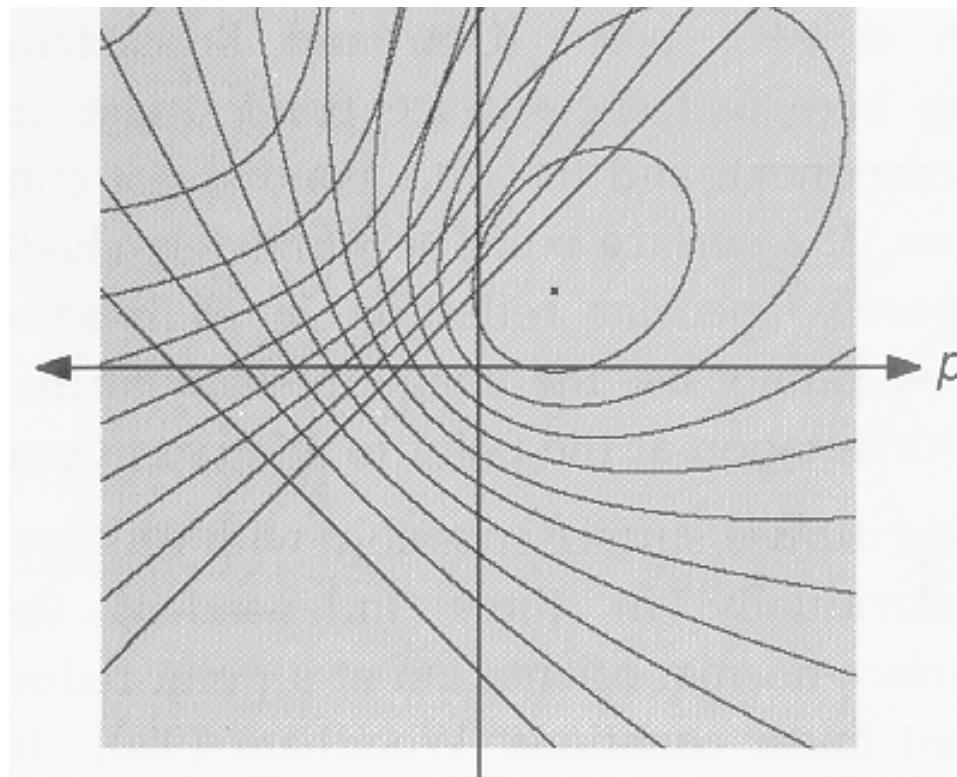
PHOTOMETRIC STEREO



Given multiple images of the same surface under different known lighting conditions, can we recover the surface shape?

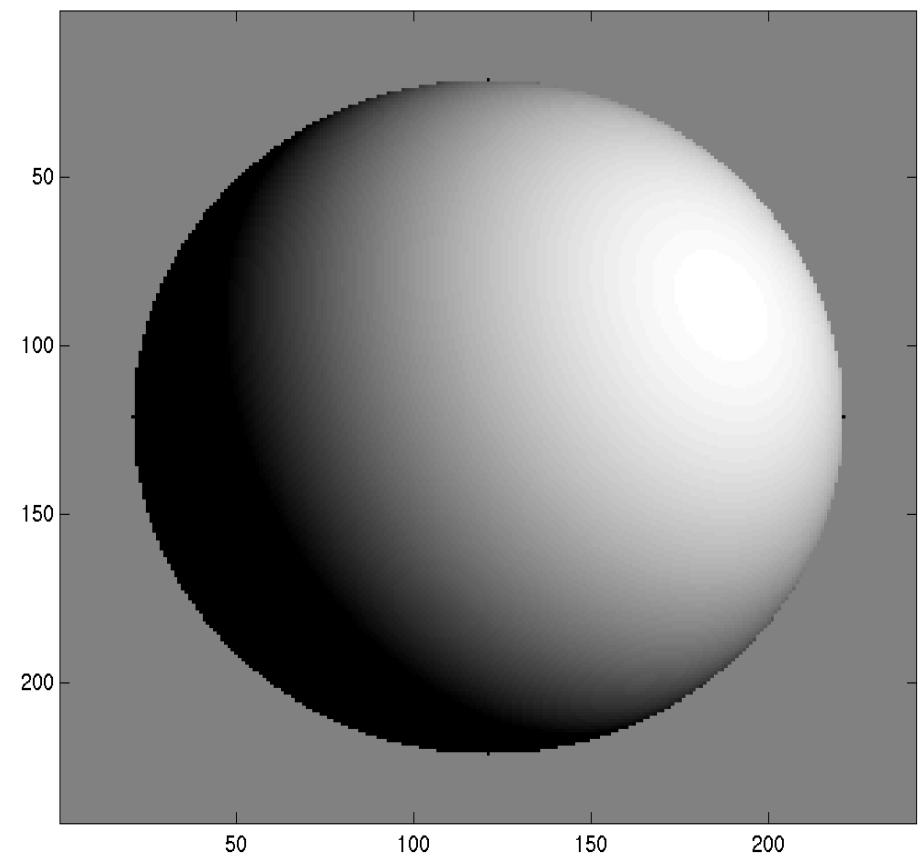
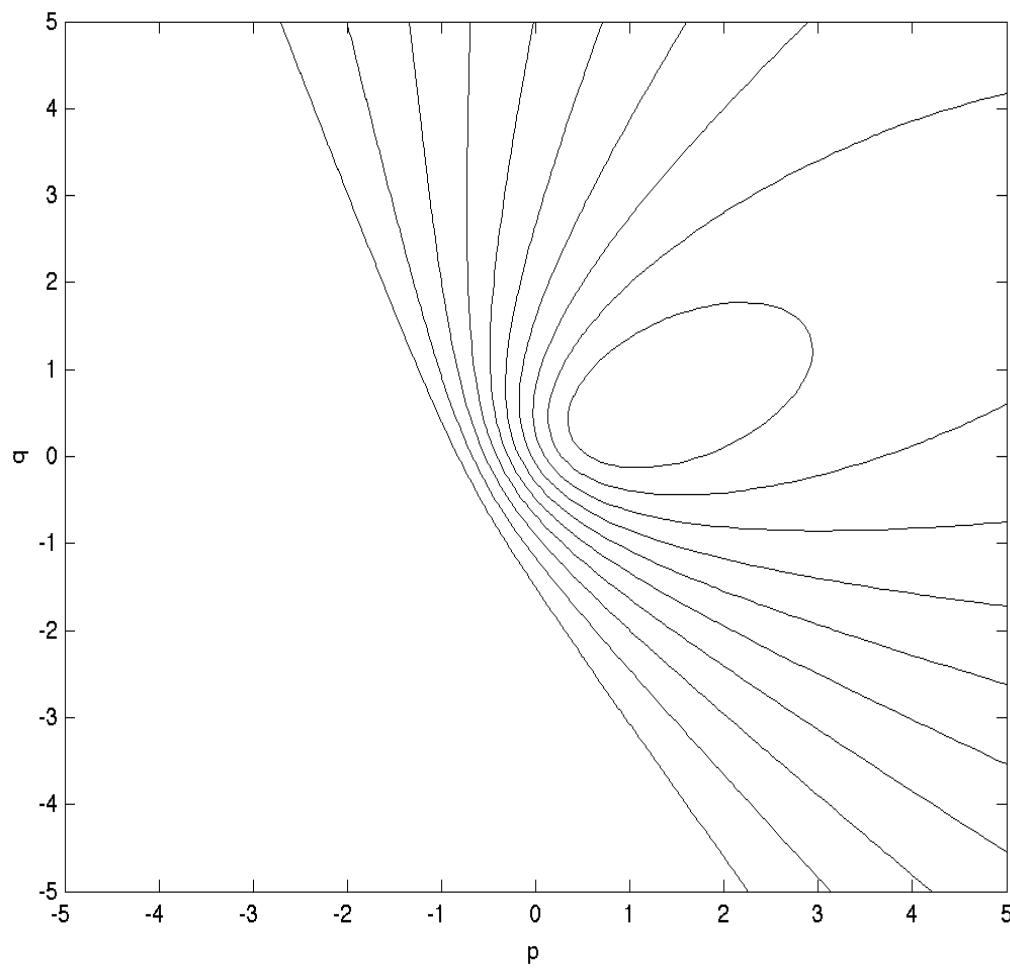
Yes! (Woodham, 1978)

PHOTOMETRIC STEREO

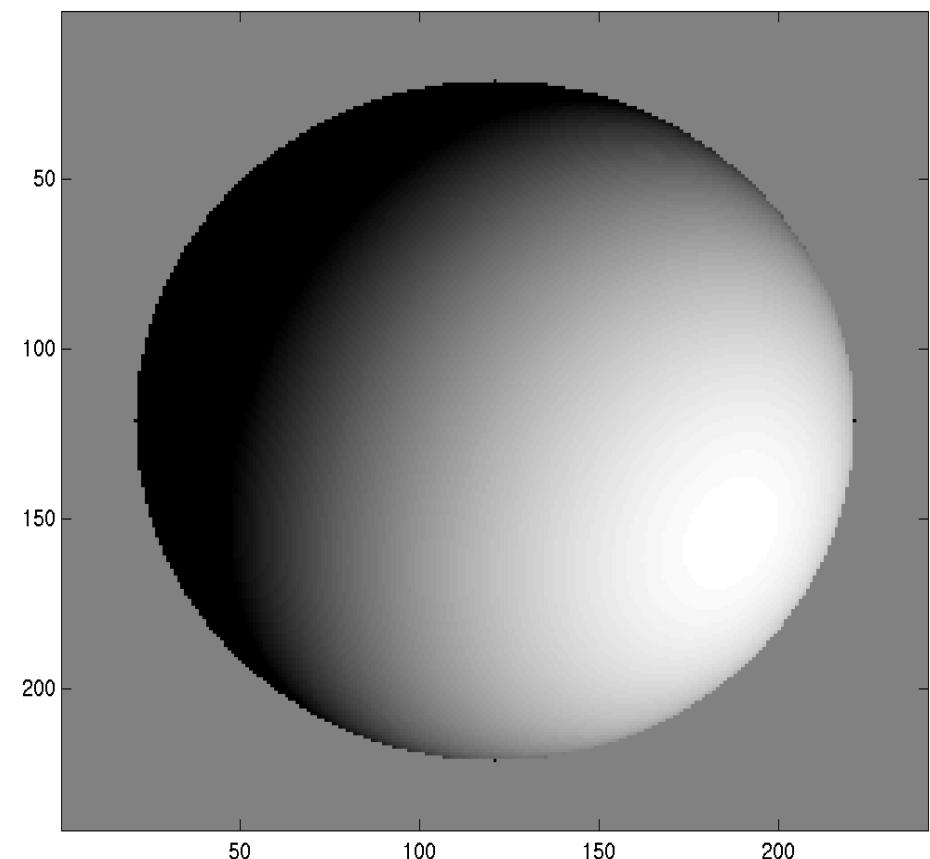
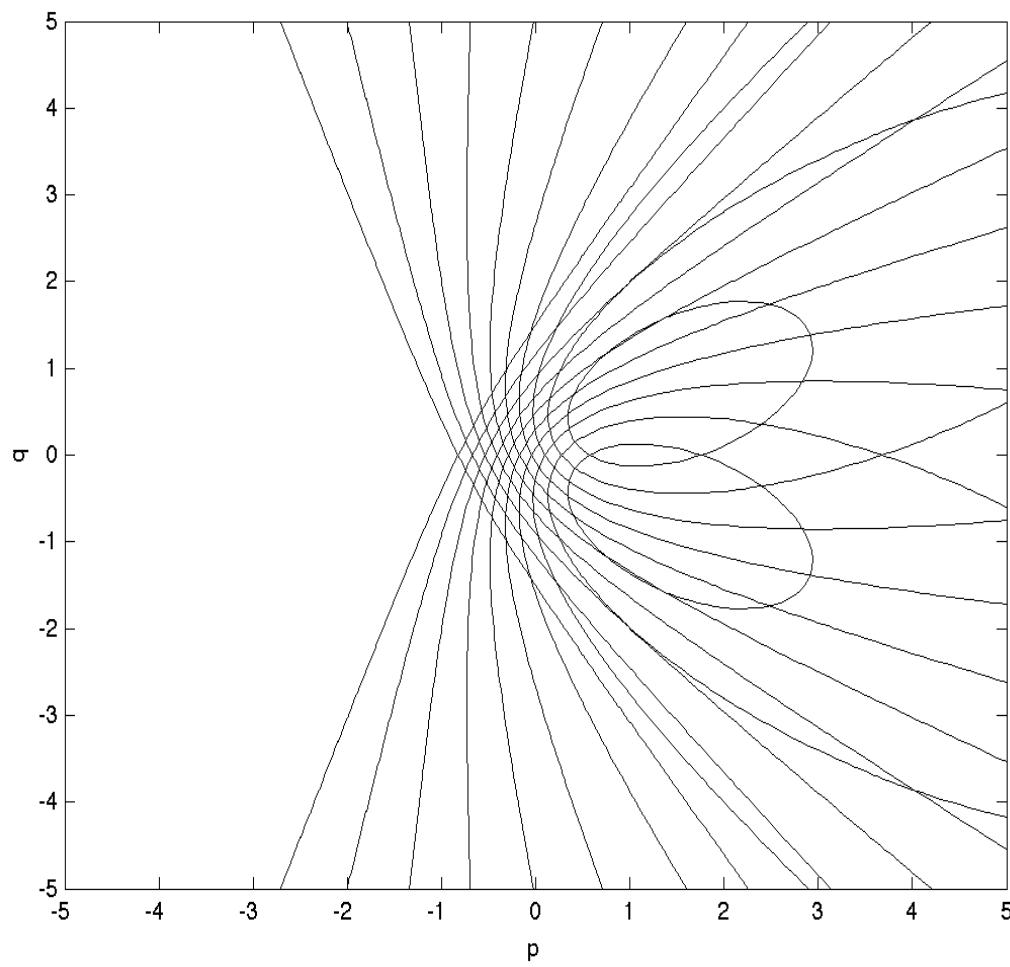


- Take several images under different lighting conditions.
- Infer the orientation from the changes in illumination.

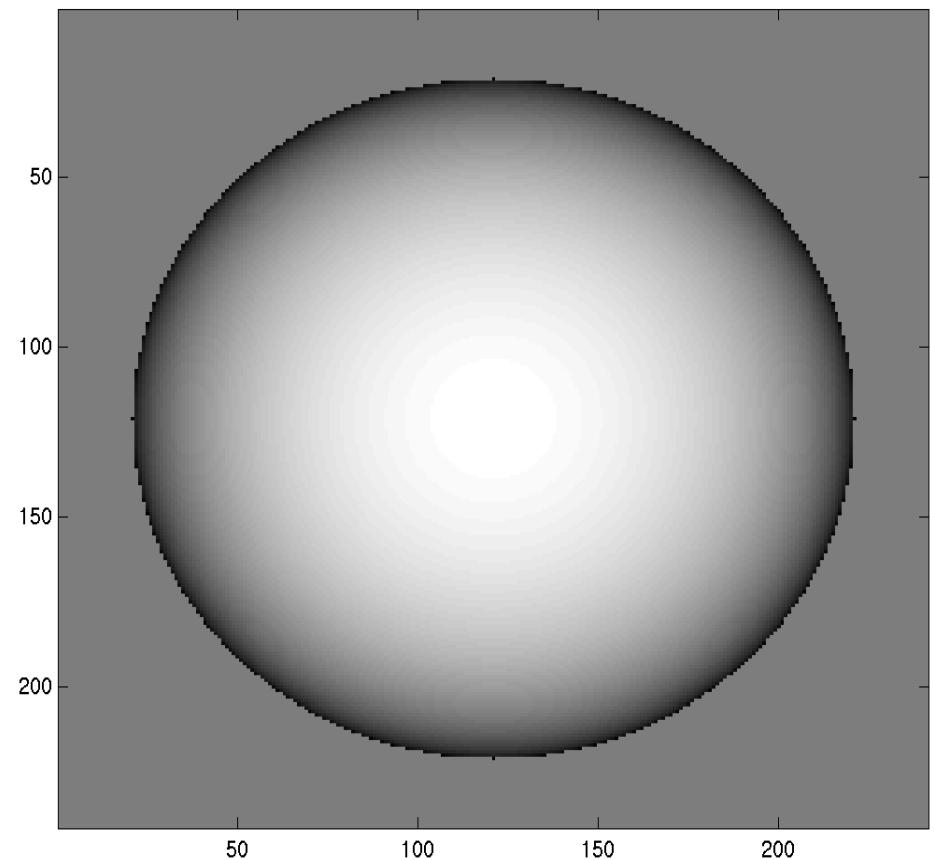
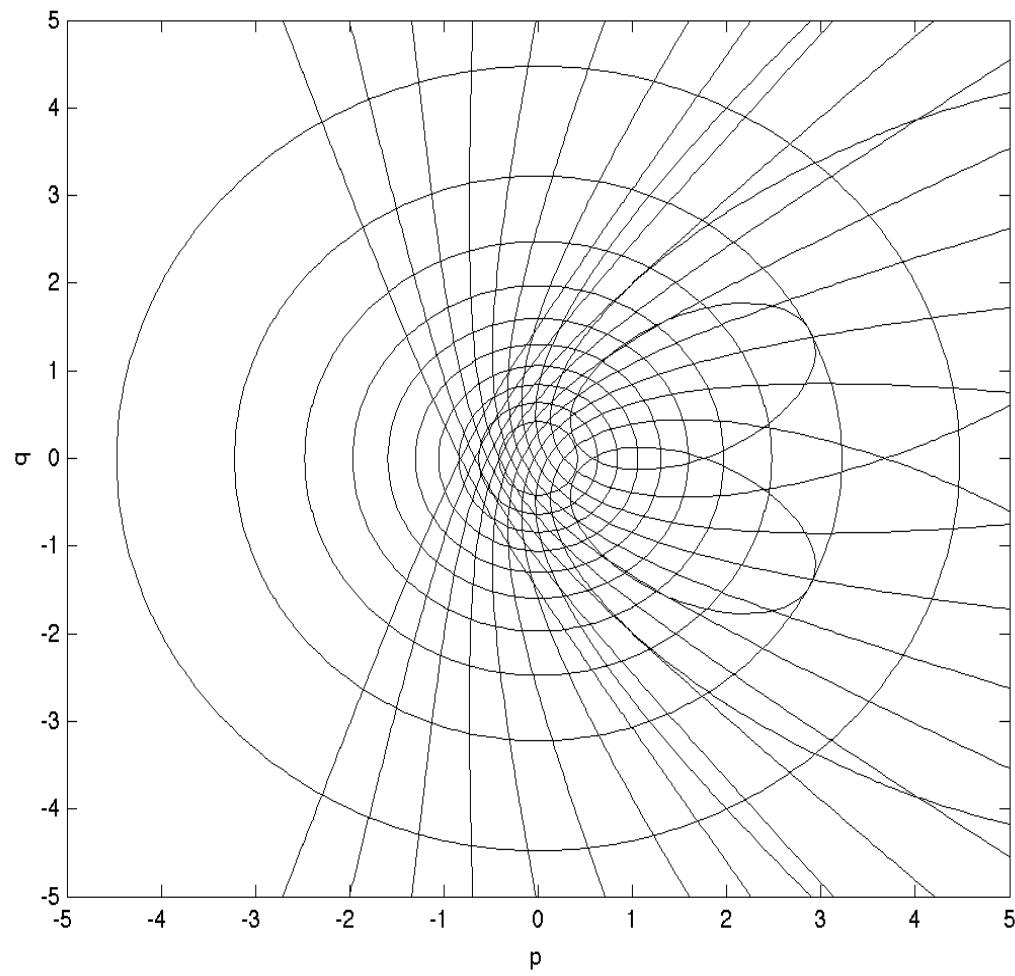
PHOTOMETRIC STEREO



PHOTOMETRIC STEREO



PHOTOMETRIC STEREO



ALGEBRAIC FORMULATION

Lambertian model: $I = \alpha(\mathbf{L} \cdot \mathbf{N}) = (\mathbf{L} \cdot \mathbf{M})$

Three light sources:

$$\begin{bmatrix} I_1 \\ I_2 \\ I_3 \end{bmatrix} = \begin{bmatrix} \mathbf{L}_1 \\ \mathbf{L}_2 \\ \mathbf{L}_3 \end{bmatrix} \mathbf{M}$$

$$\mathbf{N} = \frac{\mathbf{M}}{||\mathbf{M}||}$$

$$\alpha = ||\mathbf{M}||$$

ADDITIONAL LIGHTS

Over-constrained problem:

$$\mathbf{I} = \mathbf{LM}, \text{ with } \mathbf{I} = \begin{bmatrix} I_1 \\ \vdots \\ I_n \end{bmatrix} \text{ and } \mathbf{L} = \begin{bmatrix} \mathbf{L}_1 \\ \vdots \\ \mathbf{L}_n \end{bmatrix}$$

$$\Rightarrow \mathbf{LL^tM} = \mathbf{L^tI} \text{ (Least - squares solution)}$$

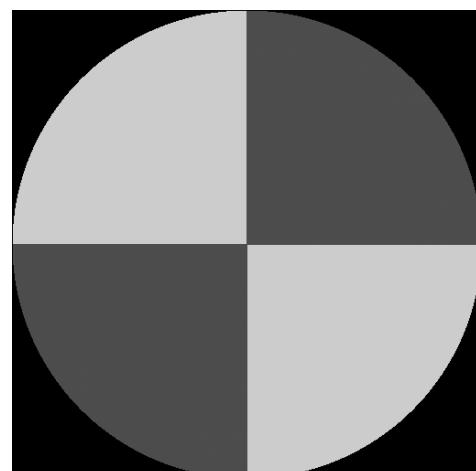
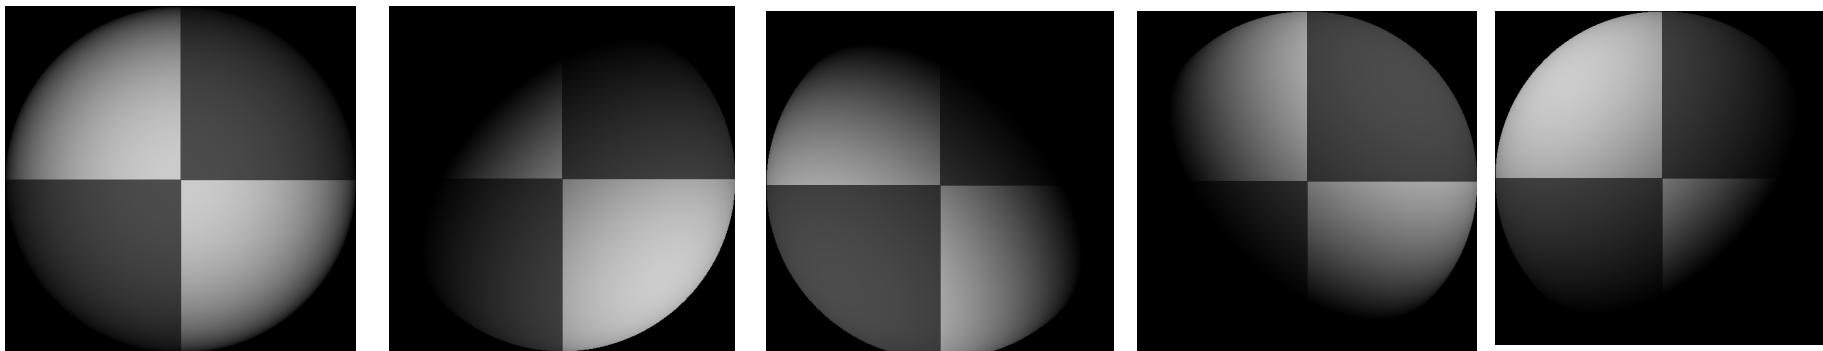
SHADOWS

- Shadowed pixels for a given light source position are outliers.
- Premultiplying by a thresholded weight matrix eliminates their contributions.

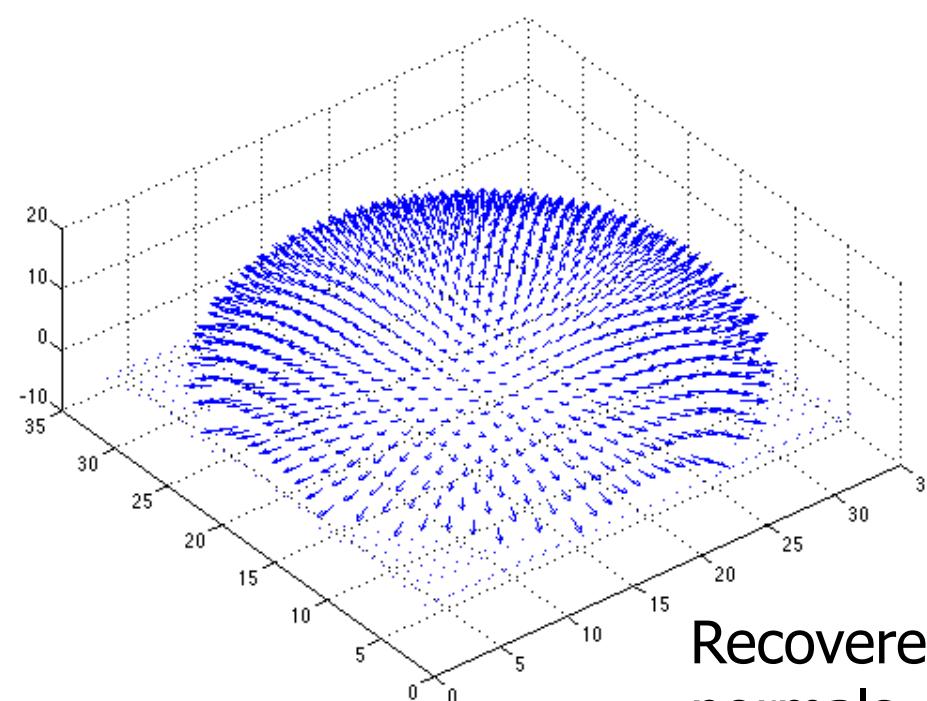
$$\begin{bmatrix} I_1 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & I_n \end{bmatrix} \mathbf{I} = \begin{bmatrix} I_1 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & I_n \end{bmatrix} \mathbf{LM}$$

SYNTHETIC SPHERE IMAGES

Five different lighting conditions:

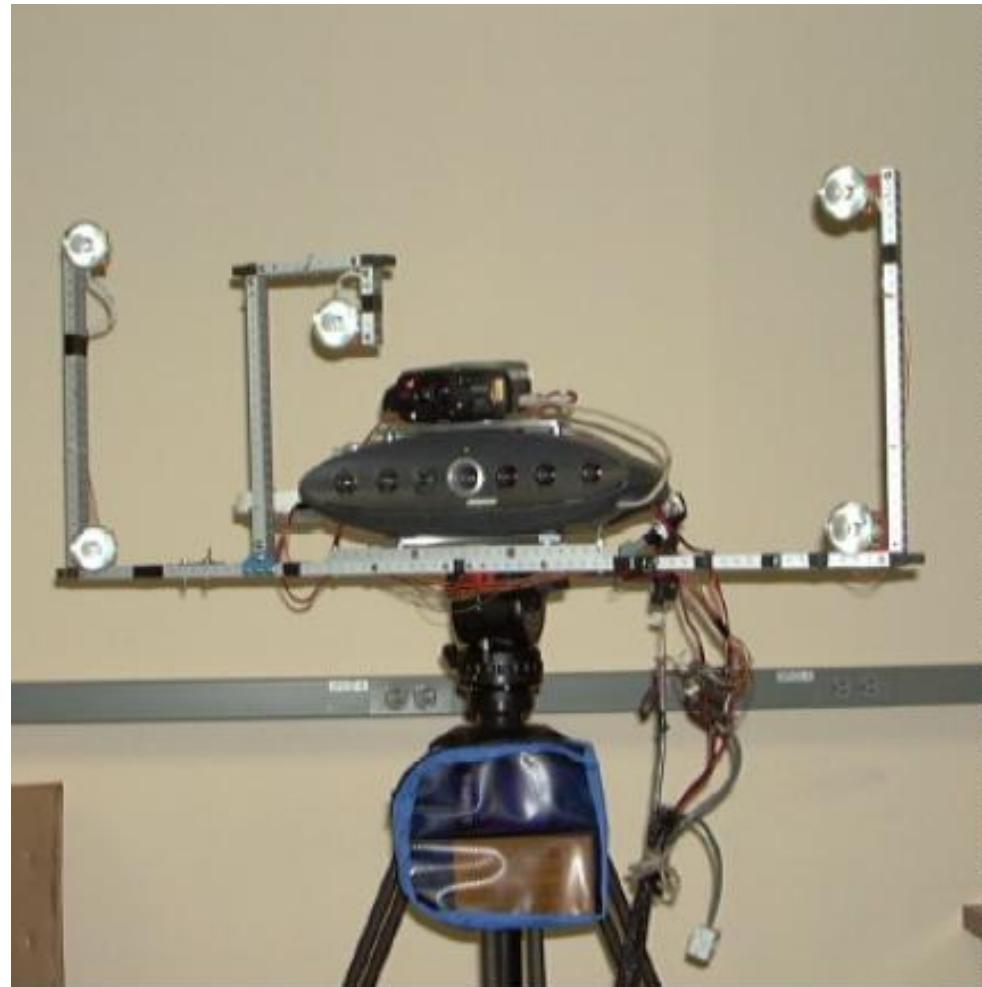


Recovered albedo



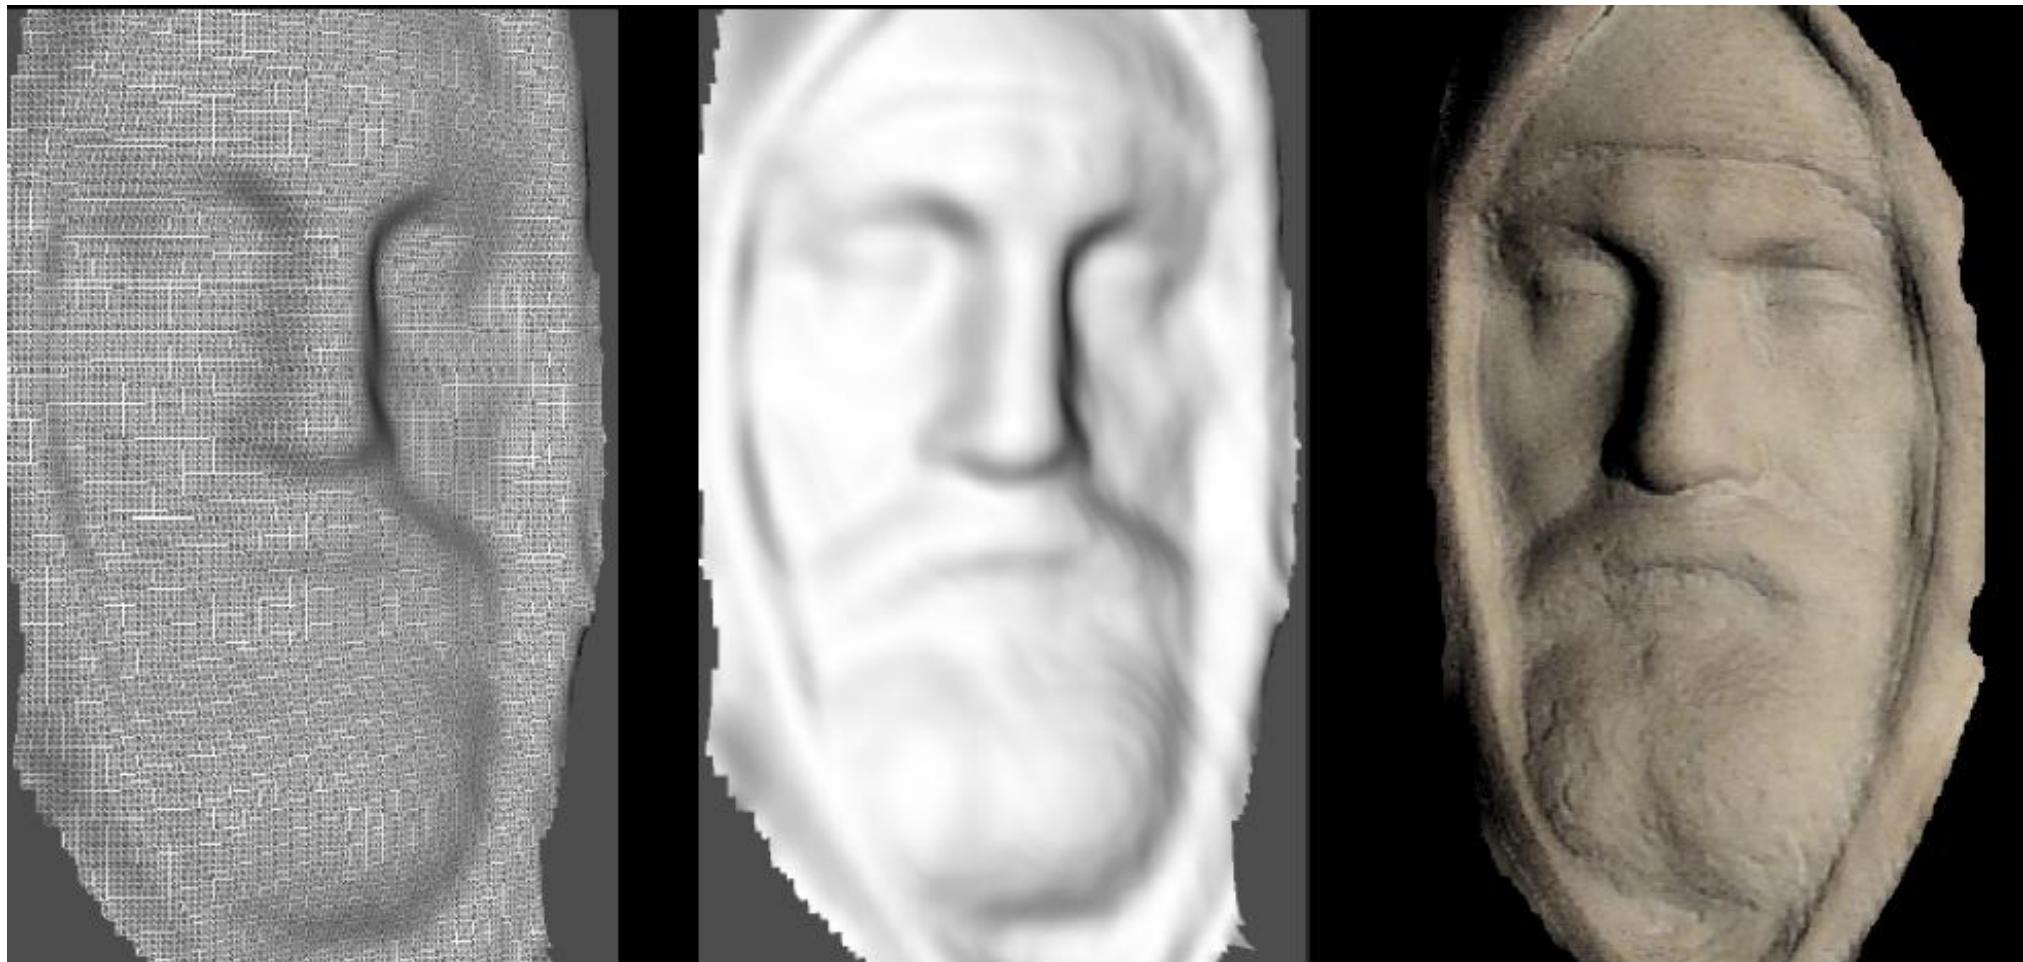
Recovered surface
normals

VIRTUOSO

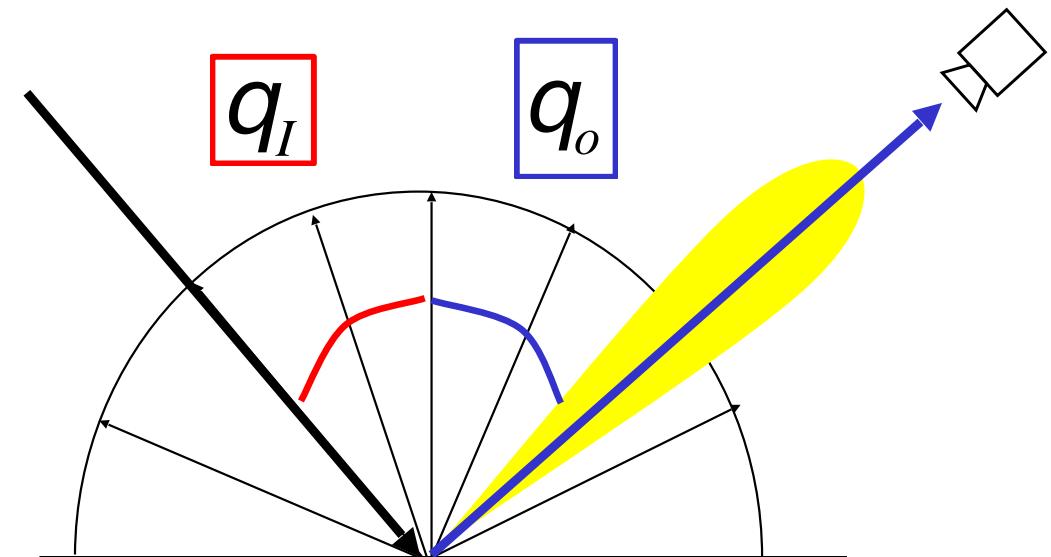


One camera and five light sources

DELIGHTED TEXTURE MAPS



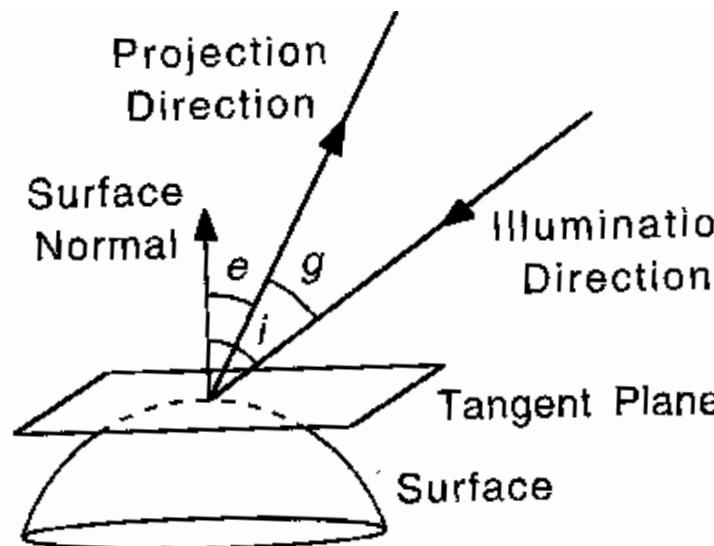
SPECULARITIES



places where the surface behaves like a mirror. normal must be such that reflection ray is symmetric

- At specular points Lambertian assumptions are violated.
- However, they can be used to infer normal information.

ANGLES

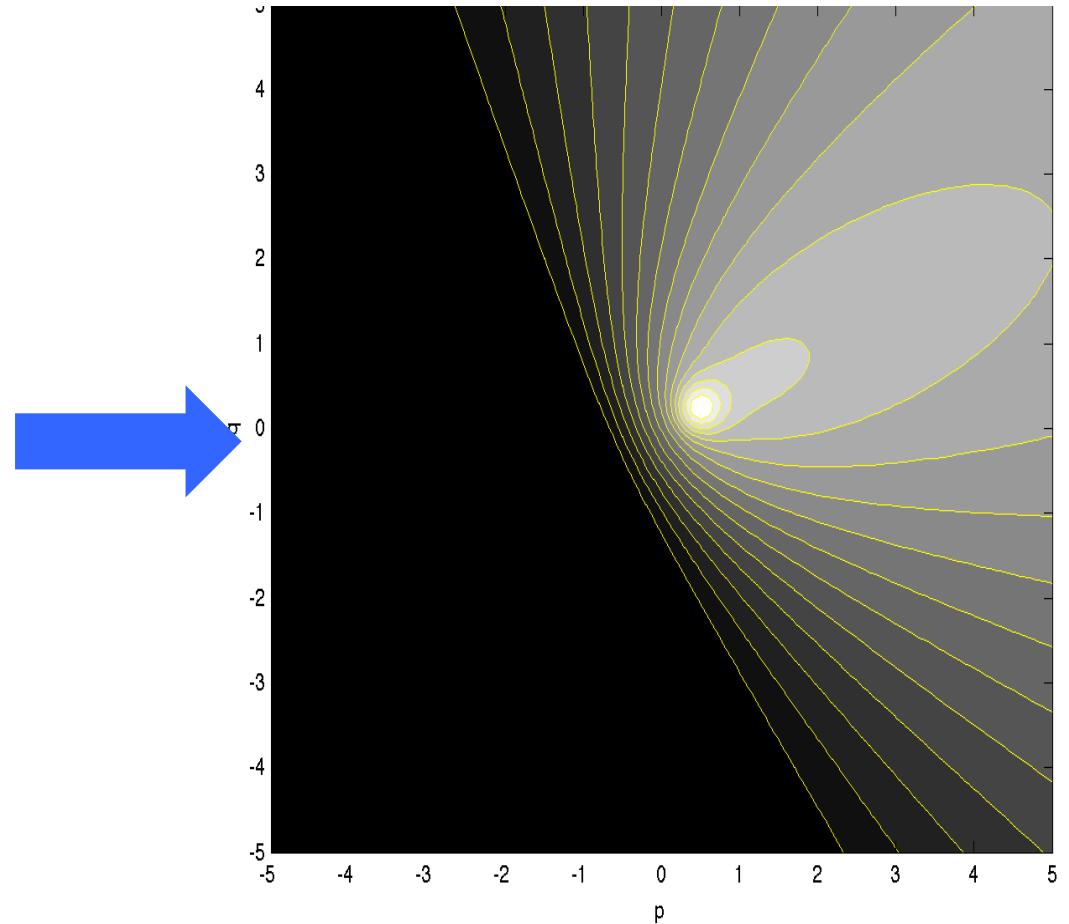
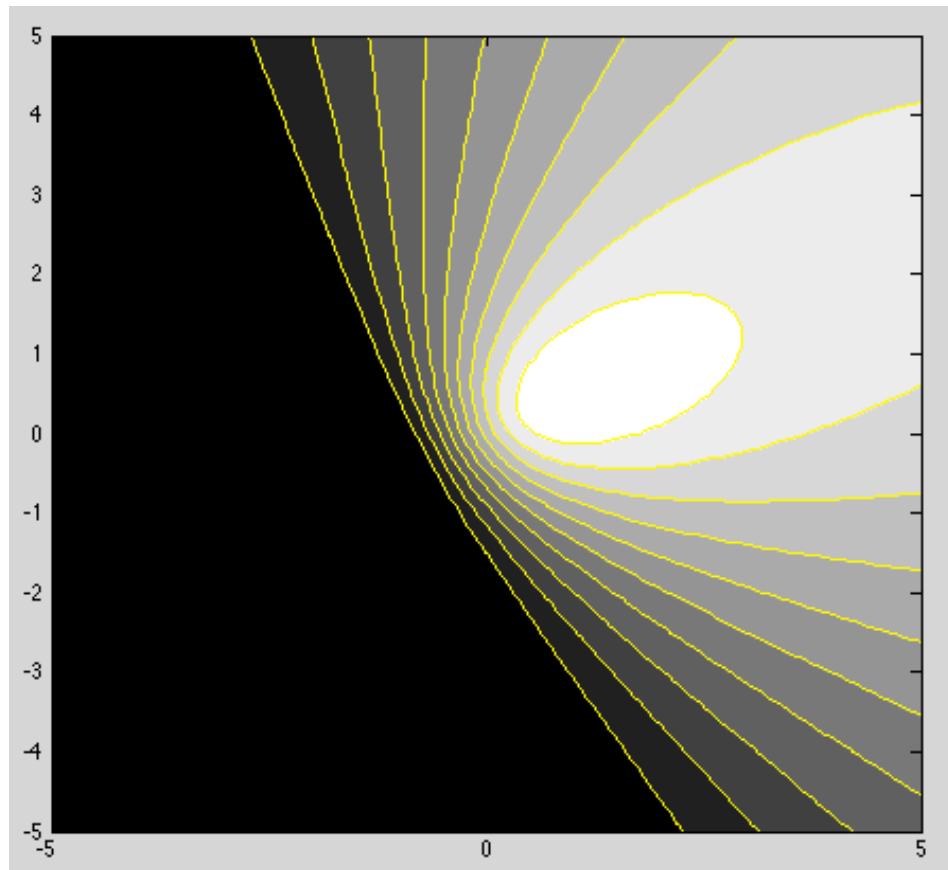


Angle of incidence (i): angle between the surface normal and the direction of the incident light ray.

Angle of emittance (e): angle between emitted light ray and surface normal.

Phase angle (g): angle between incident and emitted light ray.

LAMBERTIAN+SPECULAR REFLECTANCE MAP



Weighted average of the individual diffuse- and specular-component of glossy white paint.

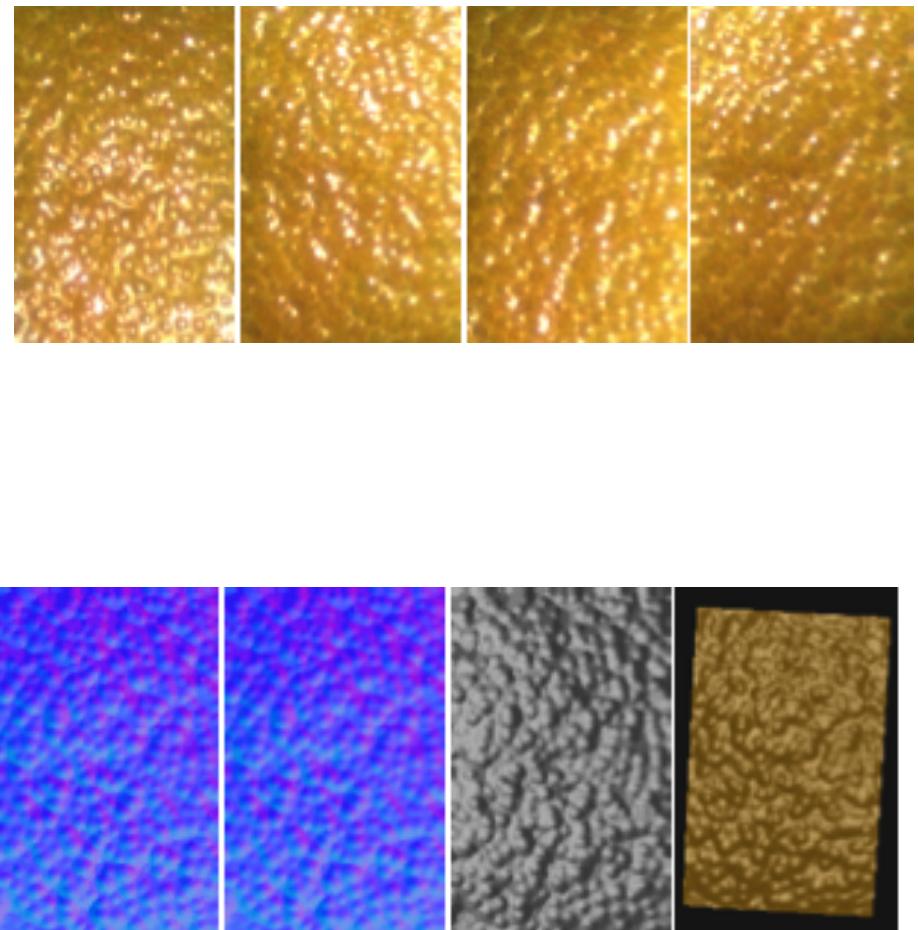
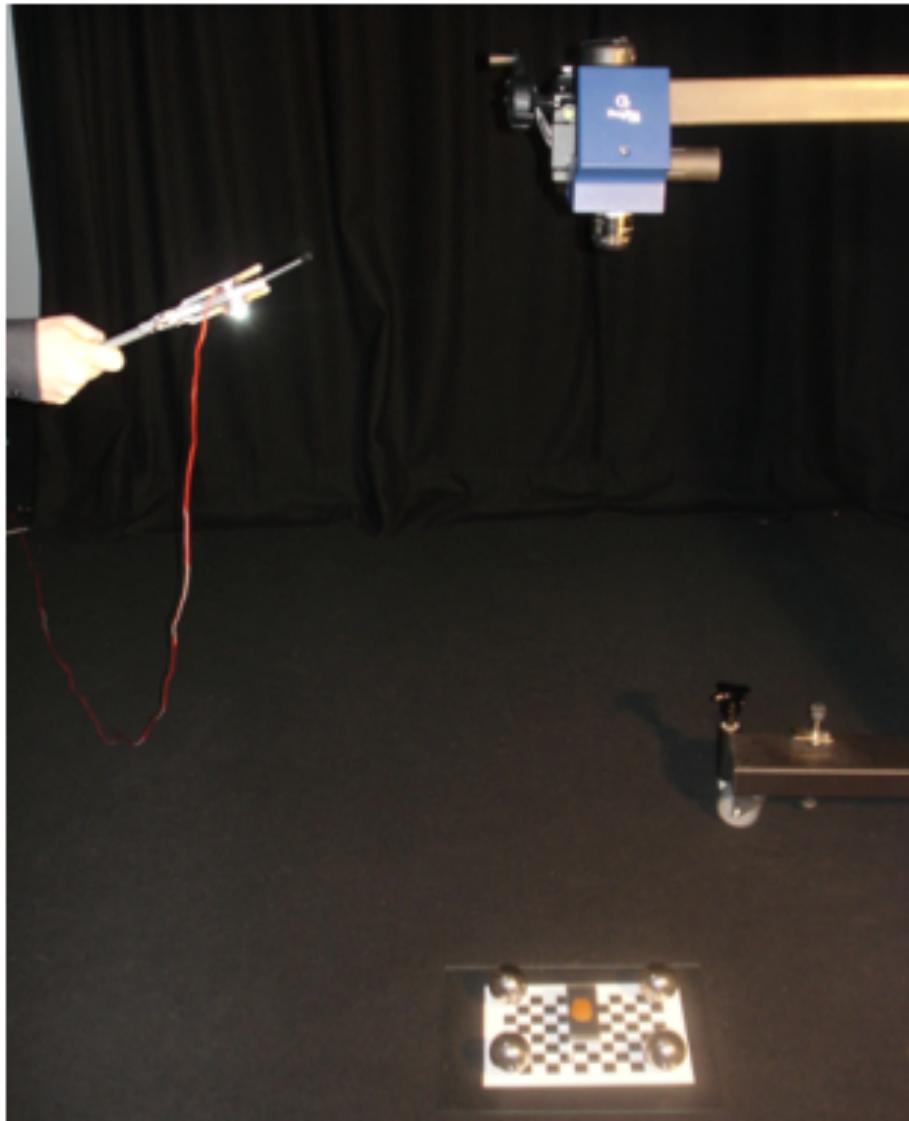
SPECULAR SURFACE



Perfect mirror: Reflects light only when $i = e$ and $i + e = g$.

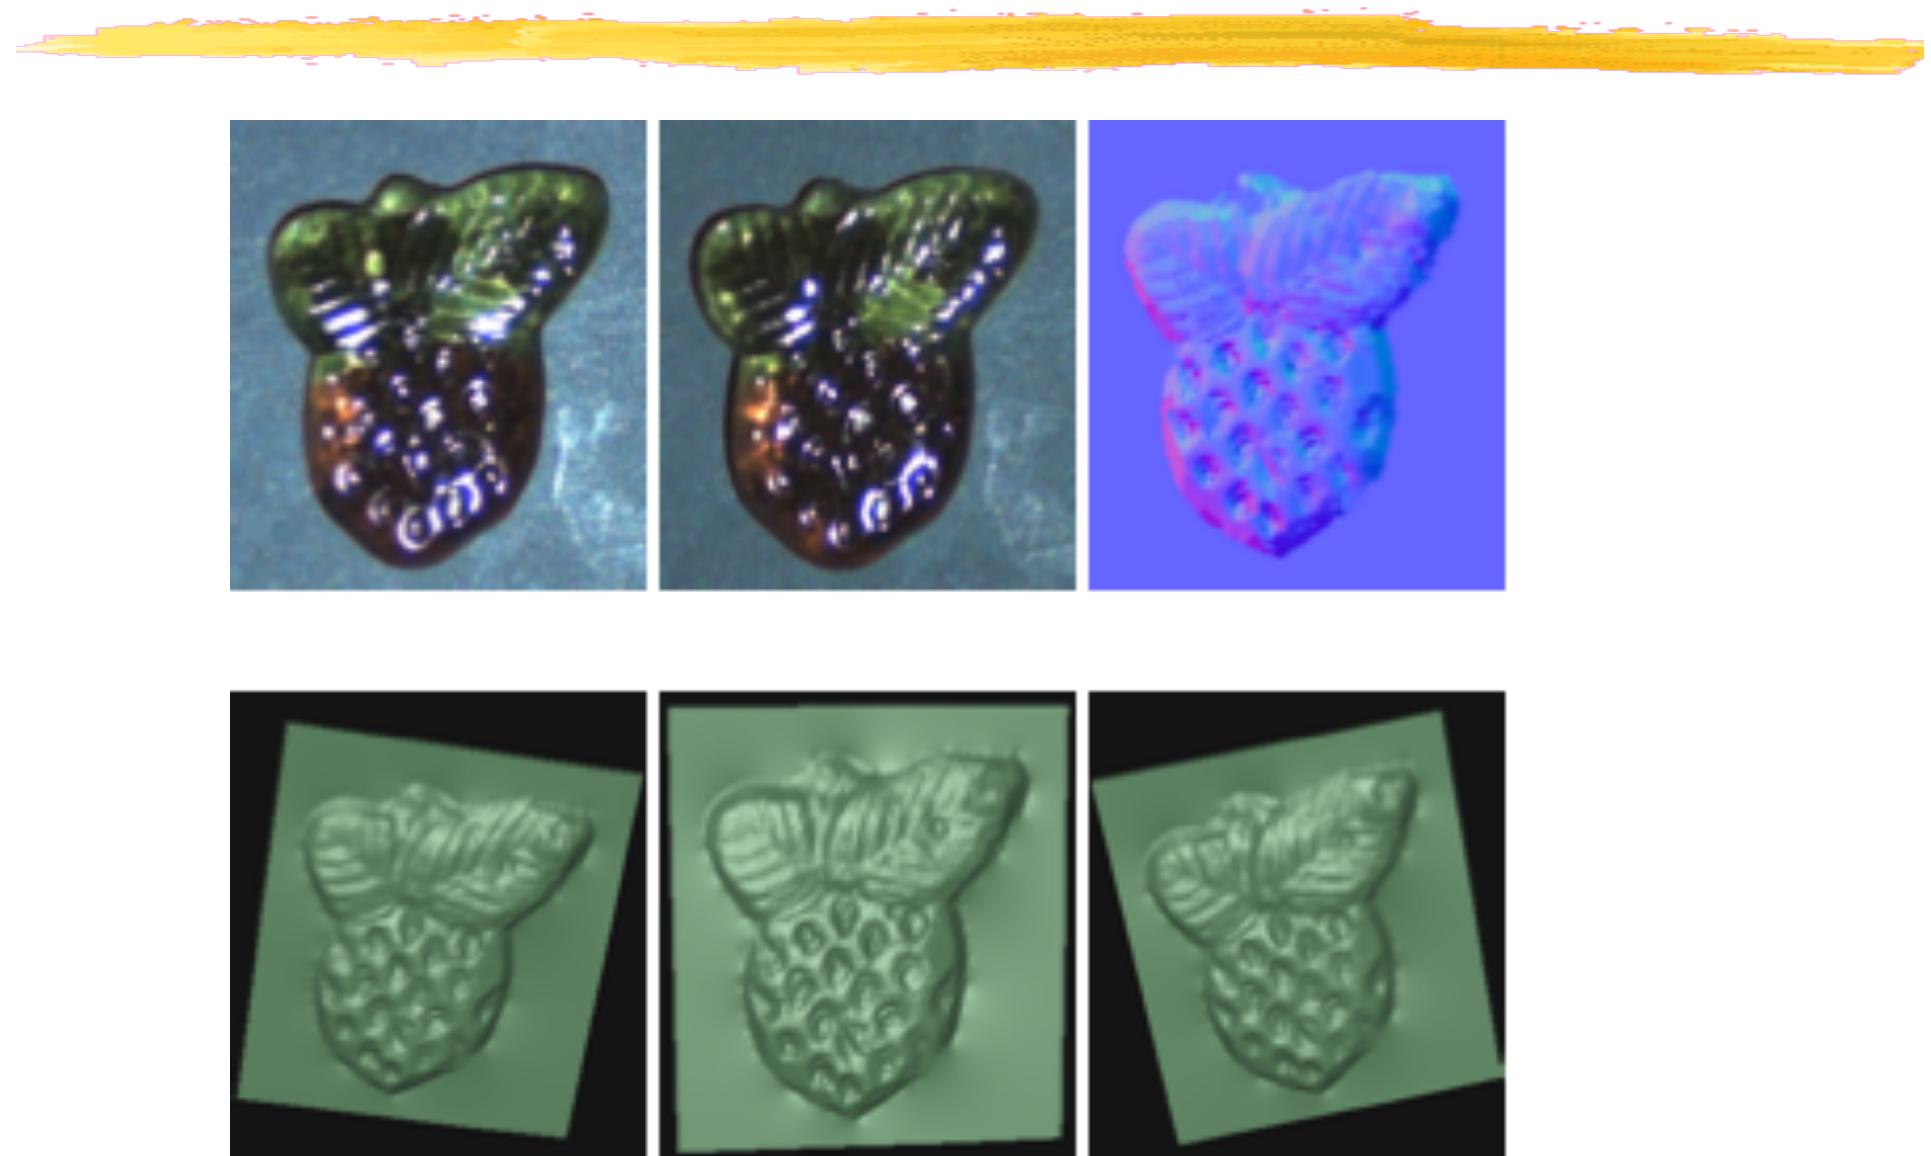
→ In practice, most surfaces are a mixture of specular and Lambertian.

SHAPE FROM SPECULARITIES

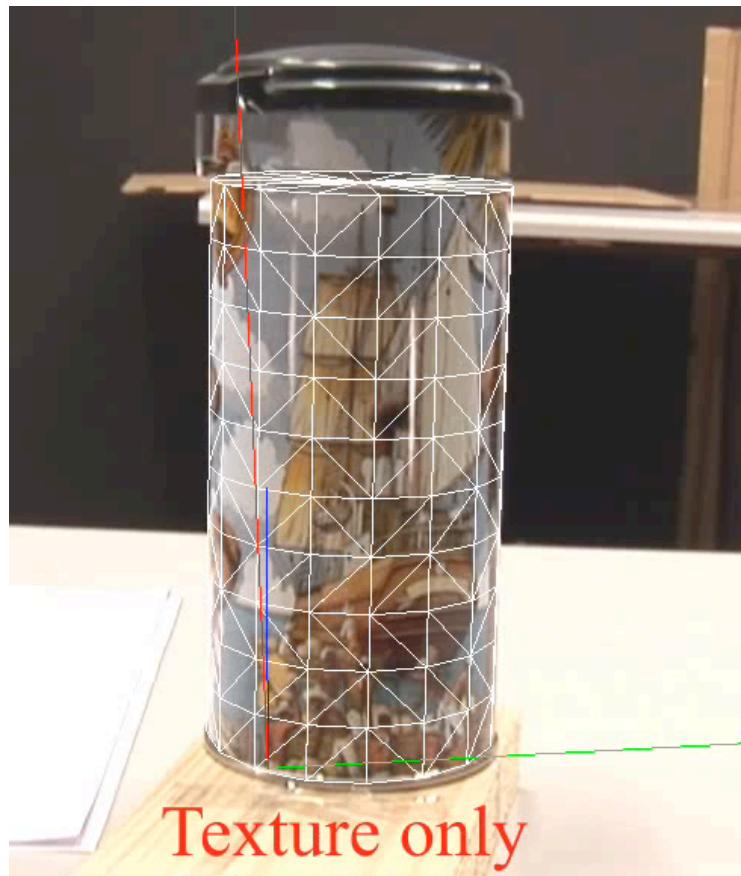


Chen et al., CVPR 2006

SHAPE FROM SPECULARITIES



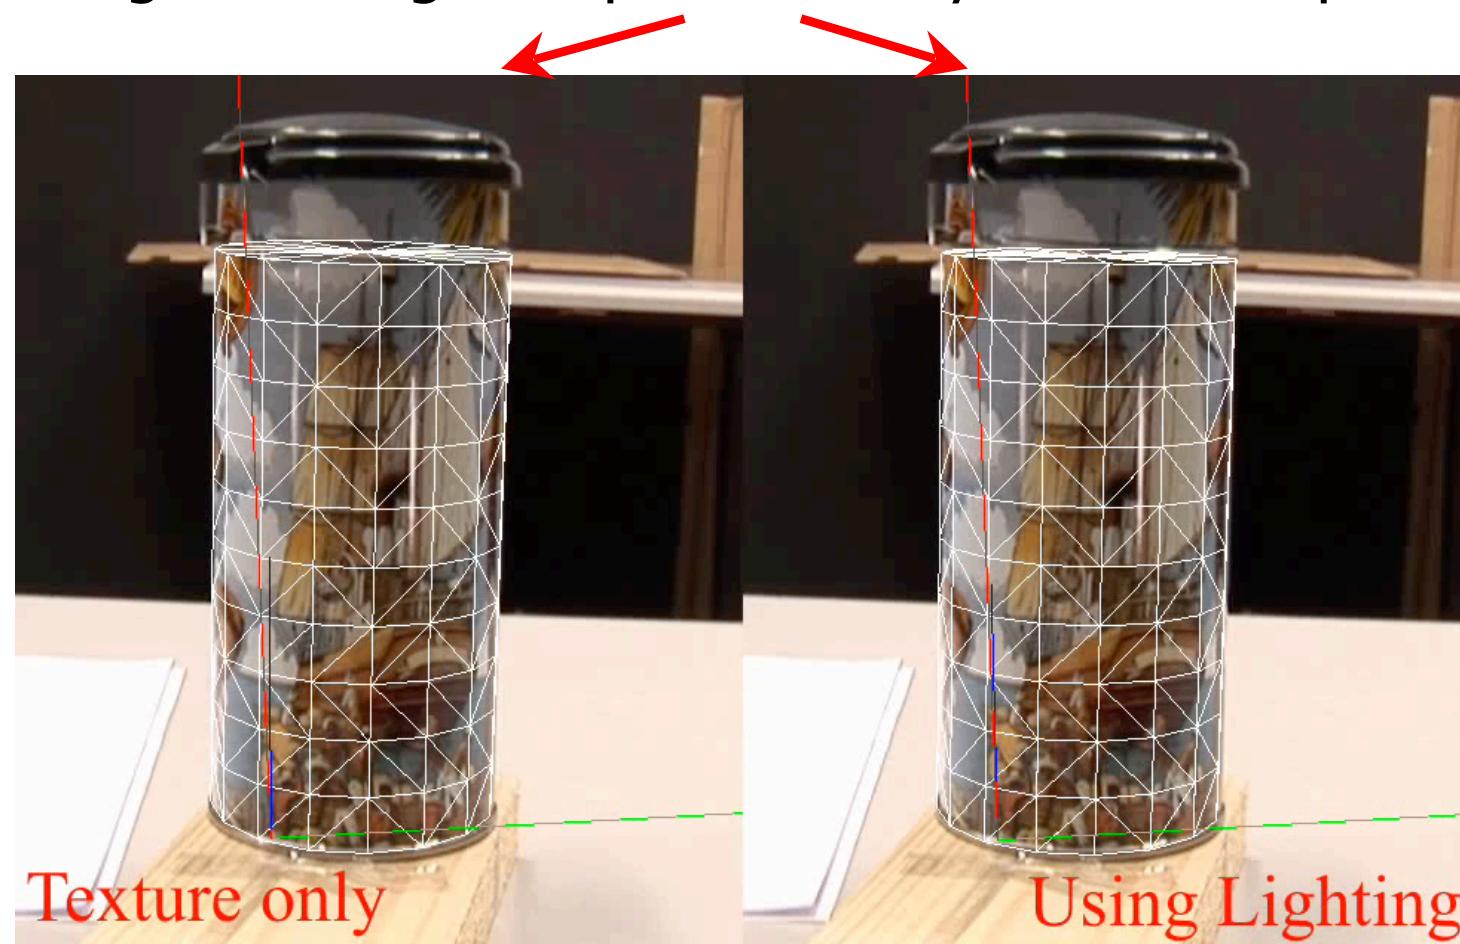
POSE FROM TEXTURE



- ❖ Texture tracking yields correct reprojection of frontal facets.
- ❖ But the jittering top shows that the tracking is nevertheless inaccurate.

POSE FROM SPECULARITIES

Taking advantage of specularities yields better poses



→ No more jittering top!

TEXTURE AND SPECULARITIES

Specularities are:

- Very sensitive to motion
- Affected differently than texture by the same motion

Specularities are not:

- Capable of constraining all the degrees of freedom

Therefore texture and specularities must be combined



IN SHORT



Traditional Shape-from-Shading requires making strong assumptions:

- Constant or piece-wise constant albedo
 - No inter-reflections
 - No shadows
 - No specularities
- In a single image context, it is most useful in conjunction with other information sources.
- stereo works best
when albedo is
changing